



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Escola Tècnica Superior d'Enginyeria
de Telecomunicació de Barcelona



Channel Knowledge Map Prediction with Deep Learning for 6G UAV-Enabled Networks

Bachelor's Degree Thesis

Submitted to the Faculty at the

Escola Tècnica Superior

d'Enginyeria de Telecomunicació de Barcelona

of the Universitat Politècnica de Catalunya

by

Guillem Moreno Garcia

In partial fulfillment

of the requirements for the

DEGREE IN TELECOMMUNICATION TECHNOLOGIES AND SERVICES ENGINEERING

Advisors: Sergi Abadal, Evgenii Vinogradov

Barcelona, May 2026

Resum

Aquest TFG estudia la predicció de mapes de coneixement de canal per a xarxes 6G amb UAVs. A partir del conjunt de dades CKM (66 ciutats reals, 16,180 mostres, mapes urbans de 513×513 píxels amb màscares LoS/NLoS i altura d'antena contínua entre 12 i 478 m) es prediuen pèrdua de trajecte, dispersió temporal / delay spread i dispersió angular / angular spread sota un protocol estricte de validació per ciutats completes. Després de provar U-Nets, cGANs, experts i priors físics, el treball convergeix en un model híbrid: priors interpretables calibrats per pèrdua de trajecte i spreads, més un model final amb cap GMM residual conjunt. En els resultats finals, aquest cap s'avalua mitjançant la predicció determinista de la mitjana de la mescla, no com una densitat predictiva calibrada.

El resultat final no és un prior pur ni una xarxa neuronal pura: el model final usa deep learning com a corrector multitasca condicionat per la geometria urbana, l'altura del UAV, les màscares LoS/NLoS i els mapes prior. En el split final de test de 14 ciutats no vistes, els priors congelats assoleixen 1.94 dB, 28.10 ns i 13.74 graus RMSE, mentre que el model final amb residuals ancorats als priors els millora fins a 1.65 dB, 26.56 ns i 11.39 graus. Això fa que els priors siguin una baseline forta per si mateixa, no només un pas intermedi: concentren la calibració per règims de densitat urbana, altura d'edificis, visibilitat i altura d'antena, i deixen al model neuronal només el residual local. El resultat és una estratègia de divideix i venceràs més que una nova llei universal de propagació.

Resumen

Este TFG estudia la predicción de mapas de conocimiento de canal para redes 6G con UAVs. A partir del conjunto de datos CKM (66 ciudades reales, 16,180 muestras, mapas urbanos de 513×513 píxeles con máscaras LoS/NLoS y altura de antena continua entre 12 y 478 m) se predice pérdida de trayecto, dispersión temporal / delay spread y dispersión angular / angular spread bajo un protocolo estricto de validación por ciudades completas. Tras probar U-Nets, cGANs, expertos y priors físicos, el trabajo converge en un modelo híbrido: priors interpretables calibrados para pérdida de trayecto y spreads, más un modelo final con cabeza GMM residual conjunta. En los resultados finales, esa cabeza se evalúa mediante la predicción determinista de la media de la mezcla, no como una densidad predictiva calibrada.

El resultado final no es un prior puro ni una red neuronal pura: el modelo final usa deep learning como corrector multitarea condicionado por la geometría urbana, la altura del UAV, las máscaras LoS/NLoS y los mapas prior. En el split final de test de 14 ciudades no vistas, los priors congelados alcanzan 1.94 dB, 28.10 ns y 13.74 grados RMSE, mientras que el modelo final con residuales anclados a los priors los mejora hasta 1.65 dB, 26.56 ns y 11.39 grados. Esto convierte a los priors en una baseline fuerte por sí misma, no solo en un paso intermedio: concentran la calibración por regímenes de densidad urbana, altura de edificios, visibilidad y altura de antena, y dejan al modelo neuronal solo el residual local. El resultado es una estrategia de divide y vencerás más que una nueva ley universal de propagación.

Summary

This thesis studies channel knowledge map prediction for 6G UAV-enabled networks. Using the CKM dataset (66 real cities, 16,180 samples, 513×513 urban maps with LoS/NLoS masks and a continuous UAV antenna height between 12 and 478 m), the project predicts path loss, delay spread, and angular spread under a strict city-holdout evaluation protocol, where entire cities are reserved for validation and test. After testing U-Nets, cGANs, expert models, and physical priors, the project converges on a hybrid design: calibrated interpretable priors for path loss and spreads, plus a joint residual GMM-head model. In the final reported results, that head is evaluated through its deterministic mixture-mean prediction, not as a calibrated predictive density.

The final system is neither a pure prior nor a pure neural network: the final model adds deep learning as a multi-task residual corrector conditioned on urban geometry, UAV height, LoS/NLoS masks, and prior maps. On the final 14-city unseen test split, the frozen priors achieve 1.94 dB, 28.10 ns, and 13.74-degree RMSE, while the final prior-anchored residual GMM-head model improves them to 1.65 dB, 26.56 ns, and 11.39-degree RMSE. This makes the priors a strong baseline in their own right, not just an intermediate step: they concentrate the calibration over urban-density, building-height, visibility, and antenna height regimes, leaving the neural model to learn only the local residual. The result is a divide-and-conquer strategy rather than a new universal propagation law.

For all those who failed throughout their journey.

Acknowledgements

I would like to express my sincere gratitude to my supervisors, Sergi Abadal and Evgenii Vinogradov, for their guidance throughout this project. Although I worked on this thesis with a high degree of independence, their advice was extremely valuable. I am especially grateful to Evgenii Vinogradov, whose propagation-domain perspective and suggestions, such as encouraging me to study the two-ray model, helped me tremendously in shaping and improving this work.

I would also like to thank ETSETB for creating the PARS programme, which allows students to begin the Master's Degree in Telecommunications Engineering before completing their Bachelor's degree. This opportunity was especially useful while working on this thesis: starting the master's before finishing the degree gave me a broader technical perspective and helped me approach the project with more maturity, depth, and ambition.

Finally, I would like to thank my family and friends for their patience and encouragement throughout the thesis.

Revision history and approval record

Revision	Date	Author(s)	Description
1.0	13/04/2026	GMG	Document creation
1.1	18/04/2026	SA, EV	Initial scope and work-plan review
1.2	22/04/2026	GMG	Scope updates after supervisor feedback
2.0	29/04/2026	GMG	First skeletal version
2.1	02/05/2026	SA, EV	Review of methodology and evaluation plan
2.2	04/05/2026	GMG	Methodology and results sections revised after review
3.0	11/05/2026	GMG	Complete technical draft
3.1	18/05/2026	GMG	Reduced version
3.2	21/05/2026	EV, SA	First correction of the final text

DOCUMENT DISTRIBUTION LIST

Role	Surname(s) and Name
[Student]	Moreno Garcia, Guillem
[Project Supervisor 1]	Abadal, Sergi
[Project Supervisor 2]	Vinogradov, Evgenii

Written by:		Reviewed and approved by:	
Date		Date	Pending review
Name	Moreno Garcia, Guillem	Name	Abadal, Sergi; Vinogradov, Evgenii
Position	Project Author	Position	Project Supervisors

Contents

Summary	2
Acknowledgements	5
Revision history and approval record	6
Contents	7
List of figures	10
List of tables	11
Abbreviations	13
1 Introduction	16
1.1 Work goals	16
1.2 Research questions, answers and contributions	17
1.3 Thesis outline	18
1.4 Requirements and specifications	18
1.5 Methods and procedures	19
1.6 Work plan	19
2 State of the Art	21
2.1 General Framework: CKM Surrogates and Evaluation Contracts	21
2.2 Path Loss Prediction	22
2.2.1 Free space and two-ray propagation	22
2.2.2 Empirical and semi empirical models	22
2.2.3 UAV and air to ground channel models	22
2.2.4 Spread definitions used later	23
2.2.5 Deep learning for path-loss maps	23
2.2.6 Physics-informed and hybrid path loss methods	25
2.3 Angular Spread Prediction	26
2.3.1 Empirical models for spread parameters	26
2.3.2 Deep learning for wideband channel statistics	26
2.4 Delay Spread Prediction	27
2.4.1 Distributional mismatch and the spike-plus-exponential problem	27
2.5 Training and Fair Comparison	28
2.5.1 Heteroscedastic uncertainty	28
2.5.2 Gaussian mixture model heads	28
2.5.3 Stochastic diffusion models	28
2.5.4 Multitask and joint prediction	29
2.5.5 Training stabilisation and regularisation techniques	29
2.5.6 Key benchmarks and fair comparison	30
2.5.7 Emerging techniques and implications for this thesis	32
3 Methodology	33
3.1 Dataset and split	33

3.2	General Framework and Data Flow	35
3.3	Final priors: overview and literature basis	36
3.3.1	High level architecture	36
3.3.2	Prior design rules	38
3.3.3	Literature ingredients and thesis adaptations	39
3.4	Path Loss Prior and Residual Prediction	40
3.4.1	LoS branch: enhanced two-ray model	40
3.4.2	Calibration of the LoS two-ray model	46
3.4.3	NLoS branch	49
3.4.4	Evaluation metric: pixel-weighted RMSE	60
3.4.5	Observed path loss-prior results	60
3.5	Angular Spread Prior and Residual Prediction	61
3.6	Delay Spread Prior and Residual Prediction	62
3.6.1	Physical motivation for log-domain modeling	62
3.6.2	Pixel-wise notation for the spread prior fit	62
3.6.3	Raw spread prior in log-domain	62
3.6.4	Calibrated spread prior	63
3.6.5	Spread prior results	69
3.7	Training and Evaluation Protocol	71
3.7.1	Experimental protocol and reporting status	71
3.7.2	Prior Anchored GMM Residual Model	71
3.7.3	Architecture: Shared UNet with FiLM	73
3.7.4	Composite Optimization Objective	76
4	Results	79
4.1	Final test contract and frozen priors	79
4.1.1	LoS path-loss prior progression	79
4.2	Path Loss Results	80
4.2.1	RMSE by UAV transmitter height	81
4.2.2	Path loss prior by topology class	82
4.2.3	Path loss: prior versus HARP-NET CKM	83
4.3	Angular Spread Results	84
4.4	Delay Spread Results	84
4.5	State of the Art Context and Runtime	85
4.5.1	Target vs final result	85
4.5.2	Comparison with closest state of the art metrics	85
4.5.3	Inference time sanity check	87
4.6	Qualitative Analysis	88
4.7	Limitations	88
5	Sustainability Analysis and Ethical Implications	90
5.1	Sustainability matrix	90
5.2	Environmental impact	90
5.3	Economic and social impact	91
5.4	Ethical implications and SDGs	92

6	Conclusions and Future Work	93
6.1	Conclusions	93
6.1.1	What the final result means	93
6.1.2	Answers to the research questions	94
6.1.3	Main contributions	95
6.1.4	Practical implications and scope	95
6.1.5	Negative results	95
6.2	Future Directions	96
	Bibliography	97
A	Development Phases and Attempt Trace	101
A.1	Expanded phase notes	101
A.1.1	From image translation to receiver-valid evaluation	102
A.1.2	Why the prior returned	102
A.1.3	Why Try 80 is the endpoint	102
A.1.4	What was moved out of the main chapters	103
B	Representative Final Model Panels	104
C	CKMGenerator Interface and Reproducibility Tool	108
C.1	Purpose and source record	108
C.2	Interface and execution modes	108
C.3	Input contract and grid fitting	112
C.4	Mask, prior and neural-model path	112
C.5	Outputs and performance controls	113
D	Additional Results and Qualitative Diagnostics	114

List of Figures

1.1	Project's Gantt diagram	20
2.1	Direct and ground-reflected paths in the two-ray model.	22
3.1	Empirical UAV transmitter height density read directly from the CKM HDF5 database using the 3 m histogram of Eq. 3.1, linearly interpolated between bin centres for display. Most training support lies in the low and mid-altitude bands.	34
3.2	Overall prior-anchored pipeline with bounded DNN corrections.	38
3.3	Two-ray geometry.	41
3.4	Effective ground-reflection amplitude as a function of UAV height.	49
3.5	FSPL versus coherent two-ray component.	49
3.6	Pipeline for the hybrid path-loss prior.	61
3.7	Feature grouping for the calibrated spread prior.	66
3.8	Pipeline for the delay and angular-spread priors.	70
3.9	Schematic of the final large-capacity shared UNet architecture with FiLM conditioning. The backbone and decode logic are shared, while each task uses its own heads and frozen prior.	75
3.10	Detailed sinusoidal height conditioning path.	75
3.11	Detailed prior-anchored GMM-head used by the final model.	76
3.12	Composite final loss family	78
4.1	Path loss prior by topology class.	83
B.1	Final residual GMM-head inspection panel for an open, low-antenna validation sample in Segovia.	105
B.2	Final residual GMM-head inspection panel for a mixed, mid-antenna validation sample in Nice.	106
B.3	Final residual GMM-head inspection panel for a dense, mid-antenna held-out test sample in Vancouver.	107
C.1	CKMGenerator Streamlit controls.	109
C.2	CKMGenerator application output after running one HDF5 sample.	110
C.3	CKMGenerator exported visual files and directory structure.	111
C.4	CKMGenerator metadata and numerical array archive.	112
D.1	Compact additional diagnostics. The GMM plot exposed different LoS/NLoS marginals and motivated distribution-aware heads; the quantisation panel records the target-storage limitation discussed in Chapter 4.	114

List of Tables

1.1	Main CKM dataset fields used in the thesis.	18
2.1	Approximate ICASSP 2023 outdoor RMSE scale.	24
2.2	Evaluation condition comparison across representative methods.	30
3.1	City-holdout assignment used by the distribution-first and final neural models.	34
3.2	External framework inputs, internally computed maps, neural model inputs and final outputs.	35
3.3	Term by term interpretation of the main frozen prior quantities.	38
3.4	Compact design rules extracted from the prior experiments.	39
3.5	Literature to design mapping used by the final prior-anchored model. . .	40
3.6	Term-by-term inventory of the LoS prior used by the final model. Unlike the NLoS branch, this is not a long pixel-feature regression vector; it is a height-interpolated two-ray calibration plus a radial residual table. . . .	46
3.7	Representative smoothed deployed LoS two-ray parameters by height bin.	48
3.8	Term-by-term inventory of the NLoS calibration vector.	54
3.9	Fixed constants used in the raw spread prior of Eq. (3.64).	63
3.10	Term by term inventory of the direct terms in the spread calibration vector.	65
3.11	Interaction terms in the spread calibration vector.	66
3.12	Delay spread: calibrated test RMSE (ns) by topology class.	69
3.13	Angular spread: calibrated test RMSE (°) by topology class.	70
3.14	Active loss settings in the final residual GMM-head model stages.	77
4.1	LoS path-loss prior progression.	79
4.2	Final hybrid path-loss prior, pixel-weighted RMSE on the held-out validation+test split.	80
4.3	Calibrated spread prior test results, pixel-weighted RMSE.	81
4.4	Path loss prior pixel-weighted RMSE by UAV transmitter height, on the held-out validation+test split.	81
4.5	Delay-spread and angular-spread calibrated pixel-weighted RMSE by UAV transmitter height, on the city-holdout test split.	82
4.6	Hybrid path-loss prior RMSE by topology class.	82
4.7	HARP-NET CKM test results on unseen cities.	84
4.8	HARP-NET CKM test RMSE separated by LoS and NLoS pixels. The test cities are unseen during training.	84
4.9	Validation-versus-test sanity check.	85
4.10	Plain language comparison with closest path loss metrics.	86
4.11	Full-pipeline runtime and ray tracing comparison.	87
4.12	Approximate RMSE floor from target quantisation.	89
5.1	Sustainability matrix summary for the thesis.	90

A.1	Compact attempt log grouped by methodological phase.	101
A.2	Decisive attempt milestones kept from the detailed attempt log.	102
B.1	Representative final-model panel manifest.	104
C.1	CKMGenerator artifacts and controls retained in the thesis appendix. . .	113

Abbreviations

2D	two-dimensional
3D	three-dimensional
3GPP	Third Generation Partnership Project
6G	sixth-generation mobile network
A2G	air-to-ground
ABS	aerial base station
AS	angular spread
BT	bachelor thesis
cGAN	conditional generative adversarial network
CKM	channel knowledge map
CNN	convolutional neural network
COST-231	European Cooperation in Science and Technology 231
CPU	central processing unit
CUDA	Compute Unified Device Architecture
DirectML	Direct Machine Learning
DL	deep learning
DNN	deep neural network
DPM	Dominant Path Model
DS	delay spread
DS-RPP	distribution shift radio path loss prediction
EMA	exponential moving average
ETSETB	Escola Tecnica Superior d'Enginyeria de Telecomunicacio de Barcelona
FiLM	feature-wise linear modulation
FLOPs	floating-point operations
FM-RME	foundation model-empowered radio map estimation
FR3	Frequency Range 3
FSPL	free-space path loss
GAN	generative adversarial network
GMM	Gaussian mixture model
GPU	graphics processing unit
GRETST	Grau en Enginyeria de Tecnologies i Serveis de Telecomunicacio
GT	ground truth
HDF5	Hierarchical Data Format version 5
HF	high-frequency
HPC	high-performance computing
ICASSP	IEEE International Conference on Acoustics, Speech and Signal Processing

IEEE Institute of Electrical and Electronics Engineers
IRT Intelligent Ray Tracing
ITU International Telecommunication Union
JSON JavaScript Object Notation
KL Kullback-Leibler divergence
LoS line of sight
MAE mean absolute error
MATLAB Matrix Laboratory
MDN mixture density network
ML machine learning
MLP multi-layer perceptron
mmWave millimetre wave
MSE mean squared error
NLL negative log-likelihood
NLoS non-line of sight
NMSE normalized mean squared error
NPZ NumPy zipped archive
OLS ordinary least squares
PARS Programa Academic de Recorregut Successiu
PCT Parallel Computing Toolbox
PDE partial differential equation
PINN physics-informed neural network
PL path loss
PMHHNet path-loss map network with high-frequency and height awareness
PMNet path-loss map prediction network
PNG Portable Network Graphics
PSNR peak signal-to-noise ratio
RGB red-green-blue
RM3D RadioMap3DSeer
RMa rural macro
RMSE root mean squared error
RQ research question
RQ1 research question 1
RQ2 research question 2
RQ3 research question 3
RSSI received signal strength indicator
RT ray tracing
RTX NVIDIA RTX graphics-processing platform
Rx receiver
SDG Sustainable Development Goal

SOA	state of the art
SSIM	structural similarity index measure
SWA	stochastic weight averaging
TR	Technical Report
TTA	test-time augmentation
Tx	transmitter
U-Net	U-shaped convolutional encoder-decoder network
U2G	UAV-to-ground
UAV	unmanned aerial vehicle
UMa	urban macro
UMi	urban micro
USC	University of Southern California
UxNB	UAV-mounted radio access node
VQGAN	vector-quantized generative adversarial network
VRAM	video random-access memory
WiCo-PG	Wireless Channel Foundation Model for Pathloss Map Generation
WINNER	Wireless World Initiative New Radio
YAML	YAML Ain't Markup Language

Introduction

Reliable radio coverage prediction is a central requirement for future 6G systems in which UAVs may act as aerial base stations, relays, or temporary coverage nodes. In 3GPP terminology, a UAV mounted radio access node is a UxNB, defined as a radio access node on board a UAV [1]. A classical ray-tracing simulator can produce detailed channel knowledge maps, but this is too expensive to repeat every time the transmitter height, the city environment, or the deployment assumptions change. In the 6G literature, a channel knowledge map is a site-specific, location-tagged repository of channel information that can support environment-aware planning and reduce online channel-acquisition overhead [2, 3]. This thesis studies whether HARP-NET CKM, a height-aware residual prior surrogate, can predict dense CKM target maps directly from urban geometry while preserving enough physical structure to generalize to cities not seen during training.

The project is formulated as dense image-to-image regression on 513×513 pixel city maps. The final model receives a building-height topology map, a LoS/NLoS visibility mask, the UAV antenna height, and the frozen prior maps described later in the thesis. In the CKM experiments the visibility mask is read from the dataset; for topology only deployment it can be generated geometrically from the building map and UAV height. The predicted outputs are path loss, delay spread, and angular spread. The neural network therefore does not infer visibility from scratch; it uses visibility as one of the physical support inputs and learns the residual radio map structure around the calibrated priors.

HARP-NET CKM is the final Height-Aware Residual Prior Network for Channel Knowledge Map prediction. The name reflects its central design choice: calibrated physical/statistical priors explain the stable propagation structure, while a height-conditioned neural network learns the remaining residual corrections.

1.1 Work goals

The initial technical targets were defined in physical units and were kept as the main success criteria throughout the project. The proposal also included a runtime objective, because a surrogate is only useful if it is clearly faster than generating the same maps through ray tracing:

- path-loss RMSE below 5 dB;
- delay-spread RMSE below 50 ns;
- angular-spread RMSE below 20° ;
- separate reporting for global, LoS, and NLoS regions;
- a 10 to 100 $\times$ reduction in channel map generation time compared with ray tracing.

These numerical targets were set in the project proposal and kept unchanged throughout the thesis. They are meant to match the order of magnitude at which a ray tracing

surrogate starts being useful for planning and coverage decisions. A few dB of path-loss RMSE is compatible with practical link budget margins, and tens of ns of delay spread and tens of degrees of angular spread stay well inside typical multipath spread values reported by 3GPP TR 38.901 for UMa/UMi scenarios. Chapter 4 checks the runtime target explicitly by comparing the measured final generator time with the measured MATLAB ray-tracing time on the same Barcelona geometry and hardware.

The most important methodological objective was not only to obtain a low average RMSE, but to do so under a strict city-holdout split. City-holdout means that entire cities, not individual pixels or tiles, are reserved for validation and testing. A test city therefore contains no samples that were seen during training. This is harder and more deployment-relevant than random pixel or random tile splitting, because the model must work on urban layouts it has never seen.

1.2 Research questions, answers and contributions

The thesis is organised around three research questions, which also define the three main scientific claims of the final pipeline:

- **RQ1.** Can a single deep learning surrogate predict per-pixel path loss, delay spread, and angular spread on 513×513 channel knowledge maps, under strict city-holdout evaluation, with errors compatible with the thesis targets?
- **RQ2.** Which part of the prediction should be left to a calibrated analytic prior and which part should be left to the neural network, so that the hybrid model is both accurate and interpretable?
- **RQ3.** How does the continuous UAV transmitter height (12 m to 478 m) affect the prediction error, and is the same architecture able to cover the full height range without per height retraining?

In simple terms, the final answer to the three questions is positive within the fixed CKM simulation setting used in this thesis. HARP-NET CKM predicts path loss, delay spread, and angular spread on cities that were not used for training. The reason it works is not that the neural network learns all propagation physics from zero. Stable effects such as distance, visibility, UAV height, and city type are first captured by calibrated priors; the neural network then corrects the remaining local errors. The hardest cases are still low-altitude and strongly blocked NLoS regions, so the results are always reported by regime instead of only as one global RMSE. Chapter 4 and the conclusions give the full technical answer.

The main contributions corresponding to these questions are:

- A reproducible CKM prediction pipeline for ground users served by a centred UAV transmitter. The evaluation uses 513×513 maps, city-holdout train/validation/test splits, valid ground receiver pixels only, and continuous UAV altitude conditioning.
- A train calibrated path-loss prior based on enhanced two-ray LoS propagation and ordinary least-squares (OLS) NLoS calibration over environment features. On held-out cities this prior already reaches 1.75 dB LoS and 3.40 dB NLoS RMSE.
- Train calibrated delay-spread and angular-spread priors based on log-domain ridge regression over geometry, visibility, UAV height, and environment features.

- HARP-NET CKM, a shared neural model that receives the frozen PL, DS, and AS priors and learns bounded residual corrections for the three maps jointly.
- A distribution shape analysis showing that earlier point regression networks tended to average away difficult NLoS and spread tail cases. This motivates using mixture style residual predictions instead of a single direct regression value everywhere.

1.3 Thesis outline

Chapter 2 reviews the state of the art and the technical background needed for the thesis: classical propagation models, radio map deep learning, hybrid physics/neural methods, distribution-aware prediction, and fair comparison criteria. Chapter 3 describes the final methodology: the CKM dataset, the ground receiver evaluation contract, the city-holdout split, the calibrated priors, and the HARP-NET CKM residual model. Chapter 4 reports the final quantitative results, the comparison with the frozen priors and the closest published benchmarks, qualitative examples, runtime, and limitations. Chapter 5 discusses sustainability, economic impact, social impact, and ethical scope. Chapter 6 answers the research questions and presents future-work. The appendices preserve the development trace, representative panels, and generator documentation.

1.4 Requirements and specifications

The project uses the CKM HDF5 dataset CKM_Dataset_270326.h5. The CKM dataset contains 66 real city scenarios and 16180 samples. Each sample is a 513×513 pixel raster with the following fields:

Table 1.1: Main CKM dataset fields used in the thesis.

Field	Shape	Type	Role
topology_map	513^2	float32	Building height in metres. Pixels with buildings are not valid receiver locations.
los_mask	513^2	uint8	Binary LoS/NLoS visibility mask.
path_loss	513^2	uint8	Ray-traced path loss in dB, quantized at approximately 1 dB.
delay_spread	513^2	uint16	Delay spread target, stored as integers and evaluated in ns after decoding.
angular_spread	513^2	uint8	Angular spread target, stored as integers and evaluated in degrees after decoding.
uav_height	scalar	float32	UAV transmitter height, used for conditioning and physical priors.

The physical receiver scenario is a ground user served by the UAV. Therefore, losses and metrics are computed only on valid ground receiver pixels. Building pixels remain in the topology input because they affect propagation, but they are excluded from error accumulation because no receiver is placed inside a building pixel. The fixed receiver height used in the analytic priors is 1.5 m. The transmitter is the UAV antenna, whose height varies continuously across the dataset.

1.5 Methods and procedures

The final method follows a hybrid pipeline. First, the city map is represented as a building-height raster with a centred UAV transmitter and a fixed ground user receiver grid. Second, geometric LoS/NLoS masks and calibrated priors are prepared for path loss, delay spread, and angular spread. Third, HARP-NET CKM receives the topology, masks, UAV height, and prior maps and predicts bounded residual corrections for the three outputs. The final maps are the prior predictions plus the learned residual corrections.

Earlier neural models are still useful scientifically, but mainly as diagnostics. They showed that direct image-to-image prediction can produce plausible-looking maps while missing NLoS and spread tail behaviour. For this reason the main text presents the final method first; the chronological attempt trace is kept in the appendix.

1.6 Work plan

The executed work plan can be grouped into five phases: literature and simulation contract review, dataset and split definition, calibrated prior construction, final residual model training, and final evaluation/writing. Figure 1.1 shows these phases at project level rather than listing all individual experiments.

The main deviation from the original proposal was methodological rather than administrative. The project moved from a generic image-to-image model toward a hybrid CKM generator because the dataset showed strong radial, visibility, and distributional structure that should not be left entirely to a black-box neural network.




Figure 1.1: Gantt diagram of the project grouped by the main thesis phases: background review, dataset contract, prior construction, final residual model, evaluation, and writing.

State of the Art and Technical Background

Radio map prediction sits at the intersection of classical wireless propagation modelling, statistical channel analysis, and dense computer vision regression. This chapter reviews the background needed to understand the final design: classical and empirical propagation models, deep learning for radio map estimation, physics informed and hybrid neural approaches, distribution-aware prediction, and fair comparison rules. The goal is to identify the limitations of existing work and the tools that will be used later, not to report the results of this thesis.

2.1 General Framework: CKM Surrogates and Evaluation Contracts

The term channel knowledge map (CKM) refers to a site-specific database of location-tagged channel knowledge, such as channel gain, angles, or delay information, used to make future radio systems environment-aware [2, 3]. This framing is more precise than a generic radio map: the map is not only an image-like coverage estimate, but a queryable prior that can reduce channel-state acquisition overhead, guide beamforming, support planning, and provide initial channel knowledge before online measurements are available.

Recent CKM literature also clarifies why the thesis uses a large offline ray-traced dataset rather than sparse field samples. Xu and Zeng [4] analyse how CKM construction error depends on spatial sample density and on the number of neighbouring samples used for online prediction, showing the explicit cost/accuracy trade off of measurement driven CKM construction. For a thesis whose goal is to learn dense 513×513 targets under city-holdout, this motivates the simulation first route: it provides full supervision everywhere, while any learned model is still best understood as a surrogate for a ray-traced CKM rather than as a replacement for field calibration.

The data ecosystem has also expanded since the first path loss challenge. RadioMap3DSeer [5] and the ICASSP 2023 challenge [6] provide the closest open benchmark family for path loss image prediction, while the directive antenna dataset of Jaensch et al. [7] and CKMImageNet [8] show the broader move toward open, location specific datasets with richer channel and environment representations. The CKM dataset used in this thesis sits in the same data driven tradition, but differs in its continuous UAV height conditioning, real city-holdout split, and simultaneous path loss, delay spread, and angular-spread targets.

2.2 Path Loss Prediction

2.2.1 Free space and two-ray propagation

The most fundamental path loss term is the *free-space path loss* (FSPL), which depends only on carrier frequency f and Tx-Rx distance d :

$$\text{PL}_{\text{FSPL}} = 20 \log_{10}(d) + 20 \log_{10}(f) + 20 \log_{10}\left(\frac{4\pi}{c}\right). \quad (2.1)$$

FSPL is the single direct-ray reference baseline in unobstructed conditions, not a lower bound on the final path loss once coherent reflections are present. In aerial scenarios a specular ground-reflection can produce constructive or destructive interference periodically with distance. Let d_{2D} be the horizontal Tx-Rx distance, h_{tx} the UAV height, and h_{rx} the receiver height. The direct and reflected path lengths are

$$d_{\text{dir}} = \sqrt{d_{2D}^2 + (h_{\text{tx}} - h_{\text{rx}})^2}, \quad d_{\text{ref}} = \sqrt{d_{2D}^2 + (h_{\text{tx}} + h_{\text{rx}})^2}. \quad (2.2)$$

With wavelength λ , the coherent two-ray model adds the direct and ground-reflected fields before converting to dB [9, 10]:

$$E_{\text{total}} = \frac{e^{-j2\pi d_{\text{dir}}/\lambda}}{d_{\text{dir}}} + \Gamma \frac{e^{-j2\pi d_{\text{ref}}/\lambda}}{d_{\text{ref}}}, \quad (2.3)$$

where Γ is the effective ground-reflection coefficient. The received field magnitude $|E_{\text{total}}|$ is then converted to a path loss like dB prediction with a fitted calibration constant. This is the physical origin of the radial rings later exploited by the path-loss prior.




Figure 2.1: Direct and ground-reflected paths in the two-ray model.

2.2.2 Empirical and semi empirical models

Empirical models such as the Okumura-Hata and COST-231/Hata formulas extend FSPL by incorporating frequency, antenna heights, and city-type correction terms [9, 11]. While computationally cheap, they are calibrated on specific frequency bands and city/environment types, limiting accuracy when the geometry departs from the fitting dataset. The 3GPP Technical Report 38.901 [12] formalises stochastic channel models for frequencies from 500 MHz to 100 GHz for urban macro (UMa), urban micro (UMi), rural macro (RMa), and indoor scenarios. It provides separate path-loss expressions for LoS and non-LoS (NLoS) conditions, breakpoint distances, and shadow-fading standard deviations, which serve as reference baselines throughout this thesis.

2.2.3 UAV and air to ground channel models

Classical terrestrial channel models assume the transmitter is at rooftop height or below. Unmanned aerial vehicle (UAV) base station scenarios introduce a third dimension: the UAV can fly from near ground level to several hundred metres. Al-Hourani et al. [13]

showed that the LoS probability increases monotonically with elevation angle and depends on city/environment type (suburban, urban, dense urban, high-rise). The A2G modelling literature [13, 14] commonly parameterises this dependence through logistic functions of the elevation angle $\theta = \arctan((h_{\text{tx}} - h_{\text{rx}})/d_{2D})$, where h_{tx} and h_{rx} are transmit and receive heights respectively and d_{2D} is the horizontal distance. 3GPP TR 38.901 [12] is used here as standardized LoS/NLoS large-scale channel model context, not as the direct source of that elevation-logistic A2G formula.

Recent work by Saboor and Vinogradov [15] makes this height dependence especially relevant for the present thesis: in simulated 26 GHz urban A2G links, LoS path loss remains close to a free space exponent, whereas NLoS path loss exponents and shadow fading evolve with ABS height and with urban layout. Vinogradov et al. [16] also propose 3D shadow projections as a fast way to generate spatially consistent A2G LoS maps without full point by point ray tracing. These results support two design choices used throughout this work: treating the LoS/NLoS mask as a physical input rather than as a cosmetic label, and conditioning every dense prediction on UAV height.

These works motivate a useful separation for the thesis methodology: LoS/NLoS visibility and continuous UAV height should be treated as physical conditioning variables, not as decorative labels appended to a generic image regression network.

2.2.4 Spread definitions used later

Beyond path loss, full channel characterisation requires second-order statistics. The *root-mean-square delay spread* τ_{rms} measures the time domain dispersion of multipath arrivals and drives inter-symbol interference in wideband systems. The *azimuth angular spread* σ_{AS} at the receiver quantifies the spatial diversity of arriving paths.

For multipath components with powers P_ℓ , delays τ_ℓ , and azimuth angles ϕ_ℓ , the RMS delay spread is

$$\bar{\tau} = \frac{\sum_\ell P_\ell \tau_\ell}{\sum_\ell P_\ell}, \quad \tau_{\text{rms}} = \sqrt{\frac{\sum_\ell P_\ell (\tau_\ell - \bar{\tau})^2}{\sum_\ell P_\ell}}. \quad (2.4)$$

For angular spread, the mean angle must be computed circularly:

$$\bar{\phi} = \text{atan2}\left(\sum_\ell P_\ell \sin \phi_\ell, \sum_\ell P_\ell \cos \phi_\ell\right), \quad \sigma_{\text{AS}} = \sqrt{\frac{\sum_\ell P_\ell \text{wrap}(\phi_\ell - \bar{\phi})^2}{\sum_\ell P_\ell}}, \quad (2.5)$$

where $\text{wrap}(\cdot)$ maps angle differences to the $(-\pi, \pi]$ interval.

3GPP TR 38.901 [12] models both quantities as log-normal random variables conditioned on the LoS/NLoS state and distance. WINNER II follows the same broad log-domain large-scale parameter convention [17]. In practice, the distributions of delay spread and angular spread are right-skewed with a heavy exponential tail and often include a spike near-zero representing the dominant LoS path. A Gaussian model cannot represent this shape faithfully. This is why the methodology later checks the empirical target histograms before choosing distribution-aware heads.

2.2.5 Deep learning for path-loss maps

Image-to-image translation baselines

The radio map prediction problem can be cast as an image-to-image regression: given an input topology map (and optionally a transmitter location channel), the model predicts

a path-loss map of the same spatial extent. The pix2pix framework [18] established a general recipe for such tasks using a conditional GAN with a U-Net generator and a PatchGAN discriminator. Early iterations of this thesis used this formulation. In practice the adversarial objective sharpens spatial texture but does not address the height-dependent radial pattern or the heavy NLoS tails that dominate the error budget in the CKM dataset.

RadioUNet [19] is the canonical convolutional path-loss map baseline before the PMNet challenge family. In its accurate-map settings it reports grey-level RMSE values around 0.0203–0.0384; using the paper’s 80 dB grey-level conversion gives an approximate 1.6 dB–3.1 dB path-loss scale. RadioGUNet [20] extends this outdoor RadioMapSeer line with group-equivariant convolutions, reporting 1.304 dB on the easiest DPM setting, 1.936 dB on the IRT setting, and 1.392 dB with cars. These values are used again in Chapter 4 as scale references for the final path-loss result.

Limitations: These papers are useful path-loss scale references, but they do not report dense delay spread or angular spread maps and do not use the CKM city-holdout, continuous UAV height contract.

PMNet and the ICASSP outdoor challenge family

PMNet [21] was the winner of the first ICASSP 2023 Outdoor Pathloss Radio Map Prediction Challenge [6]. It uses a deep encoder-decoder with skip connections, trained and evaluated on the RadioMap3DSeer dataset [5]. RadioMap3DSeer consists of city maps derived from urban 3D geometry with rooftop transmitters placed 3 m above the selected rooftop in the dataset description, and building interior pixels masked to zero error.

The challenge metric is RMSE in normalised image units. As an approximate dB-scale reference, the RadioMap3DSeer image mapping clips path gain to the range $[PL_{\text{thr}}, PL_{\text{max}}] = [-111, -75]$ dB, yielding a 36 dB dynamic window under that interpretation. Multiplying normalised RMSE by 36 gives the rough dB-scale references listed in Table 2.1.

Table 2.1: Approximate dB-scale RMSE references for ICASSP 2023 outdoor challenge methods. The normalised scores are taken from the challenge family and converted with the 36 dB RadioMap3DSeer image window described in the dataset paper [6, 5].¹

Method	RMSE (norm.)	Approx. RMSE (dB)
PMNet [21]	0.0383	1.38
PPNet baseline [6]	0.0507	1.83

Limitations: These values are scale references only. They are not directly comparable to CKM because RadioMap3DSeer uses a different dataset, fixed transmitter placement, a much narrower target image window, and a different building pixel convention. The CKM task later evaluated in this thesis uses unseen real city geometries, native 513×513 maps, continuous UAV altitude, and metrics over valid ground receiver pixels. The challenge includes multiple transmitter locations, but it does not test the continuous UAV height

¹This conversion should be read as an approximate scale reference, not as a like for like physical RMSE. It uses the path-gain clipping window $[PL_{\text{thr}}, PL_{\text{max}}] = [-111, -75]$ dB associated with the normalised target image, so the physical span is $-75 - (-111) = 36$ dB. Multiplying the normalised RMSE values reported in the challenge paper [6, 5] by this factor converts them to physical dB, which gives $0.0383 \times 36 \approx 1.38$ for PMNet.

conditioning used in CKM.

Transformers, state-space models, and attention-based encoders

The RMTransformer [22] evaluates a hybrid MaxViT/CNN encoder–decoder on the USC/PMNet dataset rather than on the ICASSP challenge test set. It reports a normalised RMSE of 0.007148 versus 0.01046 for PMNet under a random 90/10 split on 256×256 maps. The paper normalises received power/path loss from -254 to 0 dBm, so the reported pixel RMSE corresponds to roughly 1.82 dB on its own scale, and the reported channel prediction error of 0.008099 corresponds to roughly 2.06 dB. This supports the use of attention for long-range urban context, but it is not inserted in Table 2.1 because it is not measured on RadioMap3DSeer or under the ICASSP challenge contract.

Recent attention and foundation-model variants extend the effective receptive field of radio map predictors [23, 24]. This is useful when large-scale urban context affects the prediction, but these methods normally need more data and compute than convolutional models and their reported benchmarks do not match the CKM UAV height contract.

Foundation models for radio maps

The most recent shift in the literature is towards *wireless channel foundation models* pre-trained on large heterogeneous datasets before task-specific fine-tuning. WiCo-PG [23] uses dual VQGANs and a transformer with frequency-guided mixture-of-experts. It exploits RGB drone camera images as an auxiliary modality and achieves an NMSE of 0.012 on its constructed multi-modal U2G benchmark. RadioLAM [25] targets fine-grained 3D radio map estimation from ultra-low sampling rates using a large generative model pipeline. FM-RME [24] frames radio map estimation as a foundation modelling problem with self-supervised pre-training on simulated propagation data.

Limitations: These foundation approaches are relevant as long-term background, but they require training corpora and compute budgets beyond a single CKM dataset. They are therefore cited as future context, not as methods directly compared with the final thesis model.

2.2.6 Physics-informed and hybrid path loss methods

Geometry-assisted and diffraction-aware deep learning

Purely data-driven models must implicitly learn diffraction geometry from the topology map. An alternative is to pre-compute a physics-based channel prior and provide it as an additional input channel, reducing the learning problem to residual estimation.

Shadow-projection methods [16] occupy an intermediate point between stochastic LoS-probability models and full ray tracing. They use 3D building geometry to project blocked ground regions from the aerial transmitter, producing a deterministic LoS map that is spatially consistent along user trajectories. The CKM dataset already provides a ray-traced LoS mask, so this thesis does not need to generate the mask itself; nevertheless, the method is directly relevant because it explains why a binary visibility channel can carry so much of the structure that the neural network otherwise has to infer from the height map alone.

Geometry-assisted diffraction features [26] use virtual obstacles, multi screen knife-edge propagation, and scattering-aware local geometry cues to inject physically meaningful NLoS information into a deep model. This is closely related to the motivation behind

the diffraction-prior experiments in this thesis, although the implementation used later is a lightweight per-map channel rather than the full geometry model-assisted learning framework of that paper. The limitation is that a single precomputed diffraction cue cannot replace a calibrated NLoS model when dense urban shadowing dominates.

Physics-informed neural networks

The ReVeal architecture [27] operates as a physics-informed neural network (PINN) by incorporating a second-order PDE residual as an auxiliary loss term. This penalises predictions with inconsistent local spatial curvature under the derived propagation constraint, reducing corner-transition artefacts. ReVeal achieves 1.95 dB RMSE in rural scenarios using only 30 training points by leveraging the physical constraint as a strong regulariser. A lightweight analogue for dense radio maps is a masked Laplacian penalty on the residual prediction, which penalises locally non-smooth outputs without solving the full ReVeal operator.

Prior plus residual correction

A principled hybrid approach is to *freeze* a calibrated analytic prior and train the neural component to predict only the structured residual. This design reduces the effective learning problem from the full path-loss range to a narrower residual distribution, improving data efficiency and interpretability.

Chapter 3 applies this philosophy to the three thesis targets. For path loss, an enhanced two-ray model fitted per height bin provides a prior map $\hat{y}_{\text{PL}}(\mathbf{x}) = \text{TwoRay}(d, h_{\text{tx}})$, and the neural component predicts the residual $\epsilon = y - \hat{y}_{\text{PL}}$. For delay spread and angular spread, the same prior-residual idea can be implemented with a ridge-regression model in log space using TX-centred geometric features and local morphology statistics. The important background point is the decomposition itself: a frozen train-calibrated prior defines what the neural model should not have to rediscover.

2.3 Angular Spread Prediction

2.3.1 Empirical models for spread parameters

Standard channel models characterise delay spread and angular spread through log-domain large-scale parameters tied to propagation state and scenario. TR 38.901 [12] provides scenario-specific tables and formulae for these parameters, with dependencies on variables such as distance, frequency, terminal heights, and LoS/NLoS state across UMa, UMi, and RMa scenarios. These parameterisations are limited to the morphology class and frequency band of the underlying measurement campaigns.

Log-domain modelling is physically motivated: spread parameters are multiplicative quantities in physical space, so their logarithm is more Gaussian-like and additive-noise assumptions are better justified. The calibrated spread prior adopts this convention by working in $\log(1 + y)$ space throughout, consistent with the 3GPP empirical modelling framework.

2.3.2 Deep learning for wideband channel statistics

Work on deep learning for delay and angular spread is sparser than for path loss. Most published methods either predict spread globally per-sample (a single scalar) rather than per-pixel, or treat it as an ancillary output of a multitask path-loss model. Per-pixel dense

prediction of delay spread and angular spread over 513×513 maps is, to the best of the author’s knowledge, not addressed by published methods at the time of writing.

The closest external quantitative references for spreads are therefore scalar channel-characteristic predictors rather than dense CKM-map benchmarks. Huang et al. [28] are not used in the results comparison tables because their Wireless InSite campus dataset predicts scalar A2G link characteristics from coordinates and link features, not dense CKM maps from city topologies.

Limitation for this thesis. Most reviewed angular-spread work reports scalar link-level channel parameters or simulation-specific statistics rather than dense CKM maps under city holdout. Chapter 4 therefore compares angular spread mainly against the frozen prior, the thesis target, and one or two papers with similar but not equal goals.

The main challenge is the spike-plus-exponential distribution shape of spread targets. A model trained only with mean-squared error (MSE) collapses to the conditional mean, which lies in the interior of the spike and the exponential tail and is therefore a poor representative of either regime. This motivates Gaussian-mixture or other distribution-aware heads, especially when a log-domain prior is available as an anchor before residual learning.

2.4 Delay Spread Prediction

Delay spread has the same dense-map gap as angular spread, but the physical interpretation is temporal rather than directional. The target is small for many LoS-dominated receivers and can become large in streets or blocks with multiple reflected and diffracted paths. This makes the quantity sensitive to rare high-delay pixels, so a mean-error metric can hide the tail behaviour that matters for wideband links.

Limitation for this thesis. Most reviewed delay-spread work reports scalar link-level channel parameters or simulation-specific statistics rather than dense 513×513 CKM maps under city-holdout. Chapter 4 therefore compares delay spread mainly against the frozen prior, the thesis target, and one or two papers with similar but not equal goals.

2.4.1 Distributional mismatch and the spike-plus-exponential problem

The spike-plus-exponential shape, consisting of a Dirac-like concentration near-zero (the LoS-dominated case) plus a heavy exponential tail (NLoS multipath), was confirmed empirically on the CKM dataset. Skewness exceeds 5 and excess kurtosis exceeds 40 for both delay and angular spread in several topology classes. A single Gaussian or even a symmetric loss function does not fit this family.

Mixture models are the natural solution, but the required number of components depends on what is being modelled. A histogram study can select the best family for the raw target distributions, where path loss and the spread targets may need several components to match the full marginal shape. In a prior-residual setting, however, the mixture head models a different quantity: the residual distribution after a frozen prior has already explained the dominant LoS, NLoS, delay-spread, and angular-spread structure. This distinction matters because the prior can remove much of the target side multi-modality, leaving a compact mixture to represent under correction, near-zero correction, and over correction modes.

2.5 Training and Fair Comparison

2.5.1 Heteroscedastic uncertainty

Kendall and Gal [29] established the framework for heteroscedastic neural networks in computer vision: predict both a mean output $\mu(x)$ and a log-variance $\log \sigma^2(x)$, optimising the negative log-likelihood (NLL)

$$\mathcal{L} = \frac{(\mu - y)^2}{\sigma^2} + \log \sigma^2. \quad (2.6)$$

Pixels with high aleatoric uncertainty (deep NLoS shadows, building edges) receive lower effective weight, while the network is penalised for overestimating uncertainty everywhere. A PMHNet uncertainty diagnostic applied this loss to path loss prediction, with σ^2 predicted as a second output channel of PMHNet.

2.5.2 Gaussian mixture model heads

Beyond heteroscedastic Gaussians, the GMM-head allows multi-modal output distributions. Mixture-density networks have already appeared in wireless channel modelling: Saleh et al. [30] use an MDN to predict probabilistic mmWave path loss distributions, Garcia Marti et al. [31] learn a stochastic Gaussian-mixture channel model for physical-layer design, and Lee et al. [32] derive conditional received-power PDFs with a mixture-density network. These works support the use of mixture outputs for wireless propagation, but they operate at the level of scalar or channel-sample distributions rather than dense radio map generation.

A dense CKM version of a GMM-head can be implemented with a two-stage architecture. Stage A produces per-sample GMM parameters $\{\pi_k, \mu_k, \sigma_k\}_{k=1}^K$ via a convolutional encoder conditioned on UAV height; Stage B produces per-pixel component assignments $p_k(i, j)$ and within-component deviations $z(i, j)$. The prediction is

$$\hat{y}(i, j) = \sum_{k=1}^K p_k(i, j) \cdot (\mu_k + z(i, j) \cdot \sigma_k). \quad (2.7)$$

This formulation avoids the mode-collapse failure of MSE regression while maintaining a compact architecture. The important idea is therefore not the isolated use of a Gaussian mixture, but its role as a global distribution head that tells the dense decoder what value distribution the map should contain before the model solves the harder pixel-placement problem.

2.5.3 Stochastic diffusion models

Conditional diffusion models applied to radio maps [33] generate physically plausible stochastic variations in deep shadow zones rather than predicting a single blurry mean. The iterative denoising procedure can be conditioned on topology maps, transmitter locations, and obstacle prompts. RadioDiff reports lower normalised RMSE and higher SSIM/PSNR than CNN and sequence-model baselines on static and dynamic RadioMapSeer settings, but its metrics are image-normalised rather than physical dB values. These models are computationally expensive at inference time and are listed as a future-work horizon for this thesis.

2.5.4 Multitask and joint prediction

Training a single shared backbone across multiple output targets can improve sample efficiency and consistency between targets. Path loss, delay spread, and angular spread are physically correlated: a deep shadow corresponds to high delay spread and potentially wide angular spread from diffracted paths. A joint model can therefore use the intermediate feature maps learned for path loss to partially explain spread patterns, reducing the effective learning problem.

Chapter 3 uses this idea by sharing the topology and UAV height representation across the three targets, while allowing target-specific residual heads for path loss, delay spread, and angular spread. The literature point is that joint training is useful only if the shared representation is anchored by target-specific priors and losses; otherwise the easiest target can dominate the common backbone.

2.5.5 Training stabilisation and regularisation techniques

The methods in this section are included because dense radio map prediction is easy to overfit. The model sees complete city maps with strong spatial correlations, while validation and test cities are completely held-out. Small training subsets, heavy-tailed NLoS errors, and quantized targets can make the training loss improve even when the deployment-relevant city-holdout error does not. Stabilisation and regularisation techniques are therefore not cosmetic training tricks; they are safeguards against selecting a model that only fits the training cities.

Stochastic weight averaging

Stochastic weight averaging (SWA) [34] averages model weights along the training trajectory at a low, approximately constant learning rate. This finds wider optima in the loss landscape that generalise better to unseen data. The CKM city-holdout evaluation makes generalisation particularly important since validation and test cities are never seen during training. In this thesis, SWA was useful as a diagnostic stabilisation tool in earlier neural experiments, although it is not part of every final training recipe.

Exponential moving average weights

Separate from SWA, exponential moving average (EMA) tracking of the model weights at a decay rate of 0.9975 can produce smoother validation checkpoints. Earlier diagnostic branches used EMA weights for validation scoring and fed a smoothed score to the ReduceLROnPlateau scheduler. This is useful background for interpreting the training choices in Chapter 3, where the final checkpoint selection rule is stated explicitly for reproducibility.

Feature-wise linear modulation

Feature-wise linear modulation (FiLM) [35] conditions the activations of a neural network on a scalar or vector input by applying learned affine transforms $\gamma(c) \cdot \mathbf{h} + \beta(c)$ to intermediate feature maps. Sinusoidal positional encoding of the UAV height scalar (using log-spaced frequencies over the full 12 to 478 m range) is fed into FiLM layers at multiple encoder depths. This ensures that the UAV altitude continuously modulates the activation statistics rather than being handled as a discrete bin, which is critical when the effective propagation regime changes substantially with height.

2.5.6 Key benchmarks and fair comparison

Why headline numbers cannot be compared directly

The most widespread mistake in radio map literature reviews is comparing headline RMSE values without verifying whether the underlying tasks are equivalent. The following evaluation axes are the minimum required for a valid comparison:

1. **Split:** city-holdout vs. in-distribution random split vs. same-scene split.
2. **Transmitter height:** fixed rooftop, discrete bins, or continuous UAV altitude.
3. **Building pixels:** excluded from loss, excluded from metrics, set to zero error, or included.
4. **Target quantity and units:** path loss (dB), path gain (dB), received power (dBm), delay spread (ns), or angular spread (degrees).
5. **Dynamic range and clipping:** determines the RMSE floor and ceiling independently of model quality.
6. **Inference extras and calibration source:** test-time augmentation (TTA), SWA, EMA weights, and whether calibration is train-only, validation-based, or fitted from test/field measurements.

This thesis uses a strict city-holdout split, continuous UAV height, valid ground receiver metrics (building pixels excluded from loss and metrics), native 513×513 pixel maps, and no test-data or field calibration after the train-calibrated priors and the model under evaluation are frozen. Calibration is therefore not shown as a binary column in Table 2.2: the relevant question is not whether a method has fitted parameters, but whether it adapts to the evaluation scene using validation, test, or field measurements. Table 2.2 summarises the structural axes that can be compared compactly.

Table 2.2: Evaluation condition comparison across representative methods.

Method	Unseen city	Cont. UAV height	≥ 512 px	Comparable?
PMNet / ICASSP 2023 [21]	No	No	No	No
RMTransformer [22]	No	No	No	No
RadioGUNet [20]	Partial	No	No	Partial
AIRMap [36]	Partial	No	No	Partial
Gao et al. [37]	No	No	Partial	Partial
Tarhouni et al. [38]	No	No	No	Physical ref.
PathFinder [39]	Partial	No	No	Partial
This thesis task contract	Yes	Yes	Yes	Reference contract

Two consequences follow from this checklist. First, method rankings can change substantially when the split changes from random maps to unseen cities. A model that has learned the texture, street orientation, and height statistics of a known city may still fail when the same transmitter height range is evaluated on a new morphology. For this reason, this thesis treats city-holdout as a core part of the task definition rather than as an optional robustness test.

Second, path loss, delay spread, and angular spread should not be treated as three copies

of the same regression problem. LoS path loss often contains a strong deterministic radial component, whereas NLoS path loss depends on blocked-link geometry and local morphology. Delay and angular spreads are even more distributional: many pixels are near-zero, but physically important tails remain. A fair comparison therefore has to report the target, regime, masking policy, and split together; otherwise an apparently lower RMSE can simply mean that an easier quantity was measured.

Tarhouni et al.: measured suburban path loss prediction

Tarhouni et al. [38] study machine-learning-based path loss prediction in a sub-6 GHz suburban campus. It is useful because the errors are in physical dB and tied to a concrete environment, but it is terrestrial, lower-frequency, and not a dense UAV CKM with delay- and angular-spread targets. Its role here is qualitative: real-environment generalisation tends to move back toward multi-dB errors, which supports keeping city-holdout as part of the task definition.

AIRMap: adaptive digital twin radio maps

AIRMap [36] is an adaptive U-Net autoencoder processing fixed 200×200 elevation-map tensors with variable spatial resolution (2.5–15 m/pixel, corresponding to 500 m–3 km per side) for ultra-low-latency digital-twin synchronisation. It reports sub-4 dB RMSE against ray-traced path-gain maps, and a separate lightweight calibration step using 20 % of field measurements reduces the median measurement-domain error to about five percent. Its strong numbers are not directly comparable to CKM because the target is path gain rather than absolute path loss, the input is a single elevation channel, and the transmitter configuration is not the continuous UAV height setting used in this thesis. Its reported ~ 4 ms inference time on an NVIDIA L40S is nevertheless a strong deployment reference: the value of a CKM surrogate is not only accuracy, but producing dense maps fast enough for planning loops.

Gao et al.: multi-layer segmentation with corridor weighting

Gao et al. [37] propose a ResNet-based outdoor path loss predictor with Tx depth maps, Rx depth maps, a distance map, and a weighting map that emphasises the direct Tx-Rx corridor. Using the ICASSP 2023 outdoor data, Table III of the paper reports their method at 5.59 dB RMSE, compared to 8.09 dB for PPNet, 9.56 dB for RPNet, and 8.72 dB for a twelve-layer Vision Transformer baseline; Table II reports 12.19 dB RMSE on the ITU Challenge 2024 large-area subset, also ahead of the same baselines, with roughly 60% fewer FLOPs. This formalises an internal lesson of the CKM work: the important geometry is not the whole map uniformly, but the subset near the propagation corridor.

The lesson for this thesis is methodological: Tx-centred distance and obstruction features are useful but are weaker than a true per-receiver corridor description. A full Gao-style per-Rx corridor map remains future-work because generating one for every receiver pixel at 513×513 resolution is substantially more expensive than using Tx-centred distance and obstruction features.

PathFinder: multi transmitter domain adaptation

PathFinder [39] tackles distribution shift via disentangled feature encoding and mask-guided low-rank attention, reporting the lowest MSE of 0.1068 and RMSE of 0.3263 on unseen rural datasets. On the RadioMap3DSeer DS-RPP split, its normalised RMSE of 0.0331 can be approximately converted with the same 36 dB RadioMap3DSeer image-window

convention, giving about 1.19 dB. The unseen-rural score would be about 11.75 dB under the same window, but this is only an approximate scale reading because the paper and code report normalised PNG image units and do not state a separate rural physical target window. The explicit domain-generalisation focus aligns with the CKM city-holdout spirit, though PathFinder’s setup involves multi transmitter inputs and a different normalisation convention. Techniques from PathFinder such as transmitter-oriented mixup augmentation are listed as medium-priority future work. The same concern motivates the city-holdout split used in this thesis: a CKM surrogate is useful only if it works on new urban layouts, not only on shuffled samples from the same city.

ICASSP 2025 indoor challenge

The ICASSP 2025 Indoor Pathloss Challenge [40] presents a structurally different problem, dense wall-penetration NLoS effects in bounded interior spaces, but its public results are a useful warning against over-interpreting sub-2 dB outdoor challenge numbers. The official result page reports that the winning SIP2Net submission obtains a weighted-average RMSE of 9.411 dB; IPP-Net reaches 9.501 dB; TerRaIn reaches 10.325 dB; and TransPathNet reaches 10.397 dB [41, 42, 43, 44]. These are not outdoor UAV results, but they show that when NLoS physics becomes harder and less directly observable from the input map, modern neural architectures again hit a roughly 9 dB to 10 dB error band. For this thesis, the lesson is diagnostic rather than competitive: harder NLoS settings favour strong priors, regime separation, and residual distribution modelling over plain direct regression.

2.5.7 Emerging techniques and implications for this thesis

Across the benchmarks reviewed here, the recurring 9 to 10 dB error band in harder NLoS-heavy settings reflects the difficulty of inferring volumetric propagation from a 2D building-height map. The most plausible routes beyond that wall are richer geometry and diffraction priors [26], PDE-style physics losses such as ReVeal [27], sparse target-domain adaptation such as RadioPiT [45], diffusion refinement [33], and larger radio map foundation models [24].

For the present thesis, the literature points to four practical design requirements. First, the comparison contract must be explicit because datasets, splits, height protocols, and target units differ widely. Second, A2G prediction requires continuous UAV height conditioning rather than a fixed transmitter height. Third, LoS/NLoS and environment-type effects should be reported separately because one global RMSE hides the hardest regimes. Fourth, distribution-aware outputs are useful when targets contain multiple modes or long tails. These observations motivate the methodology in Chapter 3.

Methodology

This chapter is structured as the route that would be followed if the project were started again with the final evidence already known. The detailed chronology is preserved in Appendix A, but the main methodology now follows the fastest defensible path: fix the evaluation contract, use the early neural models as diagnostics, identify the target distributions, calibrate frozen priors, and train a bounded residual model around those priors. The emphasis is on decisions that affect reproducibility: masking, splitting, conditioning, and the distinction between raw deep prediction and prior-anchored residual prediction.

3.1 Dataset and split

Each sample is a complete 513×513 pixel CKM. The spatial resolution is 1 m $\times$ 1 m per-pixel, so each map covers a 513 m $\times$ 513 m square. The UAV transmitter is horizontally fixed at the centre of the map, while its antenna height h_{tx} varies from sample to sample and is treated as a primary conditioning variable. Every non-building pixel is an active receiver location, with the receiver antenna placed exactly at 1.5 m above ground. Building pixels are not receiver positions and are excluded from the training loss and reported metrics rather than being counted as easy background pixels. The model never receives a random crop as an independent training city. Instead, the split is performed at city level. For the distribution-first and final neural models the split uses `split_seed = 42`, with 15% validation and 15% test samples after city assignment. This avoids the most common optimistic error in spatial radio prediction: training on one part of a city and testing on another visually similar part of the same city.

The transmitter height samples are not uniformly distributed. Comparing the stored HDF5 histogram with the tested affine beta-mixture variants shows that the closest nominal sampler is $h_{\text{design}} \approx 10 + 490X$, where

$$X \sim 0.75 \text{Beta}(3, 23) + 0.25 \text{Beta}(4, 4).$$

However, the finite HDF5 database is the relevant distribution for training and reporting. Reading the stored `uav_height` field for all 16180 samples gives an empirical range of 11.61 m to 477.54 m, usually rounded in the text as about 12 m to 478 m; 477.54 m is the largest sampled HDF5 value, not a hard endpoint of the approximate beta law. This is statistically normal: under $10 + 490X$, $P(h_{\text{tx}} > 478 \text{ m}) = 3.19 \times 10^{-5}$ per draw, so only 0.52 such samples are expected in 16180 draws and the probability of drawing none is about 0.60. A reproducible formulaic version of the realized database distribution is the 3 m empirical density. Let $\Delta h = 3 \text{ m}$, $B_k = [10 + 3k, 13 + 3k)$, and let h_i be the stored

height of sample i . Then

$$\begin{aligned} \hat{f}_{\Delta h}(h) &= d_k, \quad h \in B_k, \\ d_k &= \frac{n_k}{N\Delta h}, \quad n_k = \sum_{i=1}^N \mathbf{1}\{h_i \in B_k\}, \quad N = 16180. \end{aligned} \quad (3.1)$$

When the density is visualized, the plotted curve linearly interpolates between the bin-centre values (c_k, d_k) , with $c_k = 11.5 + 3k$. Figure 3.1 shows this realized transmitter height density directly from the stored HDF5 samples.




Figure 3.1: Empirical UAV transmitter height density read directly from the CKM HDF5 database using the 3 m histogram of Eq. 3.1, linearly interpolated between bin centres for display. Most training support lies in the low and mid-altitude bands.

The all-database height counts are 3972 samples below 50 m, 8456 from 50 m to 150 m, 2490 from 150 m to 300 m, and 1262 above 300 m. In the actual training cities, the same bins contain 2676, 5654, 1644, and only 866 samples, respectively. This matters for interpreting the height-binned results: the higher-altitude bins contain fewer training examples, so their lower RMSE later in the thesis is evidence that those cases are physically easier, not that they are better represented in the training set.

The prior artifacts use the same city-holdout contract as the final model: `val_ratio = 0.15`, `test_ratio = 0.15`, and `split_seed = 42`. The Try 78 path-loss prior and Try 79 spread priors are fitted on the 10840 training samples. Standalone prior audits can then use the full held-out validation+test pool of 5340 samples, while final model-versus-prior comparisons use the 2590-sample final test split. This keeps training cities, validation cities, and final test cities separated throughout the prior and residual-model evaluations.

Table 3.1: City-holdout assignment used by the distribution-first and final neural models.

Split	Samples	Cities	City names
Train	10840	37	Abidjan, Abu Dhabi, Alexandria, Antigua Guatemala, Athens, Bangalore, Bangkok, Bath, Berlin, Bratislava, Brugge, Buenos Aires, Canberra, Chefchaouen, Dalian, Hanoi, Ho Chi Minh City, Hyderabad, Kinshasa, Leuven, Ljubljana, Luanda, Lyon, Manaus, Minsk, Nagpur, Paris, Rio de Janeiro, Santiago, Shanghai, Siem Reap, St. Louis, Thane, Toledo, Valparaiso, Venice, Vienna.
Validation	2750	15	Bilbao, Chiang Mai, Cleveland, Dhaka, Dubrovnik, Fortaleza, Marrakesh, Mexico City, Munich, Nazareth, Nice, Salvador, Segovia, Tunis, Victoria.
Test	2590	14	Barcelona, Beijing, Carcassonne, Copenhagen, Halifax, Jaipur, Johannesburg, Key West, Kuala Lumpur, Osaka, Phuket, Pune, Surat, Vancouver.

The receiver mask is:

$$m_{\text{ground}}(i, j) = \begin{cases} 1, & \text{if topology}(i, j) = 0 \text{ (non-building receiver pixel),} \\ 0, & \text{otherwise.} \end{cases} \quad (3.2)$$

Equivalently, the CKM supervision grid contains one receiver at every non-building pixel, all at the fixed receiver height $h_{\text{rx}} = 1.5$ m. The mask only removes building

pixels from the loss and metrics; it does not subsample the remaining ground pixels. The LoS/NLoS visibility masks used later by the priors and the final residual model are geometric support inputs rather than target radio maps. In the CKM experiments they are read from the dataset when available; for topology-only deployment they can be generated from topology and UAV antenna height using geometric ray-casting, although generated masks are not expected to be bit-identical to the CKM reference masks. This makes the visibility signal explicit without requiring a previous prediction of the target radio maps. All RMSE values in this thesis are pixel-weighted over valid ground pixels:

$$\text{RMSE} = \sqrt{\frac{\sum_{s,i,j} m_s(i,j) (\hat{y}_s(i,j) - y_s(i,j))^2}{\sum_{s,i,j} m_s(i,j)}}. \quad (3.3)$$

This is stricter than averaging per-sample RMSEs because large valid maps contribute proportionally to the number of receiver pixels.

3.2 General Framework and Data Flow

With the evaluation contract fixed, the methodology used for the final thesis version is deliberately short. First, the CKM split and receiver mask are frozen so that all reported results use unseen cities and valid ground receivers only. Second, the target distributions are inspected before selecting the output model: path loss is separated into LoS and NLoS regimes, while delay spread and angular spread are treated as right-skewed, tail-sensitive quantities. Third, train-only priors are calibrated for the three targets. Finally, HARP-NET CKM learns only bounded residual corrections around those frozen priors.

Table 3.2: *External framework inputs, internally computed maps, neural model inputs and final outputs.*

Stage	Maps / variables
External inputs to the full framework	Building height map and UAV transmitter height h_{tx} .
Internally computed support maps	LoS/NLoS visibility mask and building/ground receiver mask.
Frozen prior outputs	Combined path-loss prior, path loss LoS prior, path loss NLoS prior, delay-spread prior, angular-spread prior.
Inputs to HARP-NET CKM	Building height map, UAV height conditioning, ground-mask, LoS mask, NLoS mask, combined path-loss prior, path loss LoS prior, path loss NLoS prior, delay-spread prior, angular-spread prior.
Final outputs	Path loss map, angular spread map, delay spread map.

Therefore, the full framework has two external inputs and three final outputs. The neural residual model receives ten internal inputs because the framework has already computed visibility masks, receiver masks and frozen prior maps. This is the main distinction between the deployable framework and the internal neural model: the user provides only a city height map and the UAV height, while the pipeline constructs the remaining maps before inference.

The earlier U-Net, cGAN, PMNet, PMHHNet, and distribution-first experiments are not part of the main method. They are retained in Appendix A as development evidence

because they explain why the thesis stopped expanding generic backbones and instead moved toward prior-anchored residual learning. Their practical lessons are folded into the final pipeline as constraints: no building pixel credit, no random city leakage, explicit UAV height conditioning, separate treatment of LoS/NLoS regimes, and distribution-aware residual heads when the target is multi-modal or heavy-tailed.

The rest of this chapter therefore describes the final method rather than the chronological path used to discover it. Section 3.3 summarises the frozen path loss, delay spread, and angular-spread priors; Sections 3.4 and 3.5 give their mathematical construction; and Section 3.7 describes the joint residual network trained on top of them.

3.3 Final priors: overview and literature basis

This section provides the compact entry point for the ML calibrated physical path loss and spread priors and the final joint residual GMM-head model architecture. The detailed mathematical derivations and implementation level fitting notes follow in the same methodology chapter, so the design and the audit trail remain together.

3.3.1 High level architecture

The central design decision in the final model is *not* to learn path loss, delay spread, and angular spread from scratch. Instead:

- the path-loss prior provides a frozen prior for `path_loss`,
- the spread prior module provides frozen priors for delay spread (`delay_spread`) and angular spread (`angular_spread`),
- the final residual GMM-head model learns bounded corrections around those priors.

The deterministic or interpretable structure is kept in the prior, and the deep model is used only for the residual part that the prior cannot explain. Figure 3.2 illustrates the overall pipeline.

This prior phase is not presented as a new universal closed-form propagation law. Its value is more practical: it turns physics, morphology, visibility, and height-dependent calibration into a strong auditable baseline. The final model then follows a split of responsibility: calibrated priors remove the stable component of the targets, and the neural residual model focuses on the local errors that remain.

The term *prior* is not used here in the Bayesian sense of a probability distribution over model parameters or targets, and no posterior inference is performed in the prior stage. It denotes deterministic, frozen estimator maps computed before the neural residual correction. The word *prior* is therefore used for three concrete estimators, not for a vague physical hint:


- **LoS path-loss prior.** This is the clearest physical branch: $PL_{LoS} = FSPL + C_{2ray} + b + r$. FSPL uses the 3D direct Tx–Rx distance; C_{2ray} comes from a coherent direct plus ground-reflected field-sum; b is a fitted height-dependent dB offset; and r is a smoothed radial residual table. The UAV is always centred horizontally, so the radial coordinate is unambiguous, but h_{tx} varies by sample: $\rho(h_{tx})$, $\phi(h_{tx})$, $b(h_{tx})$ and $r(d_{2D}, h_{tx})$ are fitted in 5 m height bins, smoothed, and linearly interpolated at inference. It is not a generic free-space formula and it is not a learned linear regressor.

- **NLoS path-loss prior.** This starts from the raw blocked link propagation estimate $PL^{(0)}$, which includes $PL_{\text{NLoS,raw}}$ on obstructed pixels, then calibrates it with local morphology: building-density, building-height, NLoS support, the elevation dependent shadowing proxy σ_{shadow} , elevation angle, and interactions. It uses antenna height regimes (`low_ant`, `mid_ant`, `high_ant`) because the obstruction geometry changes qualitatively when the centred UAV rises above the local skyline. It is deliberately separate from the LoS branch because blocked pixels do not follow the same two-ray mechanism.
- **Delay and angular-spread priors.** These are fitted in $\log(1 + y)$ space. The raw estimate captures propagation state, distance, elevation, and local morphology; the final ridge calibration adds the 23-feature vector of Sec. 3.5. That vector includes continuous normalized UAV height, density-weighted transmitter clearance over nearby building pixels, and the local window-area fraction occupied by buildings taller than the transmitter. This is why the spread prior can fit the near-zero LoS mass and the NLoS tail without letting the largest tail values dominate the whole calibration, unlike a native-space global average.

The notation intentionally reuses ρ in two different blocks: $\rho(h_{\text{tx}})$ is the fitted LoS reflection amplitude, whereas ρ_k is a local building-density fraction in a $k \times k$ window. They are unrelated quantities; the argument or subscript identifies which one is meant. The same convention is used for b : $b(h_{\text{tx}})$ is a LoS dB bias, $b(x)$ is the binary building mask, b_{topology} is a spread prior log-domain offset, and b_{41} is the blocker depth feature. These symbols are local to their definitions and should not be read as one shared physical variable.

Table 3.3: Term by term interpretation of the main frozen prior quantities.

Term	Prior block	Interpretation
d_{LoS}	LoS path loss	3D direct Tx–Rx distance. It sets the FSPL baseline and uses the fixed $h_{\text{rx}} = 1.5$ m receiver height.
d_{ref}	LoS path loss	Image source reflected distance. It is the length of the single ground-reflected path used in the coherent two-ray correction.
$C_{2\text{ray}}$	LoS path loss	Direct plus reflected field correction in dB. Negative values mean constructive gain over FSPL; positive values mean destructive fading.
$\rho(h_{\text{tx}})$	LoS path loss	Height interpolated effective reflection amplitude. It is fitted from CKM rather than assumed from one fixed ground material.
$\phi(h_{\text{tx}})$	LoS path loss	Height interpolated effective reflection phase. It absorbs the phase offset of the simplified one reflection geometry.
$b(h_{\text{tx}})$	LoS path loss	Height dependent dB bias remaining after the coherent two-ray fit.
$r(d_{2D}, h_{\text{tx}})$	LoS path loss	Smoothed radial residual table. It corrects distance dependent errors that remain after FSPL, two-ray interference, and height bias.
$\text{PL}^{(0)}, \text{PL}_{\text{NLoS,raw}}$	NLoS path loss	Raw hybrid formula map and its obstructed pixel branch before map morphology calibration.
ρ_k	NLoS PL and spreads	Fraction of building pixels in a $k \times k$ neighbourhood; a local density proxy.
h_k	NLoS PL and spreads	Local building-height summary in a $k \times k$ neighbourhood; a proxy for nearby reflectors and blockers.
n_k	NLoS PL and spreads	Local NLoS support fraction; indicates how much of the neighbourhood is geometrically blocked.
σ_{shadow}	NLoS path loss	Elevation dependent shadowing proxy used by the NLoS linear calibration.
$\theta, \theta_{\text{norm}}, \theta_{\text{inv}}$	PL and spreads	Elevation angle terms. Low elevation links interact with more buildings, so θ_{inv} is especially useful for spread tails.
b_{41}, c_{41}, t_{41}	Spreads	Blocker depth, transmitter clearance, and local taller-building fraction; they separate open overhead links from obstructed street canyon like links.

**Figure 3.2:** Overall prior-anchored pipeline. The frozen priors are injected into the DNN, which learns a bounded correction δ around them.

3.3.2 Prior design rules

The main lesson from the prior phase is that a useful prior is not just a physical looking equation. It is a complete estimator: the formula, fitted calibration, feature units, LoS/NLoS mask, topology features, and evaluation mask must be identical between fitting and inference. Earlier prior experiments used many of the same physical ingredients as the final path-loss prior, but they did not always keep those pieces aligned. The final

prior became strong because it made the prior self contained and auditable under the same city-holdout contract used for the final model.

The distinction matters most for NLoS. A calibration fitted for one physical formula cannot reliably correct another; local height and density features must have the same units when fitted and when applied; and regime assignment should come from the sample geometry rather than only from a broad city label. Once those constraints were enforced, the prior could separate deterministic large-scale structure from the genuinely local residual that the final model later learned.

This does not mean that all earlier NLoS failures were only implementation mismatches. The no prior and weak prior neural branches also failed because single-output point regression was a poor match for blocked, multi-modal, heavy-tailed NLoS and spread targets. That is why the thesis keeps both lessons: the final priors fix the prior-estimator problem, while the distribution-first predecessors motivate distribution-aware heads and residual uncertainty.

Table 3.4: *Compact design rules extracted from the prior experiments.*

Rule	Effect on the final pipeline
Fit and apply the same formula	The calibration corrects the actual runtime prior, not a neighbouring variant.
Keep units explicit	Distances, heights, topology values, and normalized features have one shared definition across fitting, validation, and test.
Use the metric mask during fitting	The prior is optimized for valid receiver pixels rather than for building cells that are excluded from the final metric.
Hard separate LoS and NLoS	The two regimes keep different physics, feature sets, and calibration parameters when the dataset provides a reliable mask.
Let literature define axes, not all coefficients	Distance, height, frequency, and elevation come from propagation theory; effective reflection and local morphology corrections are fitted from CKM.
Learn bounded residuals after the prior	The final model is allowed to correct local errors without overwriting the calibrated large-scale structure.
Treat calibration artifacts as part of the model	Large JSON and coefficient artifacts are not bookkeeping only; they store the height, density, visibility, and morphology-specific decisions that make the prior a competitive baseline.

3.3.3 Literature ingredients and thesis adaptations

The final pipeline is a hybrid between classical propagation models, distribution-aware regression, and dataset-specific calibration. Table 3.5 keeps the mapping explicit; the detailed equations and fitting procedures are given in the following path-loss prior, spread prior, and residual GMM methodology sections.

Table 3.5: Literature to design mapping used by the final prior-anchored model.

Block	Literature basis	CKM adaptation	New role in this thesis
LoS path-loss prior	FSPL and coherent two-ray propagation.	Effective complex reflection, height bin interpolation, and radial residual fitting on CKM LoS pixels.	Reconstructs the radial interference pattern and supplies the frozen LoS anchor.
NLoS path loss calibration	COST-231/Hata style attenuation, A2G elevation envelopes, and LoS/NLoS separation.	Map derived density, height, NLoS support, elevation, and shadowing features with regime wise calibration.	Turns a broad empirical curve into a local, topology aware NLoS prior.
Delay and angular-spread priors	Log normal spread statistics from 3GPP/WINNER style channel models and UAV elevation trends.	Log-domain ridge regression with local morphology features and a sparse regime fallback hierarchy.	Provides stable spread anchors without pretending that spreads are smooth deterministic path-loss maps.
Final residual model	U-Net style dense prediction, FiLM conditioning, probabilistic regression, and mixture density ideas.	Prior channels, explicit masks, residual caps, anchor gates, and a shared multi-task trunk.	Learns only the bounded local correction and uncertainty left after the calibrated priors.

3.4 Path Loss Prior and Residual Prediction

Reader map

This section is intentionally detailed because it is the mathematical core of the final path-loss result. A reader only interested in the modelling logic can follow it in four passes. First, the LoS branch starts from FSPL and adds an enhanced two-ray field-sum, which explains the radial ring structure observed in the CKM maps. Second, the NLoS branch uses a calibrated regime wise model over topology and obstruction features, because a LoS shaped formula cannot explain blocked streets by itself. Third, both branches are fitted only on the training split and then evaluated with the same pixel-weighted ground-mask RMSE used elsewhere in the thesis. Finally, the prior output becomes the frozen path loss anchor consumed by the final residual GMM-head model.

The corresponding quantitative audit is reported in Chapter 4; this methodology section focuses on how the prior is constructed and why its output is frozen before residual learning.

3.4.1 LoS branch: enhanced two-ray model

The LoS prior is not a pure FSPL prior. FSPL is kept because it is the direct-ray normalization term of the coherent two-ray model: it gives the path loss that would be observed if only the direct-ray existed. The deployed LoS branch then adds the fitted ground-reflection interference term, a height-dependent bias, and a smoothed radial residual profile. This differs from the NLoS branch, which uses a longer pixel-feature vector and a regime wise linear calibration.

Geometry and distances

Consider a UAV transmitter at height h_{tx} (metres above the ground-plane reference) and a receiver at height h_{rx} . For each map pixel x at 2D horizontal distance $d_{2D}(x)$ from the transmitter, two propagation distances are defined.

The *direct path* (LoS) distance is:

$$d_{\text{LoS}}(x) = \sqrt{d_{2D}(x)^2 + (h_{\text{tx}} - h_{\text{rx}})^2} \quad (3.4)$$

This is simply the Euclidean distance between the two antennas in 3D space. Because $d_{2D}(x) \geq 0$ and the squared height difference $(h_{\text{tx}} - h_{\text{rx}})^2$ is nonnegative, $d_{\text{LoS}}(x) \geq d_{2D}(x)$, with equality only when both antennas are at the same height.

The *reflected path* (image) distance is:

$$d_{\text{ref}}(x) = \sqrt{d_{2D}(x)^2 + (h_{\text{tx}} + h_{\text{rx}})^2} \quad (3.5)$$

This is the distance from the image source (the UAV mirrored below the ground plane at $-h_{\text{tx}}$) to the receiver. Since $(h_{\text{tx}} + h_{\text{rx}}) \geq |h_{\text{tx}} - h_{\text{rx}}|$ for any nonnegative heights, we always have $d_{\text{ref}}(x) \geq d_{\text{LoS}}(x)$. In the present UAV to ground setting with $h_{\text{tx}} > 0$ and $h_{\text{rx}} > 0$, the reflected path is strictly longer; in the limiting geometric case where one antenna is exactly on the ground plane, it would be equal rather than longer.




Figure 3.3: Two-ray geometry. Direct path (blue) has length d_{LoS} ; reflected path (green) has the same total length as the image source path (gray dashed), which is d_{ref} .

Free-space path loss baseline

The free-space path loss in dB is defined using the standard formula at frequency f_{MHz} (in MHz):

$$\text{FSPL}(x) = 32.45 + 20 \log_{10} \left(\frac{d_{\text{LoS}}(x)}{1000} \right) + 20 \log_{10}(f_{\text{MHz}}) \quad (3.6)$$

Derivation. FSPL in linear power ratio is $(4\pi d/\lambda)^2$ where $\lambda = c/f$. In dB:

$$\begin{aligned} \text{FSPL}_{\text{dB}} &= 20 \log_{10} \left(\frac{4\pi d f}{c} \right) \\ &= 20 \log_{10}(4\pi) + 20 \log_{10}(d) + 20 \log_{10}(f) - 20 \log_{10}(c) \\ &\approx 20 \log_{10}(d_{\text{km}} \cdot 1000) + 20 \log_{10}(f_{\text{MHz}} \cdot 10^6) + 20 \log_{10}(4\pi) - 20 \log_{10}(3 \times 10^8) \\ &= 20 \log_{10}(d_{\text{km}}) + 20 \log_{10}(f_{\text{MHz}}) + 32.45 \quad [\text{dB}] \end{aligned} \quad (3.7)$$

The constant $32.45 = 20 \log_{10}(4\pi \cdot 10^3 \cdot 10^6 / (3 \times 10^8)) \approx 32.45$ dB. In Eq. (3.6), $d_{\text{LoS}}(x)/1000$ converts metres to kilometres.

Physical meaning. FSPL grows quadratically with distance ($\propto d^2$) and quadratically with frequency. It captures only the geometric spreading of the wavefront, with no ground interaction, scattering, or diffraction.

Enhanced two-ray correction

In the two-ray model, the received field is the superposition of the direct path and a ground-reflected path. Let $k = 2\pi f/c$ be the free space wavenumber. The complex sum of the two waves at the receiver is:

$$E_{\text{total}} \propto \frac{e^{-jk d_{\text{LoS}}}}{d_{\text{LoS}}} + \Gamma \frac{e^{-jk d_{\text{ref}}}}{d_{\text{ref}}} \quad (3.8)$$

where Γ is the complex ground-reflection coefficient. In the standard textbook derivation, Γ is computed from the Fresnel equations for the ground material. Here, it is replaced by an *effective* complex parameter:

$$\Gamma(h_{\text{tx}}) = \rho(h_{\text{tx}}) e^{-j\phi(h_{\text{tx}})} \quad (3.9)$$

where $\rho \in [0, 1.5]$ is the effective amplitude (allowing over unity values during the calibration search) and $\phi \in [-\pi, \pi]$ is the effective phase offset. This range is a robustness allowance for an effective CKM-fitted parameter, not a claim that a passive Fresnel reflection coefficient exceeds unity; the deployed smoothed values reported later remain below one. Both parameters depend on the UAV-height bin.

The two-ray path loss correction term in dB is:

$$C_{2\text{ray}}(x) = -20 \log_{10} \left| 1 + \rho(h_{\text{tx}}) \frac{d_{\text{LoS}}(x)}{d_{\text{ref}}(x)} e^{-j(k(d_{\text{ref}}(x) - d_{\text{LoS}}(x)) + \phi(h_{\text{tx}}))} \right| \quad (3.10)$$

Term-by-term interpretation.

- $\rho(h_{\text{tx}})$: amplitude of the ground-reflection. A value of $\rho = 0$ recovers pure FSPL. The fitted values are $\rho \approx 0.46$ to 0.56 for heights 10 to 100 m, consistent with a moderate ground-reflection coefficient.
- $d_{\text{LoS}}(x)/d_{\text{ref}}(x)$: path amplitude ratio. This is always ≤ 1 since the reflected path is longer; it corrects for the $1/d$ amplitude decay of the reflected-ray relative to the direct ray.
- $k(d_{\text{ref}}(x) - d_{\text{LoS}}(x))$: the phase difference accumulated due to the extra path length of the reflected wave.
- $\phi(h_{\text{tx}})$: an additional phase offset capturing phase shifts upon reflection (e.g., phase shift at the boundary). The fitted values are $\phi \approx 0$ for all height bins, meaning the fitted extra phase offset is near-zero under this sign convention.
- $-20 \log_{10} |\cdot|$: the total field amplitude at the receiver, normalized to the direct-path field. When the two waves add in phase, the field amplitude is > 1 , giving a *negative*

correction (the signal is *stronger* than FSPL alone); when they cancel, the correction is *positive* (fading loss).

This spatial interference pattern creates the characteristic ring like behavior in the LoS path-loss maps: concentric rings of constructive and destructive interference around the transmitter.

Comparison with literature two-ray models

The field-sum in Eq. (3.10) is the standard *exact* coherent two-ray equation from the literature: a direct-ray plus one ground-reflected ray summed before conversion to dB. The practical adaptation in this thesis is not an extra term inside the two-ray equation, but the way the reflection coefficient is parameterized and calibrated. We refer to this fitted variant as the *enhanced two-ray* model throughout this work: the underlying equation is the textbook coherent two-ray model, but the effective reflection coefficient is enhanced beyond the Fresnel range and allowed to take amplitudes up to $\rho \leq 1.5$ in order to absorb constructive multipath effects from the surrounding urban scene that a single Fresnel coefficient could not represent. The coefficients are fitted independently in 5 m UAV-height bins and then linearly interpolated across neighboring bins at inference time.

Physical vs. effective parameters. In classic texts such as Rappaport [9] and Goldsmith [46], the complex reflection coefficient Γ is derived from the Fresnel equations for a smooth dielectric surface:

$$\Gamma_{\perp} = \frac{\sin \theta_g - \sqrt{\epsilon_c - \cos^2 \theta_g}}{\sin \theta_g + \sqrt{\epsilon_c - \cos^2 \theta_g}}, \quad \Gamma_{\parallel} = \frac{\epsilon_c \sin \theta_g - \sqrt{\epsilon_c - \cos^2 \theta_g}}{\epsilon_c \sin \theta_g + \sqrt{\epsilon_c - \cos^2 \theta_g}} \quad (3.11)$$

where θ_g is the grazing angle and $\epsilon_c = \epsilon_r - j\sigma/(\omega\epsilon_0)$ is the complex relative permittivity of the ground. Here ϵ_r is the relative dielectric constant, σ is the ground conductivity, $\omega = 2\pi f$ is the angular frequency, and ϵ_0 is the vacuum permittivity. This requires assuming physical constants for the ground (typically $\epsilon_r \approx 15$ for urban soil). The symbol is deliberately written as θ_g here because the later A2G and morphology features use $\theta(x)$ for the elevation angle of the UAV to receiver link.

In contrast, our enhanced two-ray implementation keeps the same coherent field-sum equation but replaces the angle dependent $\Gamma(\theta)$ with a height-dependent effective coefficient $\Gamma(h_{\text{tx}}) = \rho e^{-j\phi}$. We do not assume a flat dielectric plane; instead, ρ and ϕ are fitted to the CKM dataset in 5 m UAV-height bins, with linear interpolation between bins during prediction. This allows the model to capture aggregate scattering effects from the real, non flat urban terrain that a pure Fresnel calculation would miss.

The d^4 approximation. Many literature sources [9, 46] simplify the two-ray model for large distances ($d \gg h_{\text{tx}}, h_{\text{rx}}$) into the classic fourth power law:

$$P_r(d) \approx P_t G_t G_r \frac{h_{\text{tx}}^2 h_{\text{rx}}^2}{d^4} \quad (3.12)$$

Here P_t is the transmitted power, G_t and G_r are the transmitter and receiver antenna gains, and $P_r(d)$ is the received power at distance d . This approximation assumes a perfect reflection ($\Gamma = -1$) and small grazing angles. Our implementation uses the *exact* field-sum (Eq. 3.8) without the d^4 approximation. This is still a standard literature form;

it simply avoids a far field approximation that is not appropriate when UAV height is not negligible relative to horizontal distance.

Search range vs. fitted values. The calibration grid search (Sec. 3.4.2) allows $\rho \in [0, 1.5]$ for robustness. However, the smoothed deployed values reported in Table 3.7 and the stored calibration remain below unity, so the deployed model does not rely on an over unity reflected-ray. In practice, the main difference from a purely physical textbook instantiation is therefore the fitted height-binned effective coefficient, not an additional propagation term.

Extra calibrated terms outside the two-ray model. The quantities $b(h_{\text{tx}})$ and $r(d_{2D}, h_{\text{tx}})$ in Eq. (3.13) are not part of the literature two-ray formula itself. They are post fit calibration terms added on top of the standard two-ray prediction to absorb residual systematic errors.

Calibrated LoS prior

The full calibrated LoS prior is:

$$\text{PL}_{\text{LoS}}(x) = \text{FSPL}(x) + C_{2\text{ray}}(x) + b(h_{\text{tx}}) + r(d_{2D}(x), h_{\text{tx}}) \quad (3.13)$$

where:

- $b(h_{\text{tx}})$ is a height-dependent bias (dB): a constant correction that shifts the overall two-ray prediction for a given UAV height.
- $r(d_{2D}(x), h_{\text{tx}})$ is a radial residual correction: a smooth function of horizontal distance (and height bin) that captures systematic over-/under-prediction of the two-ray model as a function of range. It is obtained by smoothing the empirical mean residual in each distance bin after the two-ray fit.

The four terms in Eq. (3.13) are applied sequentially: FSPL gives the geometric baseline; the two-ray correction adds the interference pattern; the bias corrects for a systematic global offset per height; and the radial term corrects for range dependent systematic errors. The implementation then clips the frozen LoS prior to the $[20, 180]$ dB range used by the final model. For the CKM dataset used here, valid calibrated LoS values are not expected to fall outside this interval, so this clipping is a defensive range consistency guard rather than a correction that should materially change the prior.

For compact notation, the LoS branch can be summarized by the following deterministic diagnostic vector:

$$\mathbf{x}_{78}^{\text{LoS}}(x) = \begin{bmatrix} \text{FSPL}(x) \\ d_{\text{LoS}}(x)/d_{\text{ref}}(x) \\ k(d_{\text{ref}}(x) - d_{\text{LoS}}(x)) \\ \rho(h_{\text{tx}}) \\ \phi(h_{\text{tx}}) \\ C_{2\text{ray}}(x) \\ b(h_{\text{tx}}) \\ r(d_{2D}(x), h_{\text{tx}}) \end{bmatrix}. \quad (3.14)$$

This vector is not multiplied by a learned coefficient vector at inference time. Instead, it exposes the calibrated quantities used by the nonlinear two-ray expression $PL_{LoS}(x) = g_{LoS}(\mathbf{x}_{78}^{LoS}(x)) = FSPL(x) + C_{2ray}(x) + b(h_{tx}) + r(d_{2D}(x), h_{tx})$. The entries are computed as follows:

- $FSPL(x)$ is computed from the 3D direct distance $d_{LoS}(x)$ and the carrier frequency. It is the direct-ray path loss, so it sets the absolute distance and frequency scale before any reflected path correction is added.
- $d_{LoS}(x)/d_{ref}(x)$ is computed from the direct-path length and the image source reflected path length. Since the reflected path is longer, this ratio is normally below one and gives the relative amplitude of the reflected field after $1/d$ propagation loss.
- $k(d_{ref}(x) - d_{LoS}(x))$ is the geometric phase advance of the reflected-ray, with $k = 2\pi f/c$. This is the term that makes the correction oscillatory in space: small changes in extra path length can move the direct and reflected fields from constructive to destructive interference.
- $\rho(h_{tx})$ is the effective reflection amplitude used in the coherent field-sum. During calibration, each 5 m UAV-height bin is fitted by grid search over $\rho \in [0, 1.5]$ and the phase grid; for each candidate pair, the best constant bias is removed and the pair with lowest LoS RMSE is selected. The stored ρ curve is then smoothed across neighboring height bins and, at inference time, linearly interpolated to the actual h_{tx} . Thus ρ is not a Fresnel coefficient assumed from a ground material, but a CKM-calibrated effective reflected-field amplitude.
- $\phi(h_{tx})$ is fitted in the same height-binned grid search, using a uniform grid over $[-\pi, \pi)$. It absorbs the effective phase shift of the reflection and any systematic phase offset caused by the simplified one reflection geometry. As with ρ , the stored phase curve is smoothed across height bins and linearly interpolated during prediction.
- $C_{2ray}(x)$ is then computed by substituting d_{LoS}/d_{ref} , the geometric phase, $\rho(h_{tx})$, and $\phi(h_{tx})$ into Eq. (3.10). This is the actual coherent direct plus reflected correction in dB. A negative value means constructive gain relative to FSPL; a positive value means a fading loss.
- $b(h_{tx})$ is the constant dB offset fitted together with ρ and ϕ for each height bin. For a fixed candidate (ρ, ϕ) , the calibration computes the mean residual $y - (FSPL + C_{2ray})$ over the sampled LoS pixels in that height bin and uses it as the optimal bias for that candidate. The final stored b is smoothed over height and linearly interpolated at inference time. Its role is to remove the remaining height-dependent offset after the interference pattern has been fitted.
- $r(d_{2D}(x), h_{tx})$ is fitted after the smoothed two-ray model is fixed. The residual target-minus-prediction is accumulated in integer radial bins around the transmitter, separately for each 5 m height bin; empty bins fall back to a global radial profile, the profile is smoothed, and the applied correction is clipped to ± 2.5 dB. This term removes range dependent structure that remains after the coherent two-ray model and the height bias.

Stored / computed term	Mathematical role	Interpretation
$\text{FSPL}(d_{\text{LoS}}, f)$	Direct ray baseline.	Gives the geometric free-space loss before any reflected path is added. It is a component of the two-ray prior, not the final LoS model by itself.
$\rho(h_{\text{tx}})$	Height interpolated reflected-ray amplitude.	Controls how strongly the ground-reflected wave contributes relative to the direct wave.
$\phi(h_{\text{tx}})$	Height interpolated phase offset.	Captures the effective reflection phase after fitting to CKM instead of assuming a fixed Fresnel coefficient.
$C_{2\text{ray}}(x)$	Coherent interference correction.	Converts the complex direct plus reflected field-sum into an oscillatory dB correction around FSPL.
$b(h_{\text{tx}})$	Height interpolated bias.	Removes a remaining per-height offset after the coherent two-ray fit.
$r(d_{2D}, h_{\text{tx}})$	Smoothed radial residual profile, clipped to ± 2.5 dB.	Corrects systematic distance dependent residuals not captured by the two-ray equation.
$\text{clip}_{[20,180]}$	Output guard.	Keeps the frozen prior inside the physical path-loss range used by the final model.

Table 3.6: Term-by-term inventory of the LoS prior used by the final model. Unlike the NLoS branch, this is not a long pixel-feature regression vector; it is a height-interpolated two-ray calibration plus a radial residual table.

3.4.2 Calibration of the LoS two-ray model

Data preparation

All training samples are used. For each sample, valid LoS pixels are selected as those satisfying:

1. topology label = 0 (ground pixel),
2. $m_{\text{LoS}}(x) > 0$ (marked as LoS),
3. the measured path loss is finite and ≥ 20 dB.

If a sample has more than 2 500 valid LoS pixels, a random subsample of 2 500 pixels is drawn (with a fixed seed). Samples are then grouped by UAV height into height bins of 5 m width:

$$h_{\text{bin}} = 5 \cdot \text{round}\left(\frac{h_{\text{tx}}}{5}\right) \quad (3.15)$$

Each bin is further subsampled to a maximum of 8 000 pixels total.

Grid search for (ρ, ϕ)

For each height bin, (ρ, ϕ) are found by grid search:

$$\phi_{\text{grid}} = \left\{ -\pi + \frac{2\pi k}{N_\phi} : k = 0, 1, \dots, N_\phi - 1 \right\}, \quad N_\phi = 48 \quad (3.16)$$

$$\rho_{\text{grid}} = \{0.00, 0.05, 0.10, \dots, 1.50\} \quad (3.17)$$

For each pair $(\hat{\rho}, \hat{\phi}) \in \rho_{\text{grid}} \times \phi_{\text{grid}}$, and for each pixel x_i in the height bin, define the two-ray prediction without bias:

$$\hat{p}_i(\hat{\rho}, \hat{\phi}) = \text{FSPL}(x_i) - 20 \log_{10} \left| 1 + \hat{\rho} \frac{d_{\text{LoS}}(x_i)}{d_{\text{ref}}(x_i)} e^{-j(k(d_{\text{ref}} - d_{\text{LoS}}) + \hat{\phi})} \right| \quad (3.18)$$

The optimal bias for each $(\hat{\rho}, \hat{\phi})$ pair is solved analytically as the empirical mean of residuals:

$$\hat{b}(\hat{\rho}, \hat{\phi}) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{p}_i(\hat{\rho}, \hat{\phi})) \quad (3.19)$$

where y_i is the measured path loss at pixel x_i and N is the number of pixels in the bin. The biased prediction is $\hat{p}_i + \hat{b}$, and the RMSE for the pair is:

$$\text{RMSE}(\hat{\rho}, \hat{\phi}) = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{p}_i(\hat{\rho}, \hat{\phi}) - \hat{b}(\hat{\rho}, \hat{\phi}))^2} \quad (3.20)$$

The best pair is:

$$(\rho^*, \phi^*) = \arg \min_{(\hat{\rho}, \hat{\phi})} \text{RMSE}(\hat{\rho}, \hat{\phi}) \quad (3.21)$$

Post-fit Gaussian smoothing across height bins

After fitting all bins independently, a Gaussian-weighted smooth is applied to stabilize sparse high-altitude bins. Let $\{h_k\}$ be the fitted height bin centers, $\{(\rho_k^*, \phi_k^*, b_k^*)\}$ the per-bin fitted values, and $\{N_k\}$ the pixel counts used in fitting. The implementation weights each bin by $\sqrt{N_k}$ rather than by N_k itself, so very populated bins stabilize the curve without completely dominating neighboring sparse bins. The smoothed values are:

$$\tilde{\rho}(h) = \frac{\sum_k w_k(h) \sqrt{N_k} \rho_k^*}{\sum_k w_k(h) \sqrt{N_k}}, \quad w_k(h) = \exp\left(-\frac{(h - h_k)^2}{2\sigma_h^2}\right), \quad \sigma_h = 10 \text{ m} \quad (3.22)$$

and similarly for ϕ and b . The implementation also stores a suspicious-bin mask for bins where both ρ_k^* and b_k^* differ noticeably from the smoothed global trend, typically high-altitude bins with few samples. That mask is diagnostic metadata: the runtime calibration arrays used by the final residual model are the smoothed ρ , ϕ , and b curves themselves, not a mixture of raw values for ordinary bins and smoothed values only for flagged ones. At inference time the UAV height is not forced to be exactly one of the fitted 5 m bin centers. The implementation linearly interpolates $\rho(h_{\text{tx}})$, $\phi(h_{\text{tx}})$, $b(h_{\text{tx}})$, and the radial residual profile between the two nearest height bins, and clamps to the nearest

endpoint outside the fitted range. This keeps the prior continuous with respect to UAV height instead of introducing discrete jumps every 5 m.

Radial residual correction

After fitting the two-ray model, a residual correction is computed separately per height bin and integer radius:

$$r(r_{\text{px}}, h_{\text{tx}}) = \text{clip}\left(\text{GaussSmooth}_{\sigma_r}\left(\mathbb{E}_{h\text{-bin}, r_{\text{px}}}\left[\text{clip}(y - \hat{y}_{2\text{ray}}, -2.5, +2.5)\right]\right), -2.5, +2.5\right) \quad (3.23)$$

where r_{px} is the rounded 2D pixel radius from the transmitter, $\sigma_r = 1.5$ pixels is the Gaussian smoothing radius, and the ± 2.5 dB clip is applied both to individual residual samples during accumulation and to the interpolated profile used at inference. Radius bins with no samples in a height bin fall back to the global radial profile before smoothing. At inference, the profile is linearly interpolated across the stored height bins, then indexed by the pixel-radius map and added to the coherent two-ray prediction.

Smoothed deployed parameter values

Table 3.7 shows representative LoS two-ray parameters after the post fit Gaussian smoothing across height bins. These are the deployed calibration values used by the prior at inference time, not the unsmoothed per-bin grid-search optima.

Table 3.7: Representative smoothed deployed LoS two-ray parameters by height bin.

h_{tx} (m)	$\tilde{\rho}$	$\tilde{\phi}$ (rad)	$\tilde{b}$ (dB)
10	0.5605	0.000	-0.037
20	0.5242	0.000	-0.042
30	0.4950	0.000	-0.042
50	0.4632	0.000	-0.047
80	0.4958	0.000	-0.013
100	0.4926	0.000	-0.011
140	0.5308	0.000	-0.028
460	0.8966	-0.030	+1.432

The dominant deployed pattern is $\tilde{\rho} \approx 0.46$ to 0.56 with $\tilde{\phi} \approx 0$ for heights 10 to 100 m, meaning a moderate-amplitude ground-reflection with a near-zero fitted extra phase offset. Biases are near-zero (< 0.05 dB) for these bins. At very high-altitudes, the smoothed deployed $\tilde{\rho}$ and bias increase because training data is sparser; those bins are stabilized by Gaussian smoothing before deployment.




Figure 3.4: Effective ground-reflection amplitude $\tilde{\rho}$ as a function of UAV height after post fit smoothing across 5 m height bins. Values are stable in the range 0.46 to 0.56 for 10 to 200 m, then rise at the sparsely sampled very-high-altitude bins used by the calibration.

FSPL vs. coherent two-ray component ($h_{tx} = 60$ m, $f = 7.125$ GHz)




Figure 3.5: Schematic comparison of FSPL and the calibrated coherent two-ray component for a 60 m UAV at 7.125 GHz ($h_{tx} = 1.5$ m). The plotted red curve uses the code-loaded 60 m calibration values before adding the radial residual profile, so it shows the constructive/destructive interference term rather than the full final LoS prior.

3.4.3 NLoS branch

The NLoS path-loss branch used in the final model is not a pure literature formula. It is a hybrid engineering model built in five stages.

Stage 1: raw NLoS propagation prior

Two coarse physical estimates are combined.

COST-231 term. The COST-231 / Hata extension for urban macro cells gives the empirical large-scale path-loss expression used below [11]. It is included because it is one of the standard references for urban NLoS macrocell loss, even though the CKM scenario is outside its original frequency and base station height range. None of the COST-231 coefficients is therefore treated as a physically transferred absolute predictor at 7.125 GHz; the equation is only a log distance and height-sensitive basis that the train-only NLoS calibration can bend:

$$\begin{aligned} \text{PL}_{\text{COST231}}(x) = & 46.3 + 33.9 \log_{10}(f_{\text{MHz}}) - 13.82 \log_{10}(h_{\text{tx}}) - a(h_{\text{rx}}) \\ & + (44.9 - 6.55 \log_{10}(h_{\text{tx}})) \log_{10}(d_{\text{km}}) + 3 \end{aligned} \quad (3.24)$$

where $d_{\text{km}} = \max(d_{2D}(x)/1000, 0.001)$ and the implemented mobile antenna correction is the COST-231 small/medium city correction:

$$a(h_{\text{rx}}) = (1.1 \log_{10}(f_{\text{MHz}}) - 0.7) h_{\text{rx}} - (1.56 \log_{10}(f_{\text{MHz}}) - 0.8) \quad (3.25)$$

Equation (3.24) should therefore be read as a calibrated *feature basis*, not as a trusted absolute predictor or as a claim that COST-231 itself transfers to this 7.125 GHz UAV geometry. COST-231 contributes the expected log distance growth, transmitter height dependence and receiver height correction of an urban empirical model. The subsequent train-only calibration stage is what adapts that shape to the ray tracing setting used here.

Interpretation of Eq. (3.24).

- $46.3 + 33.9 \log_{10}(f_{\text{MHz}})$: frequency dependent offset. The coefficient 33.9 (compare to 20 for FSPL) reflects stronger frequency dependence in urban scattering.
- $-13.82 \log_{10}(h_{\text{tx}})$: as the base station / UAV gets higher, path loss decreases (better clearance above rooftops). The -13.82 coefficient is a fitted constant from the Hata empirical dataset.
- $-a(h_{\text{rx}})$: accounts for the receiver height using the implemented small/medium city Hata correction. For typical pedestrian heights ($h_{\text{rx}} \approx 1.5$ m), this is a small offset relative to the distance and height terms.
- $(44.9 - 6.55 \log_{10}(h_{\text{tx}})) \log_{10}(d_{\text{km}})$: distance dependent term. The effective path loss exponent is $(44.9 - 6.55 \log_{10}(h_{\text{tx}}))/10$, which decreases as the UAV climbs (effectively reducing NLoS scattering losses at longer distances from higher altitudes).
- $+3$: the COST-231 metropolitan correction constant (dB), corresponding to the dense-urban correction in the original model [11]. In this thesis it is retained only as an implementation basis constant, together with the small/medium city receiver height correction in Eq. (3.25). This deliberate mixture is not a canonical COST-231 deployment recipe; it is one empirical axis later re-fitted by the regime wise calibration.

Elevation-aware UAV envelope. The A2G NLoS envelope is used as a conservative elevation dependent shape term for urban air to ground path loss. UAV channel models motivate elevation angle as an important modelling axis because obstruction and shadowing statistics vary with elevation [14, 12]. The particular constants below, however, are not used as a transferred monotonic law. With the positive exponential term, the raw envelope can be larger near-overhead; it is therefore only one conservative basis branch that is later combined with COST-231 through a maximum and corrected by the train-only regime wise NLoS calibration.

$$\text{PL}_{\text{A2G,NLoS}}(x) = \lambda_0(h_{\text{tx}}, f) + \beta_{\text{NLoS}} + A_{\text{NLoS}} \exp\left(-\frac{90 - \theta^\circ(x)}{\tau_{\text{NLoS}}}\right) \quad (3.26)$$

where:

$$\lambda_0(h_{\text{tx}}, f) = 20 \log_{10} \left(\frac{4\pi h_{\text{tx}} f}{c} \right) \quad (3.27)$$

The fixed A2G-envelope constants used by the implementation are: $\beta_{\text{NLoS}} = -16.16$ dB, $A_{\text{NLoS}} = 12.0436$, $\tau_{\text{NLoS}} = 7.52$ deg. These A2G constants are implementation constants in the Try 78/80 prior code rather than entries in the calibration JSON files; the JSON artifacts store the fitted LoS two-ray curves and the regime wise NLoS calibration coefficients.

Interpretation of Eq. (3.26).

- λ_0 : the free space reference level used inside the raw A2G NLoS envelope before the empirical NLoS offset and elevation angle modulation are applied [14, 10]. The CKM dataset uses $f = 7.125$ GHz (FR3 6G frequency band); for $h_{\text{tx}} = 50$ m at that frequency, $\lambda_0 \approx 83.5$ dB.
- $\beta_{\text{NLoS}} = -16.16$ dB: a fixed baseline offset relative to λ_0 in the raw A2G envelope.
- $A_{\text{NLoS}} \exp(-\dots)$: an exponential modulation with $(90 - \theta^\circ)$. When the elevation angle θ° is high (UAV nearly overhead, $\theta^\circ \rightarrow 90^\circ$), the exponent vanishes and the term goes to A_{NLoS} . When θ° is low (grazing incidence), the exponent is large and the term approaches zero. With the current positive $A_{\text{NLoS}} = 12.0436$, this raw envelope therefore adds a larger A2G offset at high elevation. The later max(COST231, A2G) switch and the regime wise NLoS calibration decide whether that branch is active in the final raw prior; it should be read as a fitted conservative envelope, not as a standalone monotonic law saying all overhead NLoS links have lower loss.
- $\tau_{\text{NLoS}} = 7.52^\circ$: the characteristic angular scale. A low value means the transition from grazing to overhead happens sharply within a small elevation range.

The raw NLoS prior is then the conservative (higher) of the two estimates:

$$\text{PL}_{\text{NLoS,raw}}(x) = \max(\text{PL}_{\text{COST231}}(x), \text{PL}_{\text{A2G,NLoS}}(x)) \quad (3.28)$$

The maximum in Eq. (3.28) is deliberately pessimistic. If the empirical COST-231 curve predicts stronger attenuation than the A2G envelope, the model keeps the urban-macro value; if the elevation aware envelope predicts stronger attenuation, the model keeps the A2G value. This prevents the raw prior from becoming unrealistically optimistic in NLoS streets where neither source model is fully valid on its own.

Stage 2: raw hybrid prior assembly

The NLoS construction above defines only the obstructed branch. The full raw formula map used by the NLoS calibration stage is assembled only after both legacy branch estimates are available: the auxiliary LoS blend from Eq. (3.32), and the conservative NLoS branch from Eq. (3.28). This auxiliary LoS blend is not the final enhanced two-ray LoS prior; it is retained only because $\text{PL}^{(0)}$ is one input feature of the calibrated NLoS branch.

The legacy LoS path term used by this auxiliary branch is:

$$\text{PL}_{\text{LoS,path}}(x) = \begin{cases} \text{FSPL}(x), & d_{\text{LoS}}(x) \leq d_c, \\ 40 \log_{10}(d_{\text{LoS}}(x)) - 20 \log_{10}(h_{\text{tx}}) - 20 \log_{10}(h_{\text{rx}}), & d_{\text{LoS}}(x) > d_c, \end{cases} \quad (3.29)$$

with

$$d_c = \max \left\{ \frac{4\pi h_{\text{tx}} h_{\text{rx}}}{\lambda}, 1 \text{ m} \right\}, \quad \lambda = \frac{c}{f}. \quad (3.30)$$

Equation (3.29) switches from FSPL to the classical asymptotic two-ray d^4 path loss law after the crossover distance d_c . It is an older helper used only to construct the raw $\text{PL}^{(0)}$ feature for the NLoS calibration, not the final enhanced coherent two-ray LoS prior of Eq. (3.13).

The legacy LoS blend is then defined as:

$$\text{PL}_{\text{A2G,LoS}}(x) = \lambda_0(h_{\text{tx}}, f) - 20 \log_{10}(\max\{\sin \theta(x), 10^{-4}\}) \quad (3.31)$$

This auxiliary LoS envelope uses the same λ_0 term as Eq. (3.27); the sine clamp matches the implementation and avoids a singularity at grazing elevation. It is a legacy helper for $\text{PL}^{(0)}$, not the final LoS prior.

$$\text{PL}_{\text{LoS,blend}}(x) = 0.7 \text{PL}_{\text{LoS,path}}(x) + 0.3 \min(\text{PL}_{\text{LoS,path}}(x), \text{PL}_{\text{A2G,LoS}}(x)) \quad (3.32)$$

The 0.7/0.3 blend in Eq. (3.32) gives 70% weight to the legacy LoS path model and 30% weight to the lower-loss of the legacy path model and the A2G LoS envelope. Because lower dB means a stronger received signal, the $\min(\cdot)$ operation should not be read as a conservative guard against low path loss peaks; it is simply the older low-loss auxiliary blend inherited by the raw $\text{PL}^{(0)}$ feature. Here $\text{PL}_{\text{LoS,path}}(x)$ denotes the *legacy* LoS path term defined in Eq. (3.29). It is not the calibrated enhanced two-ray map of Eq. (3.13).

The full raw prior combines the LoS and NLoS estimates using the binary LoS mask $m_{\text{LoS}}(x) \in \{0, 1\}$:

$$\text{PL}^{(0)}(x) = m_{\text{LoS}}(x) \text{PL}_{\text{LoS,blend}}(x) + (1 - m_{\text{LoS}}(x)) \text{PL}_{\text{NLoS,raw}}(x) \quad (3.33)$$

Since $m_{\text{LoS}}(x)$ is binary, Eq. (3.33) is a hard switch: it selects the auxiliary LoS expression on LoS pixels and the NLoS expression on NLoS pixels. The auxiliary blend in Eq. (3.32) was used in older variants to keep the LoS formula below an A2G envelope. In the final frozen path, LoS output pixels come directly from the separate enhanced two-ray prior PL_{LoS} of Eq. (3.13); the raw hybrid map $\text{PL}^{(0)}$ in Eq. (3.33) is still computed only as an input feature for the calibrated NLoS path-loss branch. The LoS/NLoS mask used here is the explicit CKM mask provided with the dataset and treated as an available side channel. A deployment version would need to obtain an equivalent visibility mask from deterministic geometry or another inexpensive pre-processing step; replacing the supplied mask is not evaluated in this thesis.

Stage 3: map-derived morphology features

The raw hybrid prior $PL^{(0)}(x)$ defined in Eq. (3.33) is the scalar formula input to the calibrated NLoS path-loss branch. It combines the legacy FSPL/two-ray LoS path, the A2G envelopes and COST-231, but it still knows nothing about the detailed local morphology around each receiver. The feature vector $\mathbf{x}_{78}(x)$ defined below is *not* a competing path-loss model and *not* an alternative two-ray formulation. It is the per-pixel input row of a calibrated linear regressor whose role is to take the raw formula prediction and absorb its residual using context that the formula cannot see (urban morphology, NLoS visibility, shadowing and elevation geometry). Concretely, $PL^{(0)}(x)$ is fed in as the dominant feature, plus a quadratic of itself so that the linear fit can bend the prior, and every other entry adds a quantity that is invisible to the raw formula.

From $PL^{(0)}(x)$ and the map topology, a 15-element feature vector is computed:

$$\mathbf{x}_{78}(x) = \begin{bmatrix} PL^{(0)}(x)^2 \\ PL^{(0)}(x) \\ \log(1 + d_{2D}(x)) \\ \rho_{15}(x) \\ \rho_{41}(x) \\ h_{15}(x) \\ h_{41}(x) \\ \rho_{41}(x) \log(1 + d_{2D}(x)) \\ n_{15}(x) \\ n_{41}(x) \\ n_{41}(x) \log(1 + d_{2D}(x)) \\ \sigma_{\text{shadow}}(x) \\ \theta_{\text{norm}}(x) \\ n_{41}(x) \theta_{\text{norm}}(x) \\ 1 \end{bmatrix} \quad (3.34)$$

The lowercase bold notation is intentional: $\mathbf{x}_{78}(x)$ is a per-pixel feature vector. When vectors from many pixels in one regime are stacked for calibration, they form the uppercase design matrix $\mathbf{X}_r$. Table 3.8 gives the role of each entry before the detailed definitions below.

Term	Mathematical role	Physical interpretation
$PL^{(0)}(x)^2$	Quadratic basis term.	Lets the linear calibration bend the raw path loss curve when the raw model is systematically curved.
$PL^{(0)}(x)$	Linear raw prior term.	Keeps the calibrated model anchored to the physical prior instead of relearning path loss from scratch.
$\log(1 + d_{2D}(x))$	Smooth distance basis.	Mirrors log distance propagation laws used in classical models [9, 11].
$\rho_{15}(x), \rho_{41}(x)$	Local building-density averages.	Count how much nearby urban structure can block, reflect or diffract the signal.
$h_{15}(x), h_{41}(x)$	Density weighted local height averages.	Distinguish low dense blocks from taller blocks with stronger obstruction.
$\rho_{41}(x) \log(1 + d_{2D}(x))$	+ Density-distance interaction.	Allows distance slope to increase in dense areas.
$n_{15}(x), n_{41}(x)$	Local NLoS support fractions.	Measure how much nearby ground is geometrically shielded from the UAV.
$n_{41}(x) \log(1 + d_{2D}(x))$	+ NLoS-distance interaction.	Allows blocked streets to accumulate loss differently with range.
$\sigma_{\text{shadow}}(x)$	Elevation-dependent shadowing proxy.	Encodes the stronger low-elevation variability reported in A2G channel models [14, 12].
$\theta_{\text{norm}}(x)$	Normalized elevation angle.	Separates grazing links from near-overhead links.
$n_{41}(x)\theta_{\text{norm}}(x)$	NLoS-elevation interaction.	Allows the NLoS correction to depend on whether shielding occurs at low or high elevation.
1	Bias term.	Lets each regime learn a constant offset after all physical terms are included.

Table 3.8: Term-by-term inventory of the NLoS calibration vector.

Each term is now described in detail.

Base masks. The ground-mask is:

$$g(x) = \begin{cases} 1, & \text{if topology}(x) = 0 \\ 0, & \text{otherwise} \end{cases} \quad (3.35)$$

The building mask is $b(x) = 1 - g(x)$. The ground-NLoS support mask is:

$$n_{\text{raw}}(x) = \mathbf{1}[m_{\text{LoS}}(x) \leq 0.5] \cdot g(x) \quad (3.36)$$

This selects ground pixels that are predominantly in NLoS, i.e., pixels that are at street level and are shielded from the UAV.

Quadratic raw prior feature. The pair $(\text{PL}^{(0)}(x)^2, \text{PL}^{(0)}(x))$ allows the linear model to fit a *quadratic* relationship in the raw prior. Including both P^2 and P as features lets the calibration curve be: $\hat{P} = \beta_1 P^2 + \beta_2 P + \dots$, which can correct for systematic curvature in the raw prior. This pair is intentionally redundant: the squared feature is completely determined by the raw prior feature. It is kept because the calibrated fit improves when the linear model is allowed to represent a curved mapping from raw prior to measured path loss.

Distance term. $\log(1 + d_{2D}(x))$: the logarithm of one plus the horizontal distance, providing a distance dependent offset. The +1 prevents a singularity at $d_{2D} = 0$. The logarithm is consistent with the log distance structure of empirical path loss laws: a multiplicative increase in distance produces an additive increase in predicted dB loss [9, 11].

Per-pixel box-mean operator. All multiscale features use the same 2D box-mean operator BoxMean_k . In the prior vector this operator is not a standalone feature; it is the common construction used to compute the neighborhood terms ρ_{15}, ρ_{41} from the building mask, h_{15}, h_{41} from density-weighted building-height, and n_{15}, n_{41} from the ground-NLoS support mask. For a generic input map $f : \mathbb{Z}^2 \rightarrow \mathbb{R}$ and a window half-width $p = \lfloor k/2 \rfloor$, the value at pixel $x = (i, j)$ is:

$$\text{BoxMean}_k(f)(i, j) = \frac{1}{k^2} \sum_{i'=i-p}^{i+p} \sum_{j'=j-p}^{j+p} f_{\text{padded}}(i', j') \quad (3.37)$$

where f_{padded} is f extended outside the map boundary by *reflect* (mirror) padding: $f_{\text{padded}}(-r, j) = f_{\text{padded}}(r, j)$ and similarly at all four edges. This ensures border pixels receive a full k^2 -pixel average, with the map reflected back inward rather than zero-padded.

Equation (3.37) is evaluated efficiently using a *2D summed-area table* (integral image). Let $I(i, j) = \sum_{i' \leq i, j' \leq j} f_{\text{padded}}(i', j')$ be the prefix-sum array. Then:

$$\begin{aligned} S_k(i, j) &= \sum_{i'=i-p}^{i+p} \sum_{j'=j-p}^{j+p} f_{\text{padded}}(i', j') \\ &= I(i+p, j+p) - I(i-p-1, j+p) \\ &\quad - I(i+p, j-p-1) + I(i-p-1, j-p-1) \end{aligned} \quad (3.38)$$

This reduces each per-pixel evaluation to four table lookups, making the full-map computation $O(HW)$ regardless of k . In the implementation, $k \in \{15, 41\}$ pixels, corresponding to ~ 15 m and ~ 41 m neighborhoods (1 m/pixel resolution).

Multiscale building-density.

$$\rho_k(x) = \text{BoxMean}_k(b(\cdot))(x) = \frac{1}{k^2} \sum_{x' \in W_k(x)} b(x') \quad (3.39)$$

where $W_k(x)$ is the $k \times k$ window centered at x . Since $b(x') \in \{0, 1\}$, $\rho_k(x)$ is the *fraction of pixels that are buildings* in the window. It is zero at a ground pixel surrounded only by other ground pixels, and one at a pixel deep inside a solid building block. ρ_{15} captures fine-scale building-density within a ~ 15 m neighborhood; ρ_{41} captures a coarser ~ 41 m neighborhood.

Multiscale building-height.

$$h_k(x) = \frac{\text{BoxMean}_k(\text{topology}(\cdot) b(\cdot))(x)}{H_{\text{norm}}}, \quad H_{\text{norm}} = 90 \text{ m} \quad (3.40)$$

This is the box mean of $\text{topology}(x') \cdot b(x')$: building pixels contribute their height in metres; ground pixels contribute 0. Dividing by k^2 (inside BoxMean_k) gives:

$$h_k(x) = \frac{1}{H_{\text{norm}} k^2} \sum_{x' \in W_k(x)} \text{topology}(x') \cdot b(x') = \frac{\rho_k(x)}{H_{\text{norm}}} \bar{h}_{\text{bldg},k}(x) \quad (3.41)$$

where $\bar{h}_{\text{bldg},k}(x)$ is the *mean height of building pixels only* within the window:

$$\bar{h}_{\text{bldg},k}(x) = \frac{\sum_{x' \in W_k(x)} \text{topology}(x') b(x')}{\sum_{x' \in W_k(x)} b(x')} \quad (\text{undefined if } \rho_k(x) = 0) \quad (3.42)$$

Critically, $h_k(x) = \rho_k(x) \cdot \bar{h}_{\text{bldg},k}(x) / H_{\text{norm}}$, so h_k combines *both* how many buildings are present and how tall they are. A window with 50% building coverage and mean height 20 m gives the same h_k as a window with 100% coverage at mean height 10 m. The $H_{\text{norm}} = 90$ m normalizer maps the joint range to approximately $[0, 1]$.

Multiscale NLoS support.

$$n_k(x) = \text{BoxMean}_k(n_{\text{raw}}(\cdot))(x) \quad (3.43)$$

Since $n_{\text{raw}}(x') \in \{0, 1\}$ (Eq. 3.36), $n_k(x)$ is the *absolute fraction of the window-area* that consists of ground pixels in NLoS. Just as $h_k(x)$ conflates building coverage and building height, $n_k(x)$ conflates ground coverage and NLoS probability. We can write:

$$n_k(x) = (1 - \rho_k(x)) \bar{n}_{\text{ground},k}(x) \quad (3.44)$$

where $\bar{n}_{\text{ground},k}(x)$ is the relative fraction of *available ground pixels* that are shielded:

$$\bar{n}_{\text{ground},k}(x) = \frac{\sum_{x' \in W_k(x)} n_{\text{raw}}(x')}{\sum_{x' \in W_k(x)} g(x')} \quad (\text{undefined if } \rho_k(x) = 1) \quad (3.45)$$

Because building pixels never contribute to n_k , this feature is zero in completely open LoS areas (where $\bar{n}_{\text{ground},k} = 0$) and deep inside solid building blocks (where $1 - \rho_k = 0$). It reaches its maximum in narrow urban street canyons, where ground coverage is still significant but entirely shielded from the UAV.

Elevation angle.

$$\theta(x) = \arctan\left(\frac{h_{\text{tx}} - h_{\text{rx}}}{\max(d_{2D}(x), 1)}\right) \quad (3.46)$$

The $\max(\cdot, 1)$ prevents division by zero directly below the UAV. Then:

$$\theta_{\text{norm}}(x) = \text{clip}\left(\frac{\theta^\circ(x)}{90}, 0, 1\right) \quad (3.47)$$

where $\theta^\circ = \theta \cdot 180/\pi$ is in degrees. $\theta_{\text{norm}} = 0$ at grazing incidence, $= 1$ directly overhead.

Shadowing proxy.

$$\sigma_{\text{LoS}}(x) = \alpha_{\text{LoS}} (90 - \theta^\circ(x))^{\mu_{\text{LoS}}} \quad (3.48)$$

$$\sigma_{\text{NLoS}}(x) = \alpha_{\text{NLoS}} (90 - \theta^\circ(x))^{\mu_{\text{NLoS}}} \quad (3.49)$$

with fitted values $\alpha_{\text{LoS}} = 0.0272$, $\mu_{\text{LoS}} = 0.7475$, $\alpha_{\text{NLoS}} = 2.3197$, and $\mu_{\text{NLoS}} = 0.2361$. The blended proxy is:

$$\sigma_{\text{shadow}}(x) = m_{\text{LoS}}(x) \sigma_{\text{LoS}}(x) + (1 - m_{\text{LoS}}(x)) \sigma_{\text{NLoS}}(x) \quad (3.50)$$

This acts as a spatially varying uncertainty proxy rather than a measured shadow-fading standard deviation. Higher values occur at low-elevation angles, where propagation conditions are more difficult; the LoS/NLoS blend lets the same feature have a different scale in the two visibility regimes.

Interaction terms. The three pairwise products $\rho_{41}(x) \log(1 + d_{2D}(x))$, $n_{41}(x) \log(1 + d_{2D}(x))$, and $n_{41}(x) \theta_{\text{norm}}(x)$ allow the linear model to capture effects where building-density or NLoS support *modifies* the slope of path loss with distance or elevation angle, rather than just shifting it. These terms are also deliberately correlated with their factors. For example, $n_{41}(x) \log(1 + d_{2D}(x))$ cannot vary independently of $n_{41}(x)$ or $\log(1 + d_{2D}(x))$. This cross-correlation makes coefficient interpretation less clean, but removing such terms can reduce performance because the fit then loses a simple way to express “distance matters differently in blocked streets”. Determining exactly which terms are dispensable would require a large ablation sweep over pairs or subsets of features. Even a pairwise screening would scale as $O(n^2)$ runs for a vector with n terms, and each run would have to regenerate and inspect large calibration artifacts. The LoS two-ray calibration JSON used by the final model, for example, is about 200 000 lines long, mostly because it stores height-dependent radial residual profiles. Taken together, the stored calibration is height-dependent, building-topology-dependent, and density-dependent: the LoS JSON contributes height-dependent radial residuals, while the NLoS side adds morphology calibration through building-density, building-height, NLoS support, topology regimes, and antenna height features. For this reason the prior should be read as a calibrated

baseline with substantial stored structure, not as a single hand written formula, and the final vector keeps some correlated terms when they improved validation behavior.

Stage 4: regime wise linear calibration

The regime r is identified by three keys:

- *Coarse topology*: one of open_lowrise, mixed_midrise, or dense_highrise,
- *Propagation state*: LoS or NLoS,
- *Antenna height bin*: low_ant, mid_ant, or high_ant.

This gives up to $3 \times 2 \times 3 = 18$ regimes.

The stored calibration for this legacy mixed-prior map still contains both LoS and NLoS regime coefficients. However, in the final frozen path loss assembly only the NLoS pixels from this calibrated map are used, because LoS pixels are replaced by the separate enhanced two-ray prior.

How topology class and antenna bin are computed. These are *per-map* scalar labels, not per-pixel quantities. They are derived once from the full topology tensor and UAV height.

The *map-level* building-density and mean building-height are:

$$\bar{\rho}_{\text{map}} = \frac{1}{HW} \sum_x b(x) \quad (3.51)$$

$$\bar{h}_{\text{map}} = \frac{\sum_x \text{topology}(x) b(x)}{\sum_x b(x)} \quad (3.52)$$

$\bar{\rho}_{\text{map}}$ is simply the fraction of map pixels that are buildings. $\bar{h}_{\text{map}}$ is the mean building-height *over building pixels only* (the denominator excludes ground pixels). This is different from the per-pixel feature $h_k(x)$ in Eq. (3.41), which divides the same numerator by the full window-area k^2 instead of by the number of building pixels.

For the calibrated path-loss prior, these two quantities are converted into three coarse city types. This is a separate regime label from the later six-class topology taxonomy used by the spread priors and by final reporting:

$$\text{class}_{78} = \begin{cases} \text{dense_highrise} & \bar{\rho} \geq d_2 \text{ or } \bar{h} \geq h_2 \\ \text{open_lowrise} & \bar{\rho} \leq d_1 \text{ and } \bar{h} \leq h_1 \\ \text{mixed_midrise} & \text{otherwise} \end{cases} \quad (3.53)$$

with thresholds $d_1 = 0.1957$, $d_2 = 0.2549$ (density fractions), $h_1 = 10.91$ m, and $h_2 = 15.95$ m. Together with the binary LoS/NLoS label and the three UAV-height bins below, this yields the 18 stored path loss calibration keys.

The antenna height bin is assigned by UAV height tertiles:

$$\text{ant_bin} = \begin{cases} \text{low_ant} & h_{\text{tx}} \leq 58.12 \text{ m} \\ \text{mid_ant} & 58.12 < h_{\text{tx}} \leq 103.85 \text{ m} \\ \text{high_ant} & h_{\text{tx}} > 103.85 \text{ m} \end{cases} \quad (3.54)$$

The calibrated prediction for regime r is:

$$\widehat{\text{PL}}_r(x) = \mathbf{x}_{78}(x)^\top \mathbf{w}_r^{78} = \sum_{j=1}^{15} x_j(x) w_{r,j}^{78} \quad (3.55)$$

How $\mathbf{w}_r^{78}$ is fitted. Across all training samples whose regime label equals r , a linear regression is run:

$$\hat{\mathbf{w}}_r^{78} = \arg \min_{\mathbf{w} \in \mathbb{R}^{15}} \sum_{x \in \mathcal{S}_r} (y(x) - \mathbf{x}_{78}(x)^\top \mathbf{w})^2 \quad (3.56)$$

where $\mathcal{S}_r$ is the set of all valid pixels in training samples belonging to regime r . This ordinary least-squares problem has the closed form solution, when $\mathbf{X}_r^\top \mathbf{X}_r$ is nonsingular:

$$\hat{\mathbf{w}}_r^{78} = (\mathbf{X}_r^\top \mathbf{X}_r)^{-1} \mathbf{X}_r^\top \mathbf{y}_r \quad (3.57)$$

where $\mathbf{X}_r$ stacks the feature vectors and $\mathbf{y}_r$ stacks the measured path-loss values. The runtime code does not refit this model; it loads the stored coefficient vectors from `try78_nlos_regime_calibration.json`. At inference, if the exact antenna height key is absent, the implementation only tries alternative antenna bins in the order current, mid, low, high while keeping the same path loss-prior city type and LoS/NLoS label. It does not use the broader cross-topology fallback hierarchy of the spread-prior calibration.

The terms in Eq. (3.57) are standard linear least-squares quantities. The matrix $\mathbf{X}_r^\top \mathbf{X}_r$ is the *Gram matrix*; its diagonal entries measure the energy of each feature inside regime r , and its off-diagonal entries measure correlations between features. The vector $\mathbf{X}_r^\top \mathbf{y}_r$ measures how strongly each feature correlates with measured path loss. The inverse maps those feature-target correlations into the fitted coefficient vector $\hat{\mathbf{w}}_r^{78}$. This is the classical ordinary least-squares solution, but applied separately per propagation regime rather than globally. Because the feature vector intentionally contains correlated terms, the important object is the least-squares coefficient vector stored in `try78_nlos_regime_calibration.json`; if a normal-equation matrix is rank-deficient or ill-conditioned, Eq. (3.57) should be understood as the full-rank notation for the corresponding least-squares or pseudo-inverse solution, not as a requirement that every calibration block be perfectly invertible.

An example fitted coefficient vector for the mixed_midrise|NLoS|high_ant regime is:

$$\mathbf{w}_{\text{mixed|NLoS|high}}^{78} = \begin{bmatrix} -0.000338, & -0.0266, & 8.233, & -1.452, & -11.659 \\ & 3.508, & -5.611, & 1.904, & 0.926, & 12.084 \\ & -1.628, & -7.487, & -12.665, & -10.959, & 121.480 \end{bmatrix}^\top$$

The large bias weight and the sizeable morphology, elevation and raw prior weights should be interpreted together as one fitted calibration vector, not as independent physical

constants. Their magnitudes mainly reflect the scale and correlation of the 15 input features in this NLoS regime. This makes the path-loss prior an interpretable calibrated surrogate rather than a pure first-principles propagation formula: the two-ray and A2G terms set the shape, while regime wise fitting adapts the coefficients to the CKM simulator and to the dataset’s stored masks.

Stage 5: region replacement and clipping

$$\text{PL}_{\text{NLoS}}(x) = \text{clip}(\mathbf{x}_{78}(x)^\top \mathbf{w}_r^{78}, 20, 180) \quad (3.58)$$

Hard clipping to $[20, 180]$ dB prevents the linear model from producing physically unreasonable values. The final hybrid path-loss prior is:

$$\text{PL}_{\text{prior}}(x) = \begin{cases} \text{PL}_{\text{LoS}}(x), & m_{\text{LoS}}(x) > 0.5 \\ \text{PL}_{\text{NLoS}}(x), & \text{otherwise} \end{cases} \quad (3.59)$$

3.4.4 Evaluation metric: pixel-weighted RMSE

All numerical results are reported as *pixel-weighted* (pw) RMSE or MAE. Building pixels ($\text{topology}(x) > 0$) are always excluded; only valid ground pixels with a finite, physically plausible target value contribute. Let $\mathcal{V}_k$ be the set of all such valid (s, x) pairs across all evaluation samples for task k . Then:

$$\text{RMSE}_k^{\text{pw}} = \sqrt{\frac{\sum_{(s,x) \in \mathcal{V}_k} \left(\hat{y}_k^{(s)}(x) - y_k^{(s)}(x) \right)^2}{|\mathcal{V}_k|}} \quad (3.60)$$

This is a single square root over all pooled valid pixels, not an average of per-sample RMSEs. Units are dB (path loss), ns (delay spread), and degrees (angular spread), all in the native non-transformed domain.

3.4.5 Observed path loss-prior results

The headline path loss numbers for the calibrated prior are reported in the results chapter (LoS waterfall in Table 4.1): the enhanced LoS two-ray prior PL_{LoS} reaches 1.7516 dB pixel-weighted RMSE on LoS pixels, the regime wise OLS NLoS head reaches 3.3967 dB on NLoS pixels, and the combined prior reaches 1.9246 dB overall on the held-out validation+test split.




Figure 3.6: Pipeline for the hybrid path-loss prior. The LoS branch relies on an explicit two-ray interference model, while the NLoS branch calibrates a raw prior against map-derived morphology features using regime wise OLS. The two branches are hard-switched based on the LoS mask to produce the final frozen prior.

3.5 Angular Spread Prior and Residual Prediction

Angular spread uses the shared spread prior described below. Its final prediction is the frozen angular-spread prior plus the bounded residual correction predicted by HARP-NET CKM. The target is reported in degrees and is evaluated only on valid ground receiver pixels.

3.6 Delay Spread Prior and Residual Prediction

Delay spread uses the same shared spread prior machinery, but it is evaluated in nanoseconds and is more sensitive to rare high delay pixels. The model therefore uses the prior as a stable baseline and learns only a bounded correction.

Shared spread prior details

3.6.1 Physical motivation for log-domain modeling

Delay spread (DS) and angular spread (AS) are nonnegative, often zero-inflated, and heavy-tailed. Their empirical distributions are commonly modeled in the logarithmic domain in standard channel models [12, 17]. For a spread variable y , define:

$$z(x) = \log(1 + y(x)), \quad \hat{y}(x) = \exp(\hat{z}(x)) - 1 \quad (3.61)$$

The constant 1 keeps exact zeros finite, avoids $\log(0)$, and makes the same transform usable for both delay spread in nanoseconds and angular spread in degrees. If two pixels have native spread values y_a and y_b , their difference in transformed space is

$$z_a - z_b = \log\left(\frac{1 + y_a}{1 + y_b}\right). \quad (3.62)$$

Thus, a fixed error in z corresponds approximately to a fixed multiplicative error in $1 + y$. This is more appropriate than native-space least-squares for spreads because large NLoS tails would otherwise dominate the fit and the near-zero LoS mass would be ignored. The calibrated spread prior is still a deterministic log-domain ridge estimator; it is not claiming to estimate the full spike-plus-tail probability distribution.

3.6.2 Pixel-wise notation for the spread prior fit

Let n index a valid ground pixel from a training map. For a selected metric $m \in \{\text{DS}, \text{AS}\}$, define the native target $y_{m,n}$, the transformed target $z_{m,n} = \log(1 + y_{m,n})$, and the raw log prior value $z_{m,n}^{(0)}$. The calibrated prior is linear in a 23-dimensional feature vector:

$$\hat{z}_{m,n} = \mathbf{x}_{m,n}^\top \mathbf{w}_{m,r(n)}, \quad \mathbf{x}_{m,n} \in \mathbb{R}^{23}, \quad \mathbf{w}_{m,r} \in \mathbb{R}^{23}. \quad (3.63)$$

The regime $r(n)$ is a discrete key containing the metric, topology class, LoS/NLoS state and UAV-height bin. In other words, the spread prior does not fit one global line through all pixels. It fits a family of small linear models, each one specialized to a propagation regime.

3.6.3 Raw spread prior in log-domain

$$\begin{aligned} z^{(0)}(x) = & \log\left(1 + b_{\text{LoS/NLoS}}^{\text{native}}(x)\right) + b_{\text{topology}} + a_1 \log(1 + d_{2D}) + a_2 \theta_{\text{inv}} \\ & + a_3 \rho_{41} + a_4 h_{41} + a_5 n_{41} + a_6 n_{41} \theta_{\text{inv}}. \end{aligned} \quad (3.64)$$

where $\theta_{\text{inv}}(x) = 1 - \theta_{\text{norm}}(x)$ (high value at low-elevation). Example coefficients for delay spread: the native anchors are $b_{\text{LoS}}^{\text{native}} = 4.0$ ns and $b_{\text{NLoS}}^{\text{native}} = 11.0$ ns, which the implementation converts to $\log(1 + 4.0)$ and $\log(1 + 11.0)$ before adding the remaining log-domain terms. The delay coefficients are $a_1 = 0.11$, $a_2 = 0.62$, $a_3 = 0.48$, $a_4 = 0.32$, $a_5 = 0.86$, $a_6 = 0.60$. After Eq. (3.64), the raw log prior is clipped to $[0, 8]$, inverse transformed with $\exp(\cdot) - 1$, and clipped to the metric upper bound. For reproducibility, the $z^{(0)}(x)$ entries used later in $\mathbf{x}_{79}$ denote the log value recovered from this clipped native raw prior, i.e., $\log(1 + \text{raw_prior}_{\text{clipped}}(x))$, not an unconstrained copy of Eq. (3.64).

The raw prior is not a 3GPP formula copied directly into CKM. Instead, it takes the 3GPP/WINNER lesson that spread statistics should be separated by propagation state and represented in log space [12, 17], then adds CKM specific map morphology terms. UAV measurement studies also motivate the elevation angle dependence: low-elevation links usually interact with more rooftops and street canyons than near-overhead links [47, 48]. For reproducibility, Table 3.9 lists the fixed raw prior constants used before the regime wise ridge calibration.

Table 3.9: Fixed constants used in the raw spread prior of Eq. (3.64).

Quantity	Delay spread	Angular spread	Meaning
$b_{\text{LoS}}^{\text{native}}$	4.0 ns	3.0°	Native LoS anchor converted to $\log(1 + b)$.
$b_{\text{NLoS}}^{\text{native}}$	11.0 ns	7.0°	Native NLoS anchor converted to $\log(1 + b)$.
$a_1, \dots, a_6$	0.11, 0.62, 0.48, 0.32, 0.86, 0.60	0.05, 0.52, 0.60, 0.38, 0.58, 0.35	Coefficients of $\log(1 + d_{2D})$, θ_{inv} , ρ_{41} , h_{41} , n_{41} , and $n_{41}\theta_{\text{inv}}$, respectively.

The topology offset b_{topology} is also fixed before ridge calibration. For delay spread, the offsets for open_sparse_lowrise, open_sparse_vertical, mixed_compact_lowrise, mixed_compact_midrise, dense_block_midrise, and dense_block_highrise are 0.00, 0.18, 0.06, 0.22, 0.24, 0.34. For angular spread, the corresponding offsets are 0.00, 0.09, 0.05, 0.16, 0.18, 0.24. These are log-domain offsets, not native ns/degree offsets.

Every term in Eq. (3.64) has a physical interpretation:

- $b_{\text{LoS/NLoS}}$ is the base spread level for the propagation state. NLoS starts higher because blocked paths arrive after additional reflections and diffractions.
- b_{topology} is a morphology offset. Dense and high-rise topologies receive larger offsets because they contain more reflectors.
- $\log(1 + d_{2D})$ gives a smooth range dependence. Longer links have more opportunity to encounter delayed or angularly separated paths.
- θ_{inv} increases at low-elevation angle, where the UAV to ground ray intersects more urban structure before reaching the receiver.
- ρ_{41} , h_{41} , and n_{41} describe local building density, mean building-height, and NLoS support in a 41×41 window. These are direct map-derived proxies for multipath richness.
- $n_{41}\theta_{\text{inv}}$ is an interaction term. It activates when low-elevation geometry and local obstruction occur together.

3.6.4 Calibrated spread prior

Stage 2: 23-feature vector and design matrix

$$\mathbf{x}_{79}(x) = \begin{bmatrix} z^{(0)}(x)^2, z^{(0)}(x), \log(1 + d_{2D}), \theta_{\text{norm}}, \theta_{\text{inv}} \\ h_{\text{norm}}, h_{\text{norm}}^2, \rho_{15}, \rho_{41}, h_{15}, h_{41} \\ n_{15}, n_{41}, n_{41} \log(1 + d_{2D}), \rho_{41}\theta_{\text{inv}}, b_{41}, 1 \\ c_{41}, t_{41}, \theta_{\text{norm}}\rho_{41}, c_{41}\theta_{\text{inv}}, t_{41}\rho_{41}, t_{41}n_{41} \end{bmatrix}^T \quad (3.65)$$

The vector in Eq. (3.65) can be grouped into five blocks:

$$\mathbf{x}_{79}(x) = \left[\underbrace{\mathbf{x}_{\text{prior}}(x)}_{\text{raw prior}}, \underbrace{\mathbf{x}_{\text{geom}}(x)}_{\text{range/elevation/height}}, \underbrace{\mathbf{x}_{\text{morph}}(x)}_{\text{local buildings}}, \underbrace{1}_{\text{bias}}, \underbrace{\mathbf{x}_{\text{int}}(x)}_{\text{interactions}} \right]^{\top}. \quad (3.66)$$

The first block lets the model bend the raw prior quadratically. The geometry block handles range and height. The morphology block measures whether there are nearby scatterers. The interaction block allows the model to express statements such as “density matters more at low-elevation” or “density-weighted transmitter clearance matters differently at low-elevation”. As in the path-loss prior vector, some of these terms are not independent. For instance, $z^{(0)}(x)^2$ is determined by $z^{(0)}(x)$, and $t_{41}(x)n_{41}(x)$ is tied to both the taller-building and NLoS support features. This cross-correlation means individual fitted coefficients should not be interpreted as isolated physical constants. The terms were kept because the calibrated prior improved with them, and a rigorous feature-pruning study would require many repeated calibration runs. A simple pairwise screening would already scale as $O(n^2)$ for an n -term vector, before considering higher-order subsets.

Term	Mathematical role	Physical interpretation
$z^{(0)}(x)^2$	Quadratic raw prior basis.	Allows the calibration to correct curvature in the first spread estimate.
$z^{(0)}(x)$	Linear raw prior basis.	Keeps the calibrated model tied to the log-domain physical prior.
$\log(1 + d_{2D})$	Smooth distance basis.	Longer links can produce wider delay and angle supports because more scatterers can contribute.
θ_{norm}	Normalized elevation angle.	Separates overhead links from grazing links.
θ_{inv}	Inverse elevation angle.	Gives a large value to low-elevation links, where A2G models report stronger obstruction sensitivity [14, 12].
$h_{\text{norm}}, h_{\text{norm}}^2$	UAV height basis and quadratic height basis.	Allows the model to learn a nonlinear effect of transmitter height on spread.
ρ_{15}, ρ_{41}	Fine/coarse local building-density.	Measures how many nearby reflectors and blockers exist at two scales.
h_{15}, h_{41}	Fine/coarse density-weighted building-height.	Distinguishes low-rise clutter from high-rise clutter.
n_{15}, n_{41}	Fine/coarse local NLoS support.	Measures how much nearby ground is geometrically shielded.
b_{41}	Blocker depth term.	Measures how far nearby buildings protrude above receiver height.
1	Bias term.	Gives each regime its own additive offset.
c_{41}	Transmitter clearance term.	Density weighted window-area mean of UAV clearance above nearby building pixels.
t_{41}	Taller building area fraction.	Fraction of the full local window occupied by buildings whose height exceeds the UAV height.

Table 3.10: Term by term inventory of the direct terms in the spread calibration vector.

Term	Mathematical role	Physical interpretation
$n_{41} \log(1 + d_{2D})$	NLoS distance interaction.	Lets range matter differently in blocked streets than in open ground.
$\rho_{41} \theta_{\text{inv}}$	Density elevation interaction.	Activates when dense morphology and low-elevation occur together.
$\theta_{\text{norm}} \rho_{41}$	Overhead density interaction.	Allows dense areas to behave differently when viewed from high elevation.
$c_{41} \theta_{\text{inv}}$	Clearance elevation interaction.	Lets density-weighted rooftop clearance matter differently at grazing elevation.
$t_{41} \rho_{41}$	Tall building-density interaction.	Highlights dense high-rise blocks where a large window-area contains buildings above the UAV.
$t_{41} n_{41}$	Tall building NLoS interaction.	Activates in blocked regions where taller-building area is also present.

Table 3.11: Interaction terms in the spread calibration vector.




Figure 3.7: Feature grouping used by the calibrated spread prior. The calibrated model is linear in $\mathbf{x}_{79}$, but the features include nonlinear transforms and interactions selected from propagation intuition.

The new features beyond those in $\mathbf{x}_{78}$ are listed below. The shared morphology notation follows the path-loss prior construction in Sec. 3.4.3. In particular, $b(\cdot)$ is the building mask, ρ_k is the building-density box mean, h_k is the density-weighted building-height box mean, n_k is the ground NLoS support box mean, and θ_{norm} is the normalized elevation angle. The same reflected padding BoxMean $_k$ operator is used here, and $H_{\text{norm}} = 90$ m is the building-height normalizer. This keeps the spread prior vector self contained while preserving the common feature definitions used by the path-loss prior.

Normalised transmitter height.

$$h_{\text{norm}} = \min \left\{ 1, \frac{\log(1 + \max(h_{\text{tx}} - h_{\text{rx}}, 0))}{\log(401)} \right\} \quad (3.67)$$

This compresses UAV height on a log scale and saturates it in $[0, 1]$, using $\log(401)$ as the scale associated with a 400 m height difference. Samples above that relative height are assigned $h_{\text{norm}} = 1$, matching the implementation. This cap is intentional: the transmitter height distribution has only a small high-altitude tail, so after splitting samples into topology-specific experts there is not enough support above 400 m to estimate a separate high-altitude response reliably. The vector also includes h_{norm}^2 , which is an intentionally correlated basis expansion that lets the linear model learn a curved height effect.

Blocker depth.

$$b_{41}(x) = \frac{\text{BoxMean}_{41}(\max(\text{topology}(\cdot) - h_{\text{rx}}, 0) b(\cdot))(x)}{H_{\text{norm}}} \quad (3.68)$$

Measures how much buildings in the 41×41 window protrude above the receiver height. Only building pixels contribute because of the multiplier $b(\cdot)$. A tall building directly next to the receiver therefore increases b_{41} , while flat open ground leaves it at zero.

Transmitter clearance.

$$c_{41}(x) = \frac{\text{BoxMean}_{41}(\max(h_{\text{tx}} - \text{topology}(\cdot), 0) b(\cdot))(x)}{H_{\text{norm}}} \quad (3.69)$$

This is a density-weighted clearance feature: the clearance above each building pixel is averaged by BoxMean_{41} over the full 41×41 window and then normalized by H_{norm} . Ground pixels contribute zero through the multiplier $b(\cdot)$. A large c_{41} therefore means both that nearby buildings are present and that the transmitter is well above them. A small value can mean either little local building coverage, or buildings that reach close to or above the UAV height.

Taller building fraction.

$$t_{41}(x) = \text{BoxMean}_{41}(\mathbf{1}[\text{topology}(\cdot) > h_{\text{tx}}] \cdot b(\cdot))(x) \quad (3.70)$$

This is the fraction of the full 41×41 window-area occupied by building pixels taller than the UAV, not the fraction of buildings after conditioning on there being a building pixel. It is zero if the UAV is above all local buildings or if the local window contains no buildings, and grows when the link is likely to be embedded in a high-rise canyon. It is especially useful for distinguishing dense but low-rise zones from dense high-rise zones.

Six class topology definition. The spread prior uses the six topology labels produced by the final Try 80 prior code. For a map, let $\bar{\rho}_{\text{map}}$ be the fraction of non ground pixels

and let $\bar{h}_{\text{map}}$ be the mean height over non ground pixels only. The class assignment is:

$$\text{class}_{79} = \begin{cases} \text{open_sparse_lowrise}, & \bar{\rho}_{\text{map}} \leq 0.12, \bar{h}_{\text{map}} \leq 12 \text{ m}, \\ \text{open_sparse_vertical}, & \bar{\rho}_{\text{map}} \leq 0.12, \bar{h}_{\text{map}} > 12 \text{ m}, \\ \text{dense_block_midrise}, & \bar{\rho}_{\text{map}} \geq 0.28, \bar{h}_{\text{map}} \leq 28 \text{ m}, \\ \text{dense_block_highrise}, & \bar{\rho}_{\text{map}} \geq 0.28, \bar{h}_{\text{map}} > 28 \text{ m}, \\ \text{mixed_compact_lowrise}, & 0.12 < \bar{\rho}_{\text{map}} < 0.28, \bar{h}_{\text{map}} \leq 12 \text{ m}, \\ \text{mixed_compact_midrise}, & 0.12 < \bar{\rho}_{\text{map}} < 0.28, \bar{h}_{\text{map}} > 12 \text{ m}. \end{cases} \quad (3.71)$$

These labels are map-level routing keys for the spread prior ridge models; they are not additional per-pixel morphology features.

Stage 3: regime wise ridge regression

The regime r for the spread prior is keyed by: metric (DS or AS), topology class (6 classes), propagation state (LoS or NLoS), antenna height bin, giving up to $2 \times 6 \times 2 \times 3 = 72$ regimes. The antenna height bins use the same tertile style thresholds as Eq. (3.54): low_ant for $h_{\text{tx}} \leq 58.12 \text{ m}$, mid_ant for $58.12 \text{ m} < h_{\text{tx}} \leq 103.85 \text{ m}$, and high_ant above 103.85 m . These calibration bins are different from the broader height reporting bins used later in the results tables.

Valid pixels are subsampled at rate $p = 0.02$ (2%). The ridge regression:

$$\hat{\mathbf{w}}_r^{79} = (\mathbf{X}_r^\top \mathbf{X}_r + \lambda \mathbf{D}^\top \mathbf{D})^{-1} \mathbf{X}_r^\top \mathbf{z}_r, \quad \lambda = 10^{-3} \quad (3.72)$$

Here $\mathbf{X}_r \in \mathbb{R}^{N_r \times 23}$ stacks one row $\mathbf{x}_{79,n}^\top$ per-sampled pixel in regime r , and $\mathbf{z}_r \in \mathbb{R}^{N_r}$ stacks the corresponding transformed targets. The coefficient vector $\hat{\mathbf{w}}_r^{79}$ minimizes the regularized least-squares problem

$$\hat{\mathbf{w}}_r^{79} = \arg \min_{\mathbf{w} \in \mathbb{R}^{23}} \|\mathbf{X}_r \mathbf{w} - \mathbf{z}_r\|_2^2 + \lambda \|\mathbf{D} \mathbf{w}\|_2^2, \quad (3.73)$$

where $\mathbf{D} = \text{diag}(d_1, \dots, d_{23})$, with $d_{23} = 0$ and $d_i = 1$ for $i \neq 23$. Thus, the final interaction term $t_{41}n_{41}$, which is the last entry of Eq. (3.65), is left unpenalized by the ridge term in the frozen calibration used by the final model. This is an implementation detail of the stored artifact, not a separate physical assumption. The regularizer is small, but it prevents unstable coefficients when a regime contains few valid pixels or when two morphology features are strongly correlated.

Term by term, $\mathbf{X}_r^\top \mathbf{X}_r$ is the feature correlation matrix for regime r , $\lambda \mathbf{D}^\top \mathbf{D}$ adds a small diagonal stabilizer to every coefficient except the final $t_{41}n_{41}$ coefficient, and $\mathbf{X}_r^\top \mathbf{z}_r$ is the feature-target correlation vector in log spread space. This is ridge regression: it is the ordinary least-squares fit plus a penalty that discourages very large coefficients. The penalty is important here because many morphology features are related; for example, ρ_{41} , h_{41} , b_{41} and t_{41} all increase in dense high-rise areas.

Sufficient statistics are accumulated incrementally:

$$\mathbf{A}_r += \mathbf{X}_{r,s}^\top \mathbf{X}_{r,s} \quad \forall s \quad (3.74)$$

$$\mathbf{v}_r += \mathbf{X}_{r,s}^\top \mathbf{z}_{r,s} \quad \forall s \quad (3.75)$$

so the full design matrix is never materialised in memory. This is important in practice because the calibration artifacts are already large even after compression into fitted coefficients and profiles. The path loss LoS calibration JSON reused by the final model

is approximately 200 000 lines, and the spread calibration must be rerun separately for delay spread and angular spread whenever the feature set changes.

Fallback hierarchy. During fitting, each sampled pixel updates the exact accumulator and the four broader fallback accumulators listed below. A coefficient vector is written only for accumulators with at least $4 \times 23 = 92$ sampled pixels. During inference, the implementation tries the same keys in this order and uses the first one present:

1. exact regime: (metric, topology class, LoS/NLoS, antenna bin),
2. same topology class + LoS/NLoS + *all* antenna heights,
3. same topology class + *all* LoS states + all antenna heights,
4. global topology + LoS/NLoS + all antenna heights,
5. fully global fallback: (metric, global, all_los, all_ant), where all_los is the implementation key for all visibility states, not LoS only pixels.

Stage 4: inverse transform and clipping

$$\hat{y}(x) = \exp(\hat{z}(x)) - 1, \quad \hat{z}(x) = \mathbf{x}_{79}(x)^\top \hat{\mathbf{w}}_r^{79} \quad (3.76)$$

Regime dependent clipping is applied after the inverse transform:

$$\hat{y}_{\text{clip}}(x) = \min(y_{\text{max},r}, \max(0, \exp(\hat{z}(x)) - 1)), \quad (3.77)$$

with delay spread clipped to $[0, 400]$ ns and angular spread clipped to $[0, 15^\circ]$ in LoS or $[0, 90^\circ]$ in NLoS. The asymmetric angular limit is deliberate: LoS angular spread is dominated by the direct path and is physically closer to a near-zero spike plus small reflections, whereas NLoS angular spread can occupy a much wider support.


3.6.5 Spread prior results

Table 3.12: Delay spread: calibrated test RMSE (ns) by topology class.

Topology class	Overall	LoS	NLoS
open_sparse_lowrise	16.88	17.01	15.43
open_sparse_vertical	40.84	40.44	44.08
mixed_compact_lowrise	17.27	16.65	18.91
mixed_compact_midrise	37.88	38.26	37.25
dense_block_midrise	26.90	29.00	25.02
dense_block_highrise	43.52	40.16	45.08
Overall	28.10	27.49	29.62

Table 3.13: Angular spread: calibrated test RMSE (°) by topology class.

Topology class	Overall	LoS	NLoS
open_sparse_lowrise	10.01	10.30	6.33
open_sparse_vertical	12.70	12.98	9.97
mixed_compact_lowrise	14.37	15.99	8.17
mixed_compact_midrise	15.26	17.85	9.57
dense_block_midrise	15.53	21.20	8.17
dense_block_highrise	14.64	21.33	9.78
Overall	13.74	15.28	8.66

**Figure 3.8:** Pipeline for the delay and angular-spread priors. Both spread modalities share this identical linear architecture, predicting expected spread in the log-domain via regime wise ridge regression over 23 spatial and geometric features.

3.7 Training and Evaluation Protocol

3.7.1 Experimental protocol and reporting status

The final residual GMM-head model follows the same strict split philosophy as the recent distribution-first models. The split is a city-holdout split with `split_seed = 42`, `val_ratio = 0.15`, and `test_ratio = 0.15`. Metrics and losses are computed only on valid ground pixels, using the target-specific no data masks and the explicit LoS/NLoS masks. The model input contains nine spatial channels:

1. normalized topology,
2. LoS mask,
3. NLoS mask,
4. ground-mask,
5. combined path-loss prior,
6. LoS path-loss prior,
7. calibrated Try 78 NLoS branch map,
8. delay-spread prior,
9. angular-spread prior.

The seventh channel is named `path_loss_nlos_prior` in the code, but it should be read carefully. It is the calibrated Try 78 mixed/NLoS map computed by the regime wise path loss calibration, not a map that has been zeroed outside NLoS pixels. In the final combined path-loss prior and in the final output selection, only the explicit LoS/NLoS masks decide which regional branch is used.

The three mask channels are passed as binary maps. The topology channel is masked to ground pixels and divided by 90 m; the three path-loss prior channels are divided by 180 dB; and the delay and angular-spread priors are passed as $\log(1 + \text{prior}) / \log(1 + 400)$ and $\log(1 + \text{prior}) / \log(1 + 90)$, respectively. These scaling constants only affect the neural network input representation; the residual targets and reported metrics remain in the native units.

The scalar UAV height is not appended as a constant image channel. It is encoded with a sinusoidal embedding and injected into the network through FiLM conditioning. The selected final configuration uses a shared encoder decoder with `base_width = 128`, `cond_dim = 160`, three GMM components per-task and region, α -gate bias -2 , and fixed residual caps specified separately for path loss, delay spread and angular spread.

The selected final model reported in Chapter 4 uses these frozen path loss and spread priors as anchors. This section therefore documents both the frozen prior inputs and the residual architecture used by the final evaluation.

3.7.2 Prior Anchored GMM Residual Model

The final model introduces a joint multi-task residual GMM-head architecture that emits a *Gaussian Mixture Model (GMM)*-parameterized residual head at each pixel, rather than only a direct scalar output. For each task $k \in \{\text{PL}, \text{DS}, \text{AS}\}$, region $r \in \{\text{LoS}, \text{NLoS}\}$, component $j \in \{1, 2, 3\}$, and valid pixel x , the model predicts a mixture weight, a mean and a standard deviation. Path loss is modeled in native dB space, while delay spread

and angular spread are modeled in transformed $\log(1 + y)$ space inside the probabilistic head and converted back to native units for deterministic metrics. In the selected final fine-tuning stage the same GMM parameterization is kept, but the probabilistic losses are disabled; therefore the final reported claim is based on the deterministic mixture mean map rather than on calibrated predictive densities.

The native residual caps depend on task and region:

$$\delta_{\max,k,r} = \begin{cases} 7 \text{ dB}, & k = \text{PL}, r = \text{LoS}, \\ 25 \text{ dB}, & k = \text{PL}, r = \text{NLoS}, \\ 20 \text{ ns}, & k = \text{DS}, r = \text{LoS}, \\ 20 \text{ ns}, & k = \text{DS}, r = \text{NLoS}, \\ 17^\circ, & k = \text{AS}, r = \text{LoS}, \\ 15^\circ, & k = \text{AS}, r = \text{NLoS}. \end{cases} \quad (3.78)$$

The bounded native residual is assembled from a global map-level offset and a local pixel-level correction:

$$\Delta_{k,r,j}(x) = \text{clip}\left(\delta_{k,r,j}^{(g)} + \delta_{k,r}^{(\ell)}(x), -\delta_{\max,k,r}, \delta_{\max,k,r}\right). \quad (3.79)$$

The mean of each mixture component is then anchored to the frozen prior:

$$\mu_{k,r,j}^{\text{native}}(x) = y_k^{\text{prior}}(x) + \alpha_{k,r}(x)\Delta_{k,r,j}(x), \quad (3.80)$$

where:

- $y_k^{\text{prior}}(x)$ is the per-task frozen prior selected inside the shared output loop, then broadcast to the LoS and NLoS residual branches,
- $\delta_{k,r,j}^{(g)}$ is a map-level residual offset for component j ,
- $\delta_{k,r}^{(\ell)}(x)$ is a per-pixel residual refinement shared across components,
- $\alpha_{k,r}(x) \in (0, 1)$ is a learned spatial gate, initialized with bias -2 , so the model starts close to the prior,
- the residual is bounded to $\pm\delta_{\max,k,r}$ to prevent the DNN from deviating catastrophically from the prior.

For the spread tasks, the native anchored mean is mapped into transformed space before the Gaussian likelihood:

$$\mu_{k,r,j}(x) = \begin{cases} \mu_{k,r,j}^{\text{native}}(x), & k = \text{PL}, \\ \log\left(1 + \max\{\mu_{k,r,j}^{\text{native}}(x), 0\}\right), & k \in \{\text{DS}, \text{AS}\}. \end{cases} \quad (3.81)$$

The DNN also predicts per-pixel mixture probabilities $p_{k,r,j}(x)$ and combines global and local standard deviation terms into the component standard deviation:

$$\sigma_{k,r,j}(x) = \sqrt{\left(\sigma_{k,r,j}^{(g)}\right)^2 + \left(\tilde{\sigma}_{k,r}^{(\ell)}(x)\right)^2}. \quad (3.82)$$

The Gaussian variance used by the density is therefore $\sigma_{k,r,j}^2(x)$. Both standard deviation terms are constrained by sigmoid transforms to lie between $\sigma_{\min} = 0.05$ and $\sigma_{\max} = 3.0$.

These bounds stabilize the probabilistic head and prevent the NLL term, when it is enabled, from collapsing to unrealistically tiny uncertainty or exploding to a uselessly flat density. The local component logits are transformed with a per-pixel softmax, so $\sum_j p_{k,r,j}(x) = 1$. The global mixture weights $\pi_{k,r,j}^{(g)}$ emitted by the global branch belong to the map-level distributional branch used when map-level KL style terms are active; they are not multiplied into the deterministic point map in Eq. (3.83). This distinction is important for the selected final checkpoint, where the distributional losses are disabled and RMSE is computed from the local weighted point prediction.

The architecture supports a full GMM Negative Log Likelihood (NLL), and the preceding large-capacity residual stage used non-zero probabilistic weights. The selected final model keeps the same GMM-head, but the last residual fine-tuning stage is RMSE/MAE dominant: its NLL, KL, moment and outlier-budget weights are set to zero in the resolved YAML configuration. The deterministic point prediction used for RMSE evaluation is the mixture mean in the modelled space:

$$\hat{z}_{k,r}(x) = \sum_{j=1}^3 p_{k,r,j}(x) \mu_{k,r,j}(x), \quad \hat{y}_{k,r}(x) = \begin{cases} \hat{z}_{k,r}(x), & k = \text{PL}, \\ \exp(\hat{z}_{k,r}(x)) - 1, & k \in \{\text{DS}, \text{AS}\}. \end{cases} \quad (3.83)$$

For path loss, this is also the native-space mixture mean because the modelled space is dB. For delay spread and angular spread, the mixture is averaged in $\log(1 + y)$ space and then mapped back with $\exp(\cdot) - 1$. Therefore the reported point map should be read as the deterministic prediction used for RMSE, not as the exact native-space expectation of the probabilistic mixture. The final task map selects $\hat{y}_{k,\text{LoS}}$ or $\hat{y}_{k,\text{NLoS}}$ using the explicit CKM LoS/NLoS masks. This prevents a visually plausible NLoS correction from leaking into pixels whose geometric path is LoS, and conversely prevents LoS ring structure from being used as an explanation for blocked pixels.

The training loss also includes a prior-preservation term:

$$\mathcal{L}_{\text{guard},k} = \mathbb{E}_{x \in \Omega_k} \left[\max \left(0, (\hat{y}_k(x) - y_k(x))^2 - (y_k^{\text{prior}}(x) - y_k(x))^2 \right) \right] \quad (3.84)$$

This *pixel-wise* prior guard penalizes the model explicitly on any pixel where its squared error exceeds the prior's squared error. The guard is not used alone; it is part of the composite optimization objective in Sec. 3.7.4.

3.7.3 Architecture: Shared UNet with FiLM

The final model processes the map inputs through a shared multi-task convolutional encoder decoder. In the final large-capacity configuration, scalar UAV height is encoded into a condition vector $\mathbf{c} \in \mathbb{R}^{160}$ and injected dynamically into the convolutional feature maps at multiple scales via *Feature wise Linear Modulation (FiLM)* layers.

For a given convolutional feature map $\mathbf{X} \in \mathbb{R}^{C \times H \times W}$, a linear projection of the height condition $\mathbf{c}$ outputs a scale vector $\mathbf{g} \in \mathbb{R}^C$ and a shift vector $\mathbf{b} \in \mathbb{R}^C$. The feature map is then modulated via an affine transformation:

$$\mathbf{X}' = \mathbf{X} \odot (1 + \mathbf{g}) + \mathbf{b} \quad (3.85)$$

where $\mathbf{g}$ and $\mathbf{b}$ are broadcasted spatially across the height and width. The FiLM mechanism

and injection points are unchanged between the wide and large-capacity variants; only the width of the feature maps changes, so the FiLM projection emits $2C$ parameters for the larger channel counts.

The word “joint” in this model refers to shared representation learning and joint optimization, not to an autoregressive or cross output decoder. Path loss, delay spread and angular spread all pass through the same encoder, the same height conditioning path, the same decoder feature map and the same prior-anchored GMM assembly rule. After these shared features are computed, however, the implementation iterates over the three tasks and applies task-specific global and local heads. Therefore PL, DS and AS can exchange information through the common spatial representation, but the final residual parameters, mixture weights and frozen prior anchor are selected independently for each task.

The architecture explicitly enforces the separation of global and local effects via a dual head design for each task:

- **Global Head (Map Level Context):** The deepest bottleneck feature map is globally average pooled to produce a single vector representing the macro context. This pooled vector, fused with the height condition, is passed through an MLP to predict the non spatial quantities used by the distributional branch: auxiliary global mixture weights $\pi_{k,r,j}^{(g)}$, global offsets $\delta_{k,r,j}^{(g)}$, and global standard deviations $\sigma_{k,r,j}^{(g)}$.
- **Local Head (Pixel Level Refinement):** The high resolution output of the decoder is passed through 1×1 convolutions to predict the spatially varying fields: local component weights $p_{k,r,j}(x)$, $\delta_{k,r}^{(\ell)}(x)$, $\alpha_{k,r}(x)$, and $\tilde{\sigma}_{k,r}(x)$.




Figure 3.9: Schematic of the final large-capacity shared UNet architecture with FiLM conditioning. The backbone and decode logic are shared, while each task uses its own heads and frozen prior.

Detailed sinusoidal height embedding and FiLM path. The scalar UAV height h_{tx} is first converted into a 32-dimensional sinusoidal embedding:

$$f_i = \exp\left(\log 12 + \frac{i-1}{15}(\log 478 - \log 12)\right), \quad i = 1, \dots, 16 \quad (3.86)$$

$$\mathbf{e}(h_{tx}) = \left[\sin\left(\frac{h_{tx}}{f_1}\right), \dots, \sin\left(\frac{h_{tx}}{f_{16}}\right), \cos\left(\frac{h_{tx}}{f_1}\right), \dots, \cos\left(\frac{h_{tx}}{f_{16}}\right) \right] \in \mathbb{R}^{32} \quad (3.87)$$




Figure 3.10: Detailed sinusoidal height conditioning path.

Detailed prior-anchored GMM-head. For each task k and for each region $r \in \{\text{LoS}, \text{NLoS}\}$, the model predicts a 3-component mixture. The global head outputs $\pi_{k,r,j}^{(g)}$, $\delta_{k,r,j}^{(g)}$ and $\sigma_{k,r,j}^{(g)}$, while the local head outputs per-pixel maps $p_{k,r,j}^{(\ell)}(x)$, $\delta_{k,r}^{(\ell)}(x)$, $\alpha_{k,r}(x)$, and

$\tilde{\sigma}_{k,r}(x)$. After softmax over components, $p_{k,r,j}^{(\ell)}(x)$ becomes the local probability $p_{k,r,j}(x)$ used in the regional model space mean. The auxiliary global weights $\pi_{k,r,j}^{(g)}$ are retained for the distributional training branch and diagnostics, but the selected RMSE/MAE dominant checkpoint uses the local probabilities for the deterministic output map.




Figure 3.11: Detailed prior-anchored GMM-head used by the final model.

3.7.4 Composite Optimization Objective

The implementation defines a composite loss that can train both the probabilistic GMM-head and the deterministic residual map. For each task, let Ω_k be the valid ground pixel set, let $q_k(z | x)$ be the predicted mixture density in the modelled space, with $z = y$ for path loss and $z = \log(1 + y)$ for delay spread and angular spread, and let $\hat{y}_k(x)$ be the native-space point prediction used for RMSE. The total objective is:

$$\begin{aligned} \mathcal{L} = \sum_{k \in \{\text{PL, DS, AS}\}} w_k \Big[& \lambda_{\text{NLL}} \mathcal{L}_{\text{NLL},k} + \lambda_{\text{KL}} \mathcal{L}_{\text{KL},k} + \lambda_{\text{mom}} \mathcal{L}_{\text{mom},k} \\ & + \lambda_{\text{anchor}} \mathcal{L}_{\text{anchor},k} + \lambda_{\text{guard}} \mathcal{L}_{\text{guard},k} + \lambda_{\text{out}} \mathcal{L}_{\text{out},k} \\ & + \lambda_{\text{RMSE}} \mathcal{L}_{\text{RMSE},k} + \lambda_{\text{MAE}} \mathcal{L}_{\text{MAE},k} \Big]. \end{aligned} \quad (3.88)$$

Table 3.14 gives the resolved weights for both the large-capacity residual stage and the final residual fine-tuning stage; the equations below define the terms used in that table.

Map negative log likelihood. The likelihood is computed separately in LoS and NLoS regions, then averaged over the non-empty regions. For a transformed target $z_k(x)$,

$$q_{k,r}(z_k(x) | x) = \sum_{j=1}^3 p_{k,r,j}(x) \mathcal{N}(z_k(x); \mu_{k,r,j}(x), \sigma_{k,r,j}^2(x)), \quad (3.89)$$

and the regional NLL is

$$\mathcal{L}_{\text{NLL},k,r} = -\frac{1}{|\Omega_{k,r}|} \sum_{x \in \Omega_{k,r}} \log q_{k,r}(z_k(x) | x). \quad (3.90)$$

The implemented version evaluates Eq. (3.90) using $\log \sum \exp$ for numerical stability.

Table 3.14: Active loss settings in the final residual GMM-head model stages.

Term group	Large-capacity residual stage	Final residual fine-tuning stage	Interpretation
NLL, KL, moment and outlier-budget	1.00 / 0.25 / 0.10 / 0.02	0 / 0 / 0 / 0	Active in the preceding large-capacity stage, then disabled in the selected deterministic fine-tuning stage.
RMSE / MAE	1.50 / 0.10	1.00 / 0.05	Supervised pixel losses in native output units; the final fine-tuning stage makes them the dominant active terms.
Anchor / prior guard	0.05 / 0.10	0.01 / 0.02	Keep residual corrections small and penalize model states that become worse than the frozen priors, with lighter regularization in the fine-tuning stage.
Task weights	PL 1.00, DS 1.00, AS 1.00	PL 1.00, DS 0.25, AS 0.25	Equal target weighting before the final fine-tuning stage; then path loss becomes the main optimization target while spread supervision remains in the joint model.

Distributional KL term. The global branch predicts residual components $\delta_{k,r,j}^{(g)}$. To keep those components aligned with the empirical residual distribution of the current batch, residuals $\epsilon_k(x) = y_k(x) - y_k^{\text{prior}}(x)$ are softly binned into an empirical probability vector $\mathbf{u}_{k,r}$. The global mixture offsets define a predicted residual histogram $\mathbf{v}_{k,r}$. The loss is

$$\mathcal{L}_{\text{KL},k} = \frac{1}{2} \sum_{r \in \{\text{LoS}, \text{NLoS}\}} \sum_{b=1}^B u_{k,r,b} \log \frac{u_{k,r,b}}{v_{k,r,b}}, \quad B = 64. \quad (3.91)$$

This term does not force a perfect pixel-wise correction. It only encourages the map-level mixture to place probability mass where the batch residuals actually live.

Moment matching. The point prediction is also constrained to have approximately the same first two moments as the target map:

$$\mathcal{L}_{\text{mom},k} = (\bar{\hat{y}}_k - \bar{y}_k)^2 + (s_{\hat{y},k} - s_{y,k})^2, \quad (3.92)$$

where bars denote masked means over Ω_k and s denotes masked standard deviation. When active, this term prevents a low-NLL solution from matching local density while drifting in map-average level. In the final fine tune it is defined by the code but has zero weight.

Anchor and guard terms. The anchor penalty keeps the correction small:

$$\mathcal{L}_{\text{anchor},k} = \mathbb{E} \left[\alpha_k(x)^2 + 0.25 \left(\delta_k^{(g)} \right)^2 + 0.50 \left(\delta_k^{(\ell)}(x) \right)^2 \right]. \quad (3.93)$$

The guard term is Eq. (3.84). Together, these two terms encode the final design philosophy: the DNN is allowed to improve the prior, but it must pay a penalty for large or harmful deviations. In Eq. (3.93), the residual variables are the same tensors used by the implementation: native dB offsets for path loss, and the corresponding $\log(1+y)$ -space offsets for delay and angular spread after the native residual caps have been applied. The guard term, by contrast, is evaluated in the native units of each task against the frozen prior map.

Outlier budget and deterministic losses. When the outlier-budget term is active, the implementation uses the last local mixture component as a proxy for broad/outlier usage. This is an indexing convention rather than a hard architectural constraint: the model is not forced to make that component the broadest one. If the mean usage of this proxy exceeds $\tau = 0.60$, a hinge penalty is applied:

$$\mathcal{L}_{\text{out},k} = \max\left(0, \frac{1}{|\Omega_k|} \sum_{x \in \Omega_k} p_{k,\text{out}}(x) - \tau\right). \quad (3.94)$$

The RMSE and MAE terms are conventional masked deterministic losses:

$$\mathcal{L}_{\text{RMSE},k} = \sqrt{\frac{1}{|\Omega_k|} \sum_{x \in \Omega_k} (\hat{y}_k(x) - y_k(x))^2}, \quad \mathcal{L}_{\text{MAE},k} = \frac{1}{|\Omega_k|} \sum_{x \in \Omega_k} |\hat{y}_k(x) - y_k(x)|. \quad (3.95)$$




Figure 3.12: Composite final loss family. In the final residual fine-tuning stage, the selected final model keeps the GMM-head and uses RMSE/MAE, anchor and guard penalties; NLL, KL, moment and outlier-budget weights are set to zero. Earlier residual stages used non-zero probabilistic weights.

Results

This chapter reports the final test results of HARP-NET CKM, the joint prior-anchored residual model used in the reduced thesis. The standalone path-loss prior audit uses the city-holdout validation+test pool: 5340 samples from 29 unseen cities. The selected final model is then evaluated on the final 2590-sample, 14-city test split. Two path-loss prior numbers therefore appear in this chapter: 1.9246 dB is the standalone calibrated hybrid prior on the held-out validation+test pool, whereas 1.9383 dB is the frozen path-loss prior baseline used by HARP-NET CKM on the final 14-city test split. All final model-versus-prior comparisons use the latter.

The prior is not a weak reference included only for comparison. It is one of the central results of the thesis: a strong, train-calibrated baseline that absorbs the stable part of the map before the residual model is asked to solve the harder local errors. Appendix A keeps the development history and negative attempts; this chapter focuses on the frozen priors and final test evaluation.

4.1 Final test contract and frozen priors

The final path-loss prior is the clearest evidence that the prior is a strong path loss baseline in its own right. The enhanced two-ray LoS model reduced LoS RMSE from roughly 3.95 dB with FSPL to 1.75 dB. The final hybrid prior reached 1.9246 dB overall under the pixel-weighted metric on the held-out validation+test pool.

4.1.1 LoS path-loss prior progression

The LoS branch of the final path-loss prior was developed as a short waterfall of three increasingly physical models fitted on the same city-holdout split, so the gain from each modelling choice can be isolated (Table 4.1).

Table 4.1: *LoS path-loss prior progression: three priors evaluated on the same held-out validation+test pool (pixel-weighted RMSE).*

Prior	LoS RMSE (dB)	Gain over previous
FSPL only	3.9540	N/A
FSPL + radial residual	3.7086	−0.25
Enhanced two-ray	1.7516	−1.96

Two lessons follow. First, adding a purely radial correction on top of FSPL only buys about 0.25 dB. The residuals that a non-coherent model can absorb are small compared to the structured deviation from FSPL. Second, switching to the *enhanced* two-ray model with a fitted effective reflection coefficient (ρ, ϕ) per 5 m UAV-height bin removes almost 2 dB of additional error. This is exactly the magnitude expected when the dataset contains a visible ground-reflection interference pattern, because the radial-residual model cannot

reproduce the oscillatory ring structure with a purely monotonic function of distance. The two-ray model is therefore not just a better fit numerically; it is a physically motivated explanation of the dominant LoS residual shape in CKM, and it is why the final model uses the two-ray output directly as a frozen anchor rather than re-learning it from pixels.

4.2 Path Loss Results

Split comparability note. The final priors and HARP-NET CKM use the same `city_holdout` contract with `val_ratio = 0.15`, `test_ratio = 0.15`, and `split_seed = 42`. The standalone path-loss prior audit uses the training cities and the full validation+test held-out pool, giving 1.9246 dB over 5340 held-out samples. The final model comparison remains test only, so the frozen path-loss prior baseline in Table 4.7 is 1.9383 dB.

The calibrated path-loss prior is the clearest example of the final design. It is not a single generic propagation formula. It is split into a LoS branch and an NLoS branch because those pixels have different error mechanisms. The LoS branch extracts direct and reflected distances, FSPL, a coherent two-ray correction, height-bin bias terms, and a smoothed radial residual profile. The NLoS branch extracts a calibrated feature vector containing the raw NLoS prior, distance, local building-density, local building-height, local NLoS support, a shadowing proxy, normalized elevation angle, and interactions between those terms:

$$\mathbf{x}_{78} = [P_0^2, P_0, \log(1 + d_{2D}), \rho_{15}, \rho_{41}, h_{15}, h_{41}, \\ \rho_{41} \log(1 + d_{2D}), n_{15}, n_{41}, n_{41} \log(1 + d_{2D}), \\ \sigma_{\text{shadow}}, \theta_{\text{norm}}, n_{41}\theta_{\text{norm}}, 1]^\top.$$

The weights are fitted using training data only and separately by propagation regime. This is why the final prior is a calibrated feature extractor, not just an analytic baseline.

Second, RMSE and MAE told different stories during the middle of the project. For example, the prior-confidence stage had moderate average absolute error on many samples, but RMSE remained high because a smaller number of hard NLoS cases produced very large misses. This is why the later work moved toward tail-aware and distribution-aware methods rather than only lowering the mean error.

Third, LoS and NLoS must be reported separately. A single overall number can hide whether the model is improving direct-visibility pixels, blocked pixels, or only the easier majority regime. For that reason the final tables below always separate LoS/NLoS where the target definition permits it.

Table 4.2: Final hybrid path-loss prior, pixel-weighted RMSE on the held-out validation+test split.

Region	RMSE	Unit
LoS	1.7516	dB
NLoS	3.3967	dB
Overall	1.9246	dB

The final spread priors showed that a large part of the spread targets can also be explained by geometry, building layout, and calibration without a neural network. The fitted calibration dictionary contains 114 case keys, but this number is not one simple Cartesian

product. For each spread metric, the exact grid has 6 topology classes $\times$ 2 LoS/NLoS states $\times$ 3 antenna height bins, or 36 exact keys. The implementation then adds fallback calibrations for sparse cases: 12 antenna collapsed topology and LoS/NLoS keys, 6 topology-wide keys, 2 global LoS/NLoS keys, and 1 fully global key. This gives 57 keys per spread metric and 114 keys across delay spread and angular spread. Equivalently, the dictionary contains 72 exact DS/AS cases plus 42 fallback cases. The point of this dictionary is not novelty by formula alone, but coverage: the baseline is strong because it stores many small regime-specific decisions about density, building-height, visibility, and antenna height instead of forcing one global estimator to explain all pixels. The calibrated spread-prior module applies this idea separately to delay spread and angular spread. The raw spread estimate is only the first feature; the calibrated model then uses a 23-term vector with raw prior curvature, range, elevation angle, UAV height, local building-density at two scales, local height at two scales, LoS/NLoS support at two scales, blocker depth, transmitter clearance, taller-building fraction, and interaction terms. Both spread targets are fitted in $\log(1 + y)$ space, so the model can handle the LoS spike near-zero and the long NLoS tail with one calibrated framework while still reporting metrics in native ns and degrees.

Table 4.3: Calibrated spread prior test results, pixel-weighted RMSE.

Metric	Overall	LoS	NLoS
Delay spread (ns)	28.10	27.49	29.62
Angular spread (°)	13.74	15.28	8.66

4.2.1 RMSE by UAV transmitter height

The CKM dataset covers stored UAV heights from about 12 m to almost 478 m, while the closest nominal beta law has support up to 500 m. The absence of samples above 478 m is statistically normal for this finite draw, because the expected count above that height is only 0.52. Therefore the conditional error at a single aggregate RMSE can hide strong altitude-dependent structure. Tables 4.4 and 4.5 therefore break the path-loss prior and the spread-prior delay-spread and angular-spread priors into four UAV-height bins, using the same pixel-weighted aggregation as the headline tables.

Table 4.4: Path loss prior pixel-weighted RMSE by UAV transmitter height, on the held-out validation+test split.

UAV-height bin	Overall	LoS	NLoS	Samples
12 to 50	2.180	1.797	3.756	1211
50 to 150	1.936	1.785	3.228	2634
150 to 300	1.680	1.666	2.259	786
300 to 500	1.627	1.619	2.255	362

The altitude dependence is monotonic and physically intuitive. Higher UAV transmitters see more of the ground as unobstructed LoS and encounter fewer deep shadow pockets, so both the NLoS pixel fraction and the worst-case tails shrink. Path loss NLoS RMSE drops from 3.76 dB at low-altitudes to 2.26 dB above 300 m, and delay-spread NLoS RMSE collapses from 35.13 ns to 2.76 ns in the same comparison. The opposite is true for the

Table 4.5: Delay-spread and angular-spread calibrated pixel-weighted RMSE by UAV transmitter height, on the city-holdout test split.

UAV-height bin	Metric	Overall	LoS	NLoS
12 to 50	Delay spread (ns)	39.90	43.04	35.13
50 to 150	Delay spread (ns)	27.17	26.72	28.29
150 to 300	Delay spread (ns)	8.70	8.79	8.13
300 to 500	Delay spread (ns)	3.70	3.80	2.76
12 to 50	Angular spread (°)	15.69	18.81	9.93
50 to 150	Angular spread (°)	14.19	15.89	8.41
150 to 300	Angular spread (°)	10.39	11.07	4.49
300 to 500	Angular spread (°)	8.85	9.34	2.85

absolute number of samples: the dataset is dominated by the 50 – 150 m band, so the global aggregate RMSE reflects that regime more than the higher-altitude extremes. Together with the empirical HDF5 histogram of Eq. 3.1 and the counts reported in Chapter 3, this means the high-altitude rows are better despite having much less training support, not because they are over-represented in the dataset. A practical consequence for the final residual model is that the neural residual head has less to correct at high-altitudes, where the path-loss prior is already close to the quantisation floor. The remaining headroom lives in the low-altitude, NLoS-heavy bin, which is exactly where a distribution-aware residual can matter.

4.2.2 Path loss prior by topology class

The headline path loss-prior number (1.92 dB overall RMSE on the held-out validation+test split, see Table 4.2) hides the same building layout structure that the development experiments surfaced. The per-sample evaluation records were re-aggregated by classifying each evaluated sample into the six topology classes with the same routing rule used by the development experiments. The pixel-weighted RMSE per topology class is reported in Table 4.6 and visualised in Figure 4.1.

Table 4.6: Hybrid path-loss prior, pixel-weighted RMSE on the held-out validation+test split, broken down by topology class. The overall row reproduces the RMSE values in Table 4.2; the sample count is the sum of samples assigned to the six topology classes. The aggregate Try 78 evaluator contains 4994 topology-assigned samples.

Topology class	LoS	NLoS	Overall	Samples
open sparse lowrise	1.39	2.92	1.43	895
open sparse vertical	1.43	3.91	1.55	370
mixed compact lowrise	1.81	2.91	1.90	1240
mixed compact midrise	1.91	3.84	2.21	1060
dense block midrise	2.19	3.20	2.40	1325
dense block highrise	2.44	4.09	2.97	104
Overall	1.75	3.40	1.92	4994

The LoS dependency on building layout is monotonic. The two-ray prior fits its own ground-reflection ring against an LoS marginal whose width grows with city density

and building-height, so the LoS RMSE of the analytic model rises from 1.39 dB in the open sparse lowrise class to 2.44 dB in the dense block highrise class. The NLoS branch is less ordered by building layout: the best Try 78 NLoS class is *mixed compact lowrise* at 2.91 dB, while the worst is *dense block highrise* at 4.09 dB. The dataset is also unbalanced in this respect: dense block highrise contributes only 104 of the 4994 topology-assigned evaluation samples, so its high RMSE has a relatively small impact on the overall pixel-weighted aggregate but is the regime most likely to limit deployment in dense urban environments.

This per-topology view also makes the central methodological claim of the thesis quantitative. The right panel of Figure 4.1 shows that the calibrated prior remains stable across the six building layout classes. This is the relevant main text result: the train calibrated prior is not only strong on average, but also auditable by city type and LoS/NLoS case. The longer comparison with the distribution-first diagnostic models is kept in Appendix A.




Figure 4.1: Path loss prior by topology class. Left: LoS, NLoS and overall pixel-weighted RMSE for each of the six topology classes. The LoS RMSE grows monotonically with building-density. Right: NLoS RMSE shown by building layout class, highlighting where the calibrated prior is easiest and hardest.

4.2.3 Path loss: prior versus HARP-NET CKM

HARP-NET CKM freezes the path loss and spread priors and learns a bounded residual correction for the three radio maps. It is tested on the city-holdout test split with 2590 samples from 14 cities that were not seen in training. The model should be read as a two-stage final pipeline, not only as the last fine-tuning loss: the large-capacity residual/GMM stage keeps NLL, KL, moment matching and outlier-budget terms active, while the selected reported checkpoint is the validation-selected state after the deterministic RMSE/MAE-dominant fine-tuning stage. The exact weights of both stages are listed in Table 3.14. The final claim is therefore an RMSE improvement over the frozen priors rather than a calibrated probabilistic prediction claim. For a target t , group G , and valid pixel set $\Omega_{G,t}$, the reported model RMSE is

$$\text{RMSE}_{\text{model}}(G, t) = \sqrt{\frac{1}{|\Omega_{G,t}|} \sum_{(i,p) \in \Omega_{G,t}} (\hat{y}_{i,p,t} - y_{i,p,t})^2}. \quad (4.1)$$

The prior RMSE uses the same formula with the frozen-prior prediction $y_{i,p,t}^{\text{prior}}$ in place of $\hat{y}_{i,p,t}$. Therefore `model_rmse` in the evaluation JSON is target-specific: it is the RMSE for path loss, delay spread, or angular spread individually, not a sum of the three RMSE values. The three targets are kept separate because their units are incompatible.

Table 4.7: HARP-NET CKM test results on unseen cities, compared with the frozen priors. Negative Δ means that the model improves the prior.

Target	Model RMSE	Prior RMSE	Δ RMSE	Model MAE	Prior MAE	Δ MAE	Unit
Path loss	1.6519	1.9383	-0.2865	0.9298	1.2319	-0.3020	dB
Delay spread	26.5570	28.1023	-1.5453	7.9897	7.3033	+0.6864	ns
Angular spread	11.3854	13.7416	-2.3562	5.0088	4.9890	+0.0198	°

Table 4.8: HARP-NET CKM test RMSE separated by LoS and NLoS pixels. The test cities are unseen during training.

Target	Scope	Model RMSE	Prior RMSE	Δ RMSE
Path loss	LoS	1.4675	1.7519	-0.2844
Path loss	NLoS	3.1482	3.5143	-0.3661
Delay spread	LoS	25.4419	27.4851	-2.0432
Delay spread	NLoS	29.2045	29.6157	-0.4112
Angular spread	LoS	12.3510	15.2818	-2.9308
Angular spread	NLoS	8.4424	8.6614	-0.2191

Path loss is the cleanest final result because both RMSE and MAE improve over the frozen prior on unseen test cities. Overall RMSE drops from 1.9383 dB to 1.6519 dB, while MAE drops from 1.2319 dB to 0.9298 dB. The LoS and NLoS rows in Table 4.8 show that the improvement is not coming from only one visibility case: LoS improves by 0.2844 dB RMSE and NLoS by 0.3661 dB.

4.3 Angular Spread Results

For angular spread, the frozen prior already meets the original project target, but HARP-NET CKM further reduces RMSE on unseen test cities. Overall RMSE drops from 13.7416° to 11.3854°. The MAE is almost unchanged (+0.0198°), which means the neural residual mainly suppresses larger errors rather than shifting every pixel uniformly. The LoS row improves strongly, from 15.2818° to 12.3510°, while NLoS also improves slightly, from 8.6614° to 8.4424°.

No reviewed external method reports the same dense CKM city-holdout angular spread map task. One or two papers have similar link level channel statistics goals, but not the same dense map input, UAV height protocol, or held-out city evaluation. For that reason, the angular-spread result is compared mainly against the frozen prior and the original thesis target.

4.4 Delay Spread Results

Delay spread also improves in RMSE, although it is the only output where RMSE and MAE move in opposite directions. Overall RMSE drops from 28.1023 ns to 26.5570 ns, while

MAE increases from 7.3033 ns to 7.9897 ns. This is not a contradiction: RMSE weights rare high-delay errors quadratically, so suppressing the largest artefacts can improve RMSE while many tiny local shifts slightly increase MAE. The LoS RMSE improves by 2.0432 ns; the NLoS RMSE improves by 0.4112 ns.

As with angular spread, no reviewed paper is directly equal to the dense CKM city-holdout delay spread map task. The nearest references predict scalar channel characteristics or use different simulation contracts, so the main comparison is the frozen delay-spread prior plus the original 50 ns project target.

The validation metrics are not the final test claim. They are an external validation re-evaluation of the same selected checkpoint on the validation cities. The same selected model is then evaluated on the separate test cities, so the table checks validation/test drift without changing checkpoint or metric. Table 4.9 reports the two scopes side by side.

Table 4.9: *Validation-versus-test sanity check for the same selected final model. The validation row re-evaluates the selected checkpoint on validation cities; the test row is the final unseen-city generalisation result.*

Split	Samples	PL RMSE	DS RMSE	DS prior RMSE	AS RMSE	Meaning
Validation	2750	1.6157	22.1061	23.4494	11.3782	External validation re-evaluation of the selected model state.
Test	2590	1.6519	26.5570	28.1023	11.3854	Separate 14-city-holdout; this is the thesis generalisation number.

The larger delay spread gap between validation and test is mainly a city mix effect, not a change of model or metric. The frozen delay-spread prior also worsens from 23.45 ns RMSE on validation to 28.10 ns on test, so the test cities contain harder delay spread distributions before the neural residual is applied. In particular, Kuala Lumpur contributes only 13.2% of the valid delay spread test pixels but 36.2% of the test delay spread squared error, with 44.00 ns model RMSE and 46.93 ns prior RMSE. This is why the final test delay-spread RMSE is higher than the validation re-evaluation row, while still improving over the prior.

Together, the three target sections show the main generalisation result: the test cities are completely unseen, and the residual model reduces RMSE for path loss, angular spread and delay spread simultaneously.

4.5 State of the Art Context and Runtime

4.5.1 Target vs final result

The selected final model clears the original numerical targets on the unseen-city test split: 1.65 dB path loss against the 5 dB target, 26.56 ns delay spread against the 50 ns target, and 11.39° angular spread against the 20° target. The same target check also holds on validation and on the all-city aggregation; even the worst unseen-test subexpert for each output remains below its target (2.66 dB, 36.89 ns, and 14.06°). On the all-city subexpert audit, the corresponding worst cases are 2.45 dB, 42.24 ns, and 13.97°, still below the same targets.

4.5.2 Comparison with closest state of the art metrics

The state of the art numbers cannot be used as a direct leaderboard because their datasets, height assumptions, city splits, and target definitions differ. Still, Table 4.10 places the

final test result next to the closest available metrics from Section 2.5.7. Path loss has several external dB-scale references, but no external paper found here reports the same dense CKM city-holdout contract for delay spread and angular spread. The table is therefore limited to path loss like external rows plus the final-model reference row; scalar channel-parameter papers are kept out of this results table. In the table, PL means path-loss RMSE and lower RMSE is better. The better/worse ratios should be read cautiously because the rows use different data, different inputs, or different evaluation contracts. Earlier internal baselines are discussed in the ablation results instead of repeated here.

Table 4.10: Plain language comparison with the closest available path loss side metrics. No scalar indoor or link level spread predictor is included as a table row, because those numbers are not dense CKM-style map benchmarks.

Reference	What it reports	Verdict for our final model	Why the verdict is limited
RadioUNet [19]	1.6 dB–3.1 dB equivalent path-loss RMSE in accurate-map settings	Our final PL RMSE lies inside the reported 1.6 dB–3.1 dB dB-scale band.	2D path-loss maps only; no UAV height conditioning or spread targets.
RadioGUNet [20]	1.304 dB DPM, 1.936 dB IRT, and 1.392 dB DPM with cars on RadioMapSeer	Our final PL RMSE is close to the reported RadioMapSeer dB-scale range, while remaining above its best DPM number.	Outdoor simulated terrestrial maps; held-out test layouts, but no UAV height or spread targets.
PMNet / ICASSP 2023 winner [21, 6]	approx. 1.38 dB converted from the RadioMap3DSeer normalized score	Final PL is close but higher under a different benchmark contract; use only as scale context, not as a direct ranking.	Not directly comparable: ICASSP/RadioMap3DSeer testing differs from our unseen CKM-city split, continuous UAV height conditioning, and joint PL/DS/AS target contract.
Gao et al. [37]	5.59 dB on ICASSP 2023; 12.19 dB on ITU Challenge 2024, with about 60% fewer FLOPs than PPNet	Our final PL RMSE is about 3.4× lower than the ICASSP number and 7.4× lower than the ITU number.	Recent outdoor PL-map model with Tx/Rx depth and corridor weighting, but no CKM variable-height UAV, no joint PL/DS/AS output, and different split/features.
ReVeal PINN [27]	1.95 dB outdoor rural/suburban radio-environment RMSE with sparse measurements	Final PL is in the same dB-scale range as ReVeal’s sparse-measurement result.	Real outdoor measurements, but sparse RSSI radio-environment mapping rather than dense UAV height-conditioned CKM generation.
AIRMap [36]	path-gain error reported below 4 dB against ray tracing	Our final PL RMSE is below 2 dB, but an exact ratio is not recoverable from the reported path-gain threshold.	Reports path gain, not the same path loss calibration.
RMTransformer [22]	0.007148 normalised pixel RMSE, approximately 1.82 dB with its 254 dB USC/PMNet scaling window; 0.008099 channel prediction error, approximately 2.06 dB	Our final PL RMSE is slightly below the converted dB-scale values under that paper’s normalisation window.	Recent dense path loss radio map reference, but on 256 px USC/PMNet maps with a random 90/10 split, no CKM city-holdout, no continuous UAV height, and no spread targets.
PathFinder [39]	0.0331 normalized RMSE on RM3D DS-RPP, approximately 1.19 dB with the 36 dB RM3D image-window convention; 0.3263 on unseen rural, approximately 11.75 dB only under the same assumption	Our final PL RMSE is above the converted DS-RPP value and far below the rural value under the same approximate 36 dB assumption.	The DS-RPP conversion follows RM3D scaling; the rural conversion is only an approximate scale check because the paper reports normalised image-space RMSE.
Final model	1.6519 dB PL , 26.5570 ns DS , 11.3854° AS on unseen CKM cities	Reference row for the ratios above.	Only row here reporting all three dense CKM targets under the final city-holdout protocol.

This comparison supports a cautious interpretation. The final path loss number is not claimed to beat every published benchmark in absolute terms; PMNet’s reported

ICASSP score is lower on a different contract, and RadioGUNet is lower on the easiest RadioMapSeer setting. The path-loss result is nevertheless in the same broad outdoor simulated-map band as RadioUNet and RadioGUNet, below the PathFinder unseen-rural value under the approximate RM3D scaling, numerically well below Gao et al.’s recent corridor-weighted segmentation numbers, close to ReVeal’s real outdoor sparse-measurement scale, and slightly below the converted RMTransformer value after applying that paper’s own 254 dB normalisation window. For spreads, the honest external reading is that the table has no direct dense counterpart: the scalar channel-parameter papers reviewed in the state of the art are too different in input, output, and evaluation protocol to rank here. The thesis claim is therefore not that the spread maps beat external scalar predictors; it is that the model reports dense delay-spread and angular-spread RMSE on unseen CKM cities, improves the closest same-protocol spread baselines, and keeps the stricter ground-mask, LoS/NLoS, and continuous-height protocol visible.

4.5.3 Inference time sanity check

The secondary runtime objective was checked on the Barcelona geometry supplied with the MATLAB ray-tracing script. The benchmark uses the same local laptop as the final generator: an AMD Ryzen 5 5600H CPU and an RTX 3050 Ti Laptop GPU with 4 GiB of VRAM. The ray tracing run was configured at 1 m receiver spacing, so each non-building pixel of the 513×513 area is represented by a receiver at 1.5 m. Two Barcelona layouts and five UAV heights per layout give ten full-resolution CKM samples. The local MATLAB installation does not include Parallel Computing Toolbox, so the ray tracer was run on CPU with the same propagation settings as the supplied script: one reflection, no diffraction, 7.125 GHz, and extraction of path loss, delay spread, and angular spread from the ray-tracing paths.

The surrogate timing uses the same ten Barcelona topologies and antenna heights. It includes the complete non-RT generation path after the checkpoint is loaded once: geometric LoS/NLoS mask generation from the height map, the final ML-calibrated path-loss prior, the final ML-calibrated delay-spread and angular-spread priors, and mixed-precision inference with the final joint prior-anchored residual GMM-head model, including the default uncompressed .npz numerical archive export. The excluded costs are the one-time checkpoint load, the one-time CUDA warm-up, and optional PNG exports for masks or rendered maps.

Table 4.11: Full-pipeline runtime comparison for one 513×513 Barcelona CKM sample. Both rows are measured on the same Ryzen 5 5600H / RTX 3050 Ti Laptop machine. The surrogate row includes geometric mask generation, prior computation, and prediction of path loss, delay spread, and angular spread, including the default numerical .npz export. Checkpoint loading, one-time CUDA warm-up, and optional PNG exports are excluded.

Method	Timed path	Resource note	Time per-sample	Dataset-scale extrapolation
MATLAB ray-tracing	Pixel level receivers outside buildings; PL, DS and AS parsed from paths	Ryzen 5 5600H CPU; GPU off, PCT unavailable	99.89 s	448.96 h (≈ 18.7 days), $1\times$
Final prior-anchored generator	Generated LoS/NLoS masks, final calibrated priors, PL/DS/AS prediction, and default .npz export	RTX 3050 Ti Laptop GPU; mixed fp16, $b = 1$	0.88 s	3.95 h, $113.7\times$ faster

The ten ray-traced samples contain 156,088 and 115,825 valid receiver pixels in the two

Barcelona layouts, respectively. The measured average is 99.89 s per full-resolution sample for ray tracing and 0.88 s per sample for the complete final prior-anchored generator, including masks, priors, predictions, and the default numerical .npz archive export, but not optional PNG exports. The resulting $113.7\times$ speed-up exceeds the original 10 to $100\times$ runtime goal and should therefore be read as a deployment-style marginal runtime result after training and calibration, not as a comparison of the full research cost of developing the model.

The component logs explain where the remaining runtime goes without needing a second table: mask generation takes about 0.30 s, torch priors about 0.04 s, and final neural prediction plus default .npz export about 0.54 s. The vectorized mask path was checked against the original GT ray-caster with zero LoS/NLoS differences on the Barcelona samples and on 100 exported CKM samples; mixed precision changed RMSE by only 0.0014 dB, 0.0028 ns, and 0.0025° relative to fp32.

4.6 Qualitative Analysis

The qualitative evidence supports the quantitative story without needing the full diagnostic gallery in the main body. The early path loss histograms showed NLoS mode-collapse: visually plausible maps could follow the radial LoS trend while missing the blocked-link tail. The selected final model removes the worst of that collapse, reduces RMSE relative to the frozen priors across the split and LoS/NLoS checks, and applies residual corrections whose sign follows the prior-error maps. The complete histogram figures, split/regime bar plots, and representative inspection panel are kept in Appendix D.

The main lesson is that LoS path loss in this dataset has a strong radial ring structure, so a calibrated two-ray prior is more efficient than asking a CNN to learn the same pattern from pixels. The remaining qualitative failures are concentrated in sparse NLoS and high-spread cases, which is why future spread work should include an explicit high-tail or outlier diagnostic.

4.7 Limitations

The dataset is produced by ray tracing, not by field measurements. The results therefore measure agreement with a simulator and not direct over-the-air deployment performance. The path-loss target is stored as uint8, which imposes approximately 1 dB quantization. The spread targets are also stored as integer maps: delay spread as uint16 in ns and angular spread as uint8 in degrees. A more suitable dataset for the final modelling stage would preserve the continuous ray-traced targets, for example as float32 or as documented scaled integers, before any display-oriented quantisation. This should not be confused with a quantized model output: the final residual GMM-head model emits continuous PL/DS/AS values. In that light, the final path loss MAE of 0.9298 dB on labels stored at approximately 1 dB resolution is already close to the path-loss label step, even though the RMSE remains 1.6519 dB because it is dominated by the harder tail errors.

The effect can be written explicitly. Let y^* be the continuous ray-traced value and $y_q = Q_\Delta(y^*) = y^* + \epsilon_q$ the stored value after rounding to step Δ . Under the usual high-resolution quantisation approximation, $\epsilon_q \sim \mathcal{U}(-\Delta/2, \Delta/2)$, so

$$\sigma_q = \sqrt{\mathbb{E}[\epsilon_q^2]} = \frac{\Delta}{\sqrt{12}}. \quad (4.2)$$

With $\Delta = 1$ in the physical units used here, the hidden uncertainty is 0.289 dB for path loss, 0.289 ns for delay spread, and 0.289° for angular spread. This is not the dominant error in the reported final model, but it is an irreducible resolution limit of the saved labels and becomes important once residual errors approach the sub-unit regime.

Table 4.12: *Approximate quantisation RMSE implied by the stored integer target precision in the current CKM HDF5 file. The table reports the actual storage used in this thesis.*

Target	Stored type	Current integer step	Current floor
Path loss	uint8	1 dB	0.289 dB
Delay spread	uint16	1 ns	0.289 ns
Angular spread	uint8	1°	0.289°

This also explains why getting very close to the best possible error is mathematically harder than the headline RMSE suggests. If model error and quantisation error are approximately independent, the observed squared error decomposes as

$$\mathbb{E}[(\hat{y} - y_q)^2] \approx \mathbb{E}[(\hat{y} - y^*)^2] + \sigma_q^2.$$

Near σ_q , a real reduction in model error produces only a small change in total RMSE because of the square-root compression. In addition, many different continuous ray tracing values collapse to the same stored integer, so the model cannot learn sub-bin structure from the labels even if the input geometry contains enough information. The heavy-tailed NLoS distributions add a second difficulty: MSE is dominated by rare large misses, so once the common LoS pixels are near the quantisation scale, further improvements depend on resolving sparse hard NLoS cases rather than on smoothing the whole map.

The LoS/NLoS mask used in these experiments should be interpreted as a geometric binary visibility mask, not as an augmented electromagnetic propagation label. In the CKM setup the transmitter position is fixed at the map centre and the UAV height is known for each sample, so the same binary mask can be derived by line-of-sight tracing through the building topology. It is therefore a deterministic geometric feature and branch key, not extra measured radio information unavailable at deployment.

The selected final model is evaluated on one city-holdout split and one training seed. This demonstrates generalisation to unseen cities, but not full robustness: future-work should repeat selection across seeds and add field measurements, especially for delay-spread tails where quantised labels make sparse high-delay artefacts difficult to interpret.

Finally, the dataset fixes frequency and simulator assumptions. Deployment would need conditioning on carrier frequency, bandwidth, antenna pattern, and environment class; the city-holdout protocol tests urban layouts, not this broader domain adaptation.

Sustainability Analysis and Ethical Implications

This chapter follows the ETSETB sustainability guide structure, but in a condensed form compatible with the reduced thesis. The work does not manufacture a physical product; its deployable result is a software pipeline for faster CKM prediction. The analysis therefore separates the one time academic development cost from the possible operational use of the trained generator.

5.1 Sustainability matrix

Table 5.1: *Sustainability matrix summary for the thesis.*

Perspective	Development of BT	Project execution	Risks and limitations
Environmental	GPU training, laptop use, storage, and writing.	Reduced need for repeated ray tracing simulations once the surrogate is trained.	Extra retraining, larger models, or poor reuse of cached priors could increase footprint.
Economic	Student and supervision time, plus shared HPC use.	Lower marginal planning cost if the surrogate replaces repeated simulations.	Real deployment needs field validation and maintenance.
Social	Skills gained in wireless modelling, ML, HPC workflows and scientific reporting.	Faster UAV coverage planning for temporary, rural or emergency networks.	Bias toward represented city morphologies and risk of overclaiming simulated results.

5.2 Environmental impact

The main environmental cost during development was computation. The final residual GMM-head model line used a four GPU RTX 2080 Ti node for approximately 12 h; the distribution-first and prior calibration phase used additional GPU/CPU time. Across the full exploratory project, the total exceeded roughly 400 GPU hours. This is a one time research overhead and includes many failed or diagnostic experiments, not only the selected final run.

Several decisions reduced avoidable impact: reusing saved histories and frozen calibration artifacts, evaluating city-holdout metrics from stored checkpoints, and keeping the final priors frozen instead of recomputing them blindly inside every neural experiment. The final design is also relatively efficient at inference. On the Barcelona benchmark, the local MATLAB ray-tracing path took 99.89 s per full-resolution sample, while the complete generator path with masks, priors, neural prediction and export took 0.88 s. The sustainability advantage therefore depends on reuse: if the same calibrated generator is queried many times, the marginal cost is much lower than repeated ray tracing; if the model is retrained from scratch for every new scenario, that advantage weakens.

The life cycle interpretation is therefore important. During the research phase, the footprint is dominated by exploration: many failed or partially successful attempts were

necessary to discover that the final path should combine strong priors with a residual GMM-head model. During the use phase, the footprint is dominated by inference and by occasional recalibration. The project is most sustainable when the trained model, frozen priors, cached masks and exported checkpoints are reused across many planning queries. It would be less sustainable if the same design were scaled by simply increasing model size or retraining from zero for every city.

The final generator also improves reproducibility, which has an environmental side. A documented .npz output, explicit mask policy, and fixed hardware/runtime benchmark make it easier to compare changes without rerunning expensive ray tracing baselines. Reproducible experiments are not only a scientific requirement; they reduce wasted computation caused by ambiguous settings, hidden preprocessing changes or repeated debugging runs.

5.3 Economic and social impact

Economically, this work is not yet a product ready to take to market. Its value is as a technical prototype showing that a dense CKM can be approximated much faster than repeated ray tracing once the target city representation and calibration protocol are fixed. The main development cost was human time: literature review, dataset handling, training/debugging, analysis, supervision and writing. The compute cost is small compared with this time cost in an academic project, but it is still relevant because the full exploration required many GPU hours. In a planning workflow, a mature version of this surrogate could reduce repeated simulation runs and enable faster scenario studies over UAV heights or city layouts.

Socially, faster CKM prediction can support temporary coverage planning for public events, rural extensions, emergency response or infrastructure restoration. The direct users would be network planners; the indirect beneficiaries would be users receiving more reliable wireless coverage. The main social risk is overconfidence: a simulated city-holdout result should not be treated as certified field performance. Errors should continue to be reported by LoS/NLoS state, UAV-height bin and morphology group rather than hidden behind one global RMSE.

The social value is strongest when the tool is used as decision support rather than as an automatic decision maker. A fast surrogate can help planners compare candidate UAV heights, identify weak coverage zones, and decide where detailed simulation or measurement should be spent. It should not replace domain review: the model has been validated on simulated CKM cities and inherits the assumptions of the ray tracer, the topology representation, and the ground-mask definition.

There is also a fairness dimension in the morphology split. The model is tested on unseen cities, but the training distribution is still finite and urban morphologies are not equally represented. Dense high-rise, unusual street canyons, informal settlements, vegetation heavy areas or non European building patterns may behave differently from the dominant CKM regimes. For that reason, future deployment should report uncertainty or at least split diagnostics by morphology and height, instead of using the global average as the only quality indicator.

5.4 Ethical implications and SDGs

The main ethical obligation is honesty about validity. The dataset is generated by ray tracing, so the thesis demonstrates simulator-level generalization to unseen cities, not real-world deployment performance. A future operational system would require field measurements, calibration and human review, especially before public service or emergency use.

A second ethical issue is transparency. The final predictor is not a pure black-box CNN: the path loss and spread priors expose the geometric and morphological assumptions before the residual network is applied. This makes the system easier to audit, because a user can inspect whether an error comes from the LoS/NLoS mask, the prior, or the learned residual correction. That audit path is particularly relevant for emergency or infrastructure planning, where a wrong map can lead to misplaced confidence in predicted coverage.

The work does not process personal data, but operational radio planning can interact with sensitive contexts such as emergency response, public events or critical infrastructure. The generator should therefore be accompanied by clear metadata: input source, resolution, antenna height, mask generation mode, checkpoint version and known limitations. Keeping that metadata with the prediction is part of the ethical responsibility of making a fast tool: speed is useful only if the output remains traceable.

The project aligns most directly with SDG 9, because it develops a faster tool for wireless infrastructure planning; SDG 11, because it can support resilient urban and temporary connectivity; and SDG 13, because a validated surrogate can reduce repeated simulation energy. These contributions are indirect: the thesis provides a reproducible method and generator, not an operational network deployment by itself.

Overall, the sustainability conclusion is conditional but positive. The research process was computation heavy, but it produced a reusable generator that can replace many repeated full-resolution simulations once validated for a target domain. The strongest sustainable use case is therefore not “train a larger model”, but “reuse a documented, prior-anchored surrogate to decide where slower simulation or real measurement is actually needed”.

Conclusions and Future Work

6.1 Conclusions

This thesis demonstrates that dense Channel Knowledge Maps for UAV-enabled networks can be predicted with varying transmitter height, provided that the problem is treated as a structured propagation task rather than as generic image-to-image translation. The final system predicts path loss, delay spread, and angular spread on 513×513 city maps under a strict city-holdout protocol, with the UAV located at the map centre and its height changing from sample to sample. The work therefore answers the original practical question: a single surrogate can produce dense CKM targets for unseen cities while conditioning on continuous UAV height.

The central lesson is that the strongest model is not a pure black-box neural network. It is a hybrid system in which calibrated physical/statistical priors explain the stable part of the channel, and a neural residual model corrects what remains. This division of labour is what made the final result reliable: the priors capture radial propagation, visibility, height effects, building layout, and spread scale statistics; the learned model then focuses on local residual structure and distribution mismatch.

The finished pipeline meets the original numerical targets. The standalone path-loss prior achieved 1.9246 dB RMSE on the held-out validation+test split, and 1.9383 dB when reused as the frozen prior on the final 14-city unseen test split. The spread priors reached 28.10 ns delay-spread RMSE and 13.74° angular-spread RMSE on the same final test split. The final prior-anchored residual GMM-head model improved those anchors to 1.6519 dB, 26.5570 ns, and 11.3854° , respectively. These numbers are meaningful because the baseline being improved is already strong and train calibrated, not a weak baseline.

6.1.1 What the final result means

The clearest contribution is a working CKM predictor for variable height UAV maps. The final model receives geometry, masks, frozen priors, and a continuous height condition, and produces dense path loss, delay spread, and angular spread maps jointly. The height dependence is not cosmetic: the same city footprint changes its visibility structure, two-ray interference pattern, and spread regimes as the UAV altitude changes. The final architecture handles this by combining height-aware priors with a sinusoidal height embedding injected through FiLM conditioning.

The path-loss result is the strongest success case. Most LoS path loss structure is radial, height-dependent, and coherent enough that an enhanced two-ray prior can explain much of it before any neural correction is applied. The residual model still matters, but its role is narrower and more realistic: it improves local blocked streets, transition regions, and geometry to radio misalignments that the analytic prior cannot fully capture. Reducing the test RMSE from 1.9383 dB to 1.6519 dB after such a strong prior is therefore

a substantial result.

Delay spread and angular spread are harder to judge by a single RMSE. Their maps contain sparse high value regions, spike like behaviour, and long tailed outliers. Qualitatively, the final predictions usually recover the main spatial structure: low spread and high spread regions appear in the right general parts of the city, and the LoS/NLoS geometry is respected. The remaining mistakes are mostly local outlier failures. Some groups of pixels with unusually high DS or AS are missed, while some false positive high spread patches are predicted where the target remains moderate. This explains why the spread result is less uniform than the path-loss result, while still being useful: the model improves the global structure and many local regimes, but rare spread hot spots remain the hardest cases. These RMSE gains should not be read as universal typical error gains: delay-spread MAE increases, and angular spread MAE is nearly unchanged, so the spread claim remains an RMSE and structure claim rather than a calibrated error claim.

The most important methodological conclusion is that CKM prediction improves when the model head follows the target distribution. Earlier point regression models could produce plausible-looking maps while missing NLoS path loss modes or suppressing spread tails. The distribution-first stage fixed that diagnosis by using GMM style outputs and histogram aware evaluation. The final model keeps that lesson in its head design: it emits mixture parameters around calibrated anchors and exports their local weighted point prediction, while the final reported claim remains the deterministic RMSE improvement selected after the RMSE/MAE dominant fine-tuning stage.

6.1.2 Answers to the research questions

RQ1. Can a single deep learning surrogate predict per-pixel path loss, delay spread, and angular spread on 513×513 channel knowledge maps, under strict city-holdout evaluation, with errors compatible with the thesis targets? Yes, within the fixed CKM contract used in this thesis. The final model predicts PL, DS, and AS jointly on unseen cities and clears the original targets: 1.6519 dB path-loss RMSE, 26.5570 ns delay-spread RMSE, and 11.3854° angular-spread RMSE. The result should be understood under one carrier frequency, one centred UAV, valid ground receivers only, and city level train/validation/test separation.

RQ2. Which part of the prediction should be left to a calibrated analytic prior and which part should be left to the neural network, so that the hybrid model is both accurate and interpretable? The calibrated priors should handle the most predictable propagation structure: two-ray LoS path loss, height-binned LoS calibration, morphology aware NLoS corrections, and log-domain ridge models for DS and AS. These priors are strong enough to be useful on their own. The neural model is then valuable because it learns bounded residual corrections around them, reducing the frozen-prior test errors from 1.9383 dB, 28.1023 ns, and 13.7416° to 1.6519 dB, 26.5570 ns, and 11.3854° .

RQ3. How does the continuous UAV transmitter height (12 m to 478 m) affect the prediction error, and is the same architecture able to cover the full height range without per height retraining? Height affects the difficulty of the prediction through visibility, blocking and spread tail behaviour, so it must be represented explicitly. In the final model, height enters twice: through height-aware calibrated priors and through a sinusoidal FiLM conditioning path inside the neural residual model. Low-altitude cases

contain more blocked links, stronger transition zones and harder spread tails, while high-altitude samples are easier even though they are less represented in the training data. This indicates that the height trend is physical, not only a sampling artifact.

6.1.3 Main contributions

The first contribution is the final variable height CKM predictor itself: a single prior-anchored residual GMM-head model that jointly predicts path loss, delay spread, and angular spread for unseen city layouts while conditioning on continuous UAV height.

The second contribution is the calibrated prior family. The path-loss prior combines free space/radial structure, coherent two-ray LoS behaviour, height-dependent calibration, NLoS morphology features, and fallback regimes. The spread priors extend the same philosophy to DS and AS through log-domain ridge calibration, topology/visibility grouping, height regimes, and robust fallbacks. Together, these priors put a large amount of physical and empirical knowledge into a reproducible estimator that is already competitive before the neural residual is trained.

The third contribution is the distribution-first modelling lesson. The thesis shows that target histograms, LoS/NLoS marginals, and spike-plus-tail spread behaviour are not diagnostic extras; they determine which output head is appropriate. This is why GMM style residual prediction and histogram aware checks became part of the final design.

The fourth contribution is the evaluation contract: city-holdout splitting, ground-only masking, LoS/NLoS reporting, topology diagnostics, and UAV height stratified reporting. This contract is essential because random splits or unmasked building pixels can make the same model appear stronger without improving the deployment-relevant task.

6.1.4 Practical implications and scope

The practical implication is that future CKM systems should be compared under the same contract before architectural claims are made. A new backbone is useful only if it improves held-out city metrics after applying the same receiver mask, height conditioning, LoS/NLoS reporting, topology analysis, and prior baselines. Otherwise, a lower error may simply come from solving an easier problem.

The result is strongest as a reproducible surrogate for the CKM simulation contract used here. It is not a universal replacement for ray tracing across all frequencies, antenna patterns, or cities. It does show, however, that once the dataset contract is fixed, the slow simulator can be compressed into a fast and interpretable predictor. In the Barcelona runtime benchmark, pixel-level MATLAB ray tracing took 99.89 s per-sample, while the complete final generator path took 0.88 s, a $113.7\times$ reduction.

6.1.5 Negative results

The negative results are useful mainly as design evidence, so the detailed attempt history is kept in Appendix A rather than repeated in the main conclusion. The short lesson is that pure image translation, larger CNN backbones, topology experts, multi scale outputs, and uncertainty heads each exposed a different failure mode, but none removed the need for strong frozen priors and distribution-aware residual learning.

This is why the thesis ends with HARP-NET CKM: the final design is not just a larger neural network, but a bounded residual predictor built around the parts of the channel that can already be explained physically or statistically.

6.2 Future Directions

Short-medium term. The next step is to strengthen the current residual pipeline without changing the evaluation contract, then extend it toward frequency and radio configuration conditioning. Useful additions include transmitter-oriented mixup, explicit corridor losses, and controlled ablations of each residual head and loss weight. The current dataset fixes the carrier and antenna assumptions, while a more general outdoor CKM model should learn how path loss and spreads vary with frequency, bandwidth, antenna pattern, and possibly polarization. A natural dataset for this stage would simulate the same cities across several radio configurations so that the model can separate geometry effects from configuration effects.

A nearer term training direction is adaptation to new or under represented topology types. The per topology analysis suggests that dense high-rise, compact mid-rise, open low-rise, mixed urban layouts, and future morphology classes do not fail in exactly the same way. A practical continuation would therefore be to fine tune the final prior-anchored model on newly introduced urban typologies, or to train lightweight specialist heads on top of the shared backbone when only limited examples are available. This would test whether the model can recover local structure in topology regimes beyond the original training distribution, rather than only specialising to already defined classes.

For delay and angular spreads, another short to medium term direction is to optimize not only pixel-wise RMSE but also spatial structure. Spread maps can look physically coherent even when small shifts across quantized target boundaries are penalized harshly by RMSE. Future fine-tuning should therefore test structural objectives, for example gradient consistency, region wise correlation, SSIM like terms, LoS/NLoS boundary consistency, or topology aware ranking losses, so that the model preserves the correct spread patterns rather than only minimizing the average numeric error.

Medium to long-term. The final extension is field calibration. The path loss and spread priors are interpretable, which makes them suitable for simulator to reality correction using a limited number of real measurements. In the medium to long-term, this could lead to a CKM predictor that starts from ray-traced or prior generated maps, adapts to a new city with sparse measurements, and still reports uncertainty by LoS/NLoS regime, UAV-height bin, and topology class. Those safeguards should remain part of the method before considering operational use in emergency coverage, UAV network planning, or rapid deployment studies.

Bibliography

References used for the propagation models, neural-network baselines, training techniques, and evaluation methods cited in this thesis.

- [1] 3GPP. *Uncrewed Aerial System (UAS) Support in 3GPP; Stage 1*. Tech. rep. TS 22.125. Specification details page, accessed 21 May 2026. 3rd Generation Partnership Project, 2026. URL: <https://portal.3gpp.org/desktopmodules/Specifications/SpecificationDetails.aspx?specificationId=3545>.
- [2] Yong Zeng and Xiaoli Xu. "Toward Environment-Aware 6G Communications via Channel Knowledge Map". In: *IEEE Wireless Communications* 28.3 (2021), pp. 84–91. DOI: [10.1109/MWC.001.2000327](https://doi.org/10.1109/MWC.001.2000327).
- [3] Yong Zeng et al. "A Tutorial on Environment-Aware Communications via Channel Knowledge Map for 6G". In: *IEEE Communications Surveys & Tutorials* 26.3 (2024), pp. 1478–1519. DOI: [10.1109/COMST.2024.3364508](https://doi.org/10.1109/COMST.2024.3364508).
- [4] Xiaoli Xu and Yong Zeng. "How Much Data Is Needed for Channel Knowledge Map Construction?" In: *IEEE Transactions on Wireless Communications* 23.10 (2024), pp. 13011–13021. DOI: [10.1109/TWC.2024.3397964](https://doi.org/10.1109/TWC.2024.3397964).
- [5] Ç. Yapar et al. *Dataset of Pathloss and ToA Radio Maps With Localization Application*. arXiv:2212.11777. 2022. URL: <https://arxiv.org/abs/2212.11777>.
- [6] Çağkan Yapar et al. *The First Pathloss Radio Map Prediction Challenge*. arXiv:2310.07658. 2023. URL: <https://arxiv.org/abs/2310.07658>.
- [7] Fabian Jaensch, Giuseppe Caire, and Begüm Demir. *Radio Map Estimation: An Open Dataset with Directive Transmitter Antennas and Initial Experiments*. arXiv:2402.00878. 2024. URL: <https://arxiv.org/abs/2402.00878>.
- [8] Zijian Wu et al. *CKMImageNet: A Dataset for AI-Based Channel Knowledge Map Towards Environment-Aware Communication and Sensing*. arXiv:2504.09849. 2025. URL: <https://arxiv.org/abs/2504.09849>.
- [9] Theodore S. Rappaport. *Wireless Communications: Principles and Practice*. 2nd ed. Prentice Hall, 2002. URL: <https://www.pearson.com/en-us/subject-catalog/p/wireless-communications-principles-and-practice/P200000003199>.
- [10] Chia-Chuan Chiu et al. "Channel Modeling of Air-to-Ground Signal Measurement with Two-Ray Ground-Reflection Model for UAV Communication Systems". In: *2021 30th Wireless and Optical Communications Conference (WOCC)*. 2021, pp. 251–256. DOI: [10.1109/WOCC53213.2021.9603250](https://doi.org/10.1109/WOCC53213.2021.9603250).
- [11] COST Action 231. *Digital Mobile Radio Towards Future Generation Systems: Final Report*. Tech. rep. Luxembourg: Office for Official Publications of the European Communities, 1999. URL: <https://op.europa.eu/en/publication-detail/-/publication/f2f42003-4028-4496-af95-beaa38fd475f>.
- [12] 3GPP. *Study on Channel Model for Frequencies from 0.5 to 100 GHz*. Tech. rep. TR 38.901. Specification archive, accessed 15 May 2026. 3rd Generation Partnership Project, 2026. URL: https://www.3gpp.org/ftp/Specs/archive/38_series/38.901/.

- [13] A. Al-Hourani, S. Kandeepan, and S. Lardner. “Optimal LAP Altitude for Maximum Coverage”. In: *IEEE Wireless Communications Letters* 3.6 (2014), pp. 569–572. DOI: [10.1109/LWC.2014.2342736](https://doi.org/10.1109/LWC.2014.2342736).
- [14] W. Khawaja et al. “A Survey of Air-to-Ground Propagation Channel Modeling for Unmanned Aerial Vehicles”. In: *IEEE Communications Surveys & Tutorials* 21.3 (2019), pp. 2361–2391. DOI: [10.1109/COMST.2019.2915069](https://doi.org/10.1109/COMST.2019.2915069).
- [15] Abdul Saboor and Evgenii Vinogradov. *Millimeter-Wave UAV Channel Model with Height-Dependent Path Loss and Shadowing in Urban Scenarios*. arXiv:2511.10763. 2025. URL: <https://arxiv.org/abs/2511.10763>.
- [16] Evgenii Vinogradov et al. “Spatially Consistent Air-to-Ground Channel Modeling and Simulation via 3D Shadow Projections”. In: *2026 International Conference on Computing, Networking and Communications*. 2026. DOI: [10.1109/ICNC68183.2026.11416865](https://doi.org/10.1109/ICNC68183.2026.11416865).
- [17] P. Kyösti et al. *WINNER II Channel Models*. Tech. rep. IST-4-027756 WINNER II Deliverable, 2007. URL: <https://www.cept.org/files/8339/winner2%20-%20final%20report.pdf>.
- [18] Phillip Isola et al. “Image-to-Image Translation with Conditional Adversarial Networks”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2017, pp. 1125–1134. URL: <https://arxiv.org/abs/1611.07004>.
- [19] Ron Levie et al. “RadioUNet: Fast Radio Map Estimation With Convolutional Neural Networks”. In: *IEEE Transactions on Wireless Communications* 20.6 (2021), pp. 4001–4015. DOI: [10.1109/TWC.2021.3054977](https://doi.org/10.1109/TWC.2021.3054977). URL: <https://arxiv.org/abs/1911.09002>.
- [20] Ziyue Yang et al. *Group Equivariant Convolutional Networks for Pathloss Estimation*. arXiv:2511.17841. 2025. DOI: [10.48550/arXiv.2511.17841](https://doi.org/10.48550/arXiv.2511.17841). URL: <https://arxiv.org/abs/2511.17841>.
- [21] Ju-Hyung Lee et al. “PMNet: Large-Scale Channel Prediction System for ICASSP 2023 First Pathloss Radio Map Prediction Challenge”. In: *Proc. IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 2023, pp. 1–2. DOI: [10.1109/ICASSP49357.2023.10095257](https://doi.org/10.1109/ICASSP49357.2023.10095257). URL: <https://doi.org/10.1109/ICASSP49357.2023.10095257>.
- [22] Yuxuan Li et al. *RMTransformer: Accurate Radio Map Construction and Coverage Prediction*. arXiv:2501.05190. 2025. URL: <https://arxiv.org/abs/2501.05190>.
- [23] Mingran Sun et al. *WiCo-PG: Wireless Channel Foundation Model for Pathloss Map Generation via Synesthesia of Machines*. arXiv:2511.15030. 2025. URL: <https://arxiv.org/abs/2511.15030>.
- [24] Dong Yang et al. *FM-RME: Foundation Model Empowered Radio Map Estimation*. arXiv:2602.22231. 2026. URL: <https://arxiv.org/abs/2602.22231>.
- [25] Zhiyuan Liu et al. *A Fine-Grained 3D Radio Map Construction Paradigm with Ultra-Low Sampling Rates by Large Generative Models*. arXiv:2509.11571. 2025. URL: <https://arxiv.org/abs/2509.11571>.
- [26] Wangqian Chen and Juntao Chen. “Diffraction and Scattering Aware Radio Map and Environment Reconstruction Using Geometry Model-Assisted Deep Learning”. In: *IEEE Transactions on Wireless Communications* 23.12 (2024), pp. 19804–19819. DOI: [10.1109/TWC.2024.3487225](https://doi.org/10.1109/TWC.2024.3487225). URL: <https://arxiv.org/abs/2403.00229>.

- [27] Mukaram Shahid et al. *ReVeal: A Physics-Informed Neural Network for High-Fidelity Radio Environment Mapping*. arXiv:2502.19646. 2025. URL: <https://arxiv.org/abs/2502.19646>.
- [28] Borui Huang et al. "Transformer-Based Air-to-Ground mmWave Channel Characteristics Prediction for 6G UAV Communications". In: *Sensors* 25.12 (2025), p. 3731. DOI: [10.3390/s25123731](https://doi.org/10.3390/s25123731). URL: <https://www.mdpi.com/1424-8220/25/12/3731>.
- [29] Alex Kendall and Yarin Gal. "What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?" In: *Advances in Neural Information Processing Systems*. 2017. URL: <https://arxiv.org/abs/1703.04977>.
- [30] Taha Saleh et al. "Probabilistic Path Loss Predictors for mmWave Networks". In: *Proceedings of the IEEE 93rd Vehicular Technology Conference*. 2021, pp. 1–6. DOI: [10.1109/VTC2021-Spring51267.2021.9448967](https://doi.org/10.1109/VTC2021-Spring51267.2021.9448967). URL: <https://doi.org/10.1109/VTC2021-Spring51267.2021.9448967>.
- [31] Dolores Garcia Marti et al. "A Mixture Density Channel Model for Deep Learning-Based Wireless Physical Layer Design". In: *Proceedings of the 23rd International ACM Conference on Modeling, Analysis and Simulation of Wireless and Mobile Systems*. ACM, 2020, pp. 53–62. DOI: [10.1145/3416010.3423229](https://doi.org/10.1145/3416010.3423229). URL: <https://dSPACE.networks.imdea.org/handle/20.500.12761/886>.
- [32] Youngmin Lee, Xiaomin Ma, and Andrew S. I. D. Lang. *Modeling of Time-varying Wireless Communication Channel with Fading and Shadowing*. arXiv:2405.08199. 2024. DOI: [10.48550/arXiv.2405.08199](https://doi.org/10.48550/arXiv.2405.08199). URL: <https://arxiv.org/abs/2405.08199>.
- [33] Xiucheng Wang et al. "RadioDiff: An Effective Generative Diffusion Model for Sampling-Free Dynamic Radio Map Construction". In: *IEEE Transactions on Cognitive Communications and Networking* 11.2 (2025), pp. 738–750. DOI: [10.1109/TCCN.2024.3504489](https://doi.org/10.1109/TCCN.2024.3504489). URL: <https://arxiv.org/abs/2408.08593>.
- [34] Pavel Izmailov et al. "Averaging Weights Leads to Wider Optima and Better Generalization". In: *Proceedings of the 34th Conference on Uncertainty in Artificial Intelligence*. 2018. URL: <https://arxiv.org/abs/1803.05407>.
- [35] Ethan Perez et al. "FiLM: Visual Reasoning with a General Conditioning Layer". In: *Proceedings of the AAAI Conference on Artificial Intelligence*. 2018. DOI: [10.1609/aaai.v32i1.11671](https://doi.org/10.1609/aaai.v32i1.11671). URL: <https://arxiv.org/abs/1709.07871>.
- [36] Ali Saeizadeh et al. *AIRMap: AI-Generated Radio Maps for Wireless Digital Twins*. arXiv:2511.05522. 2025. URL: <https://arxiv.org/abs/2511.05522>.
- [37] Yuan Gao et al. *Effective Outdoor Pathloss Prediction: A Multi-Layer Segmentation Approach with Weighting Map*. arXiv:2601.08436. 2026. URL: <https://arxiv.org/abs/2601.08436>.
- [38] Ferdaous Tarhouni, Muneer AlZubi, and Mohamed-Slim Alouini. *Machine Learning-based Path Loss Prediction in Suburban Environment in the Sub-6 GHz Band*. arXiv:2510.00696. 2025. URL: <https://arxiv.org/abs/2510.00696>.
- [39] Zhijie Zhong et al. *PathFinder: Advancing Path Loss Prediction for Single-to-Multi-Transmitter Scenario*. arXiv:2512.14150. 2025. URL: <https://arxiv.org/abs/2512.14150>.
- [40] Stefanos Bakirtzis et al. *The First Indoor Pathloss Radio Map Prediction Challenge*. arXiv:2501.13698. 2025. URL: <https://arxiv.org/abs/2501.13698>.

- [41] ICASSP 2025 Indoor Radio Map Challenge. *The First Indoor Pathloss Radio Map Prediction Challenge: Results*. 2025. URL: <https://indoorradiomapchallenge.github.io/results.html>.
- [42] Wenlihan Lu et al. "SIP2Net: Situational-Aware Indoor Pathloss-Map Prediction Network for Radio Map Generation". In: *ICASSP 2025 - 2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 2025, pp. 1–2. DOI: [10.1109/ICASSP49660.2025.10890319](https://doi.org/10.1109/ICASSP49660.2025.10890319).
- [43] Bin Feng et al. "IPP-Net: A Generalizable Deep Neural Network Model for Indoor Pathloss Radio Map Prediction". In: *ICASSP 2025 - 2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 2025, pp. 1–2. DOI: [10.1109/ICASSP49660.2025.10890663](https://doi.org/10.1109/ICASSP49660.2025.10890663). URL: <https://arxiv.org/abs/2501.06414>.
- [44] Xin Li et al. "TransPathNet: A Novel Two-Stage Framework for Indoor Radio Map Prediction". In: *ICASSP 2025 - 2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. 2025, pp. 1–2. DOI: [10.1109/ICASSP49660.2025.10887939](https://doi.org/10.1109/ICASSP49660.2025.10887939). URL: <https://arxiv.org/abs/2501.16023>.
- [45] Zeyao Sun et al. *RadioPiT: Radio Map Generation with Pixel Transformer Driven by Ultra-Sparse Real-World Data*. arXiv:2512.01451. 2025. URL: <https://arxiv.org/abs/2512.01451>.
- [46] Andrea Goldsmith. *Wireless Communications*. Cambridge University Press, 2005. DOI: [10.1017/CB09780511841224](https://doi.org/10.1017/CB09780511841224).
- [47] X. Cai et al. *An Empirical Air-to-Ground Channel Model Based on Passive Measurements in LTE*. arXiv:1901.07930. 2019. URL: <https://arxiv.org/abs/1901.07930>.
- [48] T. Izydorczyk et al. "Angular Distribution of Cellular Signals for UAVs in Urban and Rural Scenarios". In: *Proc. EuCAP*. 2019. URL: <https://ieeexplore.ieee.org/document/8739910>.
- [49] Guillem Moreno Garcia. *CKMGenerator: Generator of Path Loss, Delay Spread and Angular Spread*. Public GitHub repository. Snapshot commit 8b6c5fd26c30. 2026. URL: <https://github.com/unworthyzeus/CKMGenerator> (visited on 05/06/2026).
- [50] Guillem Moreno Garcia. *Final_Code_TFG: Final Training and Evaluation Code*. Public GitHub repository. Snapshot commit 63d12a3d1a05. 2026. URL: https://github.com/unworthyzeus/Final_Code_TFG (visited on 05/06/2026).
- [51] Guillem Moreno Garcia. *TFGAllProgress_Tries_and_Attempts: Development History and Experimental Attempts*. Public GitHub repository. Snapshot commit 7c13efd51708. 2026. URL: https://github.com/unworthyzeus/TFGAllProgress_Tries_and_Attempts (visited on 05/06/2026).
- [52] Guillem Moreno Garcia. *TFG_ActualText_LaTeX: Thesis Source*. Public GitHub repository. Snapshot tag final-thesis-2026-05-15. 2026. URL: https://github.com/unworthyzeus/TFG_ActualText_LaTeX (visited on 05/06/2026).

Development Phases and Attempt Trace

The original appendix listed the internal attempt numbers in several long tables. For the reduced thesis, the compiled appendix keeps phases instead of eighty row level entries. The full attempt number map remains in `appendices.tex` and in the repository history, so the trace is not lost; it is simply no longer part of the main PDF.

Table A.1: *Compact attempt log grouped by methodological phase.*

Attempt range	Phase	Why it matters in the final thesis
1–22	Direct image-to-image baselines	Established the CKM pipeline, cGAN/U-Net feasibility, bilinear decoder fixes and multiscale supervision. The phase ended because visual realism did not imply physical RMSE.
23–40	Spread branches and mask reset	Added delay/angular experiments and, most importantly, the ground-only receiver mask. This changed the evaluation contract.
41–50	Prior residual PMNet phase	Showed that raw formulas were weak, train-only calibration helped, and residual learning was useful but still NLoS-limited.
51–66	PMHHNet and topology experts	Introduced six topology experts, height FiLM, high-frequency paths, obstruction channels and no-prior controls. LoS improved, NLoS remained the bottleneck.
67–75	Stabilisation and distribution diagnosis	Tested knife-edge channels, quad heads, uncertainty and histogram-focused diagnostics. The key result was the visible NLoS and spread-tail distribution mismatch.
76–77	Distribution-first models	GMM/spike-plus-tail heads showed that modelling the target distribution before pixel assignment could remove much of the collapse.
78–79	Frozen calibrated priors	Produced the final train-only path loss and spread priors, including height-aware LoS two-ray calibration and log-domain spread ridge calibration.
80	Final prior-anchored residual GMM model	Combined frozen priors with a shared residual GMM-head model and selected the RMSE-dominant fine-tuning stage reported in Chapter 4.

The main methodological lesson is that the project did not progress by simply making a larger network. The decisive changes were the evaluation contract, regime separation, distribution diagnostics and the return to strong frozen priors once their role had become clear.

A.1 Expanded phase notes

The compact table above is intentionally phase based, but the thesis still needs enough chronology to show that the final model was not chosen in one step. The following notes keep the development evidence that most directly explains the final design.

PMNet transfer check. A separate diagnostic placed official PMNet checkpoints on CKM samples. This confirmed that PMNet family models are relevant as architectural references but cannot be imported as ready-made solutions for this thesis contract. The training and target distributions are structurally different: CKM uses continuous UAV height, full 513×513 city maps, city-holdout evaluation, and valid ground receiver metrics. For that reason the PMNet result is kept as a transfer check in the appendix, not as a main method result.

Table A.2: *Decisive attempt milestones kept from the detailed attempt log.*

Attempt	Public role	Lasting lesson
22	Clean early image-to-image baseline	Bilinear decoding and multiscale supervision reduced visual artifacts, but direct image realism was still a weak proxy for calibrated radio-map error.
33–40	Ground mask reset	Building pixels had to be excluded from loss, metrics and visual error maps; scores before and after this contract are historical, not directly comparable.
41–45	First explicit prior-residual models	Raw formulas were weak, train-only calibration helped strongly, and residual learning was useful only when the prior carried real physical structure.
50	NLoS prior sandbox	Scalar NLoS corrections, tabular deltas and simple hybrid formulas did not solve blocked-link tails. The remaining error was spatial and distributional.
54–66	PMHHNet/topology experts	Height FiLM, physical channels and six topology experts improved the easy LoS regime, but NLoS stayed dominant even after substantial backbone engineering.
67–75	Stabilisation and histogram diagnosis	Knife-edge channels, quad heads and uncertainty improved diagnosis more than final RMSE. Histogram plots made the NLoS and spread-tail mismatch explicit.
76	Distribution-first PL GMM	Predicting the target distribution before assigning pixels gave the first convincing route for hard path-loss tails.
77	Distribution-first spread GMM	Delay and angular spread behaved like spike-plus-tail targets rather than smooth images, motivating distribution-aware output heads.
78–79	Frozen calibrated priors	The strongest non-neural evidence came from height-aware LoS calibration and log-domain spread ridge calibration, both fitted only on training cities.
80	Final prior-anchored residual GMM model	The selected system is a calibrated prior-estimator plus bounded learned residual distribution head, not a black-box direct regressor.

A.1.1 From image translation to receiver-valid evaluation

The first part of the project treated CKM prediction as aligned image-to-image regression. This was useful for building the dataset pipeline: the topology map, LoS/NLoS support, distance field and targets all live on the same 513×513 grid. The early cGAN and U-Net variants also exposed practical issues such as checkerboard artifacts, excessive adversarial pressure and the need for multiscale supervision.

The decisive correction was methodological rather than architectural. Once the receiver mask was enforced, the task became prediction on ground receivers only: buildings and padded pixels are not valid receiver locations. This is why the final thesis does not mix early scores with the final leaderboard. They remain useful as a development trace, but the valid comparison contract starts once the ground-mask is fixed.

A.1.2 Why the prior returned

The first physical prior attempts were not successful as formulas alone. A raw prior could be far worse than a neural model, especially in NLoS. However, train-only calibration made the prior informative enough to act as an anchor. This distinction matters: the prior is not presented as a magic replacement for learning, but as a structured baseline that gives the network the correct coordinate system.

The negative evidence was equally important. PMNet style residual models, topology experts and NLoS focused objectives improved parts of the map, but the hard tails survived. The project therefore moved away from simply enlarging the encoder and toward separating regimes, modelling distributions and freezing strong calibrated estimators.

A.1.3 Why Try 80 is the endpoint

The last attempts are the only ones that form the final methodology. Try 78 provides the path-loss prior, with a height-aware LoS branch and calibrated NLoS residual structure. Try 79 provides delay-spread and angular-spread priors with separate LoS/NLoS, height and topology calibration.

Try 80 then learns bounded residual corrections over these frozen priors with a shared multi-task GMM-head model.

In other words, the final route is evolutionary but not arbitrary: direct regression established feasibility, masked evaluation fixed the contract, PMHHNet identified the remaining NLoS bottleneck, distribution-first models showed why tails mattered, and the frozen priors made the final residual model small enough to improve rather than overwrite the physical estimator.

A.1.4 What was moved out of the main chapters

The distribution-first PL and spread experiments remain part of the thesis evidence, but their detailed per-expert tables are kept here conceptually rather than in Chapter 4. The main chapters use them only for the lesson they provide: direct neural maps can look plausible while averaging away rare tails, and the final model should therefore keep calibrated priors and a distribution-aware residual head. This keeps the thesis shorter while preserving the negative evidence that motivated the final design.

Representative Final Model Panels

The full generated panel set contains one candidate for each (open, mixed, dense) $\times$ (low, mid, high) audit group. The expanded reduced thesis keeps three full-size qualitative panels: open/low, mixed/mid and dense/mid. The remaining panels and their manifest are kept in `img/thesis_figures/try80_appendix_panels/` and can be regenerated with the panel-generation script. The removed open/high Nazareth panel is not used because its angular-spread map is almost flat and the robust colour scale collapses to a misleading near-zero range.

Table B.1: *Representative final-model panel manifest.*

Audit group	City	Sample	Status	Reason
open/low	Segovia	sample_13213	PDF	Open low-altitude case with visible geometry and residual correction.
mixed/mid	Nice	sample_10828	PDF	Mixed morphology and mid-height regime.
dense/mid	Vancouver	sample_15265	PDF	Dense held-out test morphology; also used as the main paper/TFG qualitative panel.

Each panel has three columns (PL, DS and AS) and six rows: ground truth, frozen prior, final prediction, support/context, prior error and model correction. The last two rows are the diagnostic rows: a useful residual model should correct the sign and location of the prior error.




Figure B.1: Final residual GMM-head inspection panel for an open, low-antenna validation sample in Segovia.

CHAPTER B. REPRESENTATIVE FINAL MODEL PANELS




Figure B.2: Final residual GMM-head inspection panel for a mixed, mid-antenna validation sample in Nice.




Figure B.3: Final residual GMM-head inspection panel for a dense, mid-antenna held-out test sample in Vancouver.

CKMGenerator Interface and Reproducibility Tool

This appendix documents the standalone CKMGenerator delivered with the final prior-anchored residual GMM-head pipeline. The complete user manual belongs in the repository README, but the thesis keeps the scientific and reproducibility contract: which inputs are accepted, how they are converted to the model grid, how receivers and LoS/NLoS masks are defined, what the tool exports, and which performance switches do not change the mathematical model.

C.1 Purpose and source record

CKMGenerator is a deployment wrapper around the final thesis pipeline. It accepts topology inputs, prepares them with the same geometric conventions used for training and evaluation, computes or imports LoS/NLoS masks, evaluates the frozen calibrated path loss, delay-spread and angular-spread priors, optionally runs the final neural residual model, and exports the numerical arrays used for reproducible analysis. The Streamlit interface is intended for inspection and demonstrations; the command-line interface uses the same backend and is intended for scripted runs.

The reproducibility record is split across public repositories. The standalone generator source is in CKMGenerator; final training and evaluation code is in Final_Code_TFG; the broader experimental trace is preserved in TFGAllProgress_Tries_and_Attempts; and the thesis source is kept in TFG_ActualText_LaTeX [49, 50, 51, 52]. The final checkpoint is loaded from models/best_model.pt by default, or from an explicit `--checkpoint` path.

C.2 Interface and execution modes

Figures C.1–C.4 show the retained CKMGenerator screenshot set. The control panel is shown as a single cropped sidebar capture so the PDF documents the actual interface without wasting a full page on the blank app workspace.

The same backend can be run from the terminal. A typical HDF5 command is:

```
python ckm_generate.py ^
  --input inputs\samples\Abidjan_sample_00001.h5 ^
  --run-name abidjan_check ^
  --mixed-precision
```

For image-only topology inputs, the transmitter height must be supplied unless it is stored in a structured sidecar file:

```
python ckm_generate.py ^
  --input C:\path\topology_map.png ^
  --height 56.65 ^
  --topology-max-m 120 ^
  --meters-per-pixel 1.0 ^
  --mask-source auto ^
  --mixed-precision
```


Local input path

Or upload topology / data / HDF5

Upload

8GB per file • PNG, JPG, TIF, BMP, NPY, ...

Antenna height (m)

56,00

Device

auto

Runtime: torch=2.4.1+cpu | selected=directml

Rx height is fixed at 1.5 m for the CKM calibration and Try 78/79/80 priors.

Mask used by priors/model

auto

☒ Run Try 80 model

Model batch size

1

☐ CUDA mixed precision

☒ Save numeric arrays (.npz)

☒ Save LoS/NLoS mask PNGs

☒ Save visual map PNGs

Image max height (m)

90,00

☒ Image pixels are already metres

Model batch size

1

☐ CUDA mixed precision

☒ Save numeric arrays (.npz)

☒ Save LoS/NLoS mask PNGs

☒ Save visual map PNGs

Image max height (m)

90,00

☒ Image pixels are already metres

Image max height (m)

90,00

☒ Image pixels are already metres

Distance between pixels (m)

1,00

HDF5 city filter

HDF5 sample filter

Output folder


New run in outputs/runs

Run name


streamlit

Generate

Figure C.1: CKMGenerator Streamlit controls. The sidebar records the assumptions used by the generator: receiver height fixed at 1.5 m, variable antenna height, mask selection, optional Try 80 inference, mixed-precision availability, pixel scale and export format.



(a) Application run summary, resolved mask policy and exported sample metadata preview.




(b) Joint path loss, delay-spread and angular-spread prediction output.

Figure C.2: CKMGenerator application output after running one HDF5 sample.



(a) Saved visual products: masks, final prediction maps and joint panel.



(b) Generated output tree with masks, priors, predictions, arrays and metadata folders.

Figure C.3: CKMGenerator exported visual files and directory structure.

Metadata and numeric array contents

```
{
  "sample_id": "Abidjan_sample_00001",
  "antenna_height_m": 56.647,
  "mask_source": "provided",
  "requested_mask_source": "auto",
  "topology_class_6": "dense_block_midrise",
  "antenna_bin": "low_ant",
  "npz_contains": [
    "topology",
    "ground_mask",
    "los_mask",
    "nlos_mask",
    "prior_path_loss",
    "prior_delay_spread",
    "prior_angular_spread",
    "pred_path_loss",
    "pred_delay_spread",
    "pred_angular_spread"
  ]
}
```

Figure C.4: CKMGenerator metadata and numerical array archive. The .npz file is the authoritative output for reproducible analysis; PNG files are inspection artifacts.

C.3 Input contract and grid fitting

The generator accepts raster topology images (.png, .jpg, .tif, .bmp), numeric arrays (.npy, .npz, .csv, .txt, .json) and HDF5 files. HDF5 and NPZ inputs may provide topology, LoS/NLoS masks and antenna height directly. Raster inputs normally store intensities rather than metre-valued building-height values, so the selected maximum image height is used to convert pixel values to metres unless the option `image-values-are-metres` is enabled.

All inputs are fitted to the 513×513 grid used by the final model. The model grid represents 1 m per-pixel, with the transmitter antenna located at the centre of the map and with variable UAV height. If an input is supplied with another physical spacing, `-meters-per-pixel` is used before the centred crop or padding step. Topology arrays are resampled with bilinear interpolation; masks are resampled with nearest-neighbour interpolation and then binarised again.

The receiver convention is the same one used throughout the thesis. Ground pixels are zero-height topology pixels, and every ground pixel is a receiver at $h_{\text{rx}} = 1.5$ m. Building pixels and padded outside pixels are not receiver locations. The receiver-valid mask is therefore

$$m_{\text{ground}}(x) = \mathbf{1}\{h_{\text{topology}}(x) \leq 10^{-6}\}.$$

This convention matters because the model predicts dense maps only over valid ground receivers, not over roofs or artificial padding.

C.4 Mask, prior and neural-model path

The model consumes one LoS mask and one NLoS mask as input channels, and the same masks are also used by the priors. CKMGenerator supports three policies: `provided` requires masks stored in the input or sidecar files; `generated` ray-casts masks from topology and antenna height;

and auto uses provided masks when present and otherwise generates them. For CKM HDF5 samples, the NLoS support is derived from the raw LoS mask and the ground-mask,

$$m_{\text{LoS}} = m_{\text{LoS,raw}} m_{\text{ground}}, \quad m_{\text{NLoS}} = (1 - m_{\text{LoS,raw}}) m_{\text{ground}}.$$

The multiplication by m_{ground} is essential: otherwise building pixels would be counted as NLoS receiver pixels.

Generated masks make topology-only inference possible, but they are not claimed to be bit-identical to the CKM ray-tracing masks. They are a deployment fallback that preserves the model interface when a user does not have the original HDF5 mask fields.

After preprocessing, the tool evaluates the frozen priors from the final pipeline. These are the same height-aware, calibrated priors described in Chapter 3: a LoS/NLoS path-loss prior with centre-antenna geometry and antenna-height calibration, and spread priors with separate calibrated structure for delay spread and angular spread. When neural inference is enabled, the prepared topology, masks, priors and height context are passed to the final residual GMM-head model. The generator therefore changes neither the model definition nor the validation contract; it packages the final method for reproducible use.

C.5 Outputs and performance controls

Each run writes to a timestamped output folder, or to the explicit folder passed with `-out`. The authoritative output is the NumPy archive in `arrays/`. PNG files are optional visual inspection products.


Table C.1: *CKMGenerator artifacts and controls retained in the thesis appendix.*

Area	Reproducibility detail
Numerical archive	Stores topology, ground-mask, LoS/NLoS masks, generated or reference masks when available, PL/DS/AS priors, PL/DS/AS predictions and metadata on the final 513×513 grid.
Metadata	Records sample id, antenna height, source file, checkpoint path, requested and resolved mask policy, topology class, antenna bin, save options, timings and generated paths.
Visual outputs	Optional PNG exports include binary masks, prior maps, prediction maps and the joint prediction panel used for inspection.
Runtime	CUDA is preferred; CPU and DirectML are fallback paths. The <code>ckm_doctor.py</code> check verifies PyTorch, CUDA/DirectML availability, checkpoint presence and memory before inference.
Mixed precision	CUDA autocast can be enabled for speed. It was checked to change RMSE only at a negligible scale compared with validation error.
Batch speed	Disabling visual map PNGs, mask PNGs and compressed arrays avoids Matplotlib rendering and compression overhead while preserving the numerical arrays.


The benchmarked fast path reported in Chapter 4 uses vectorized LoS/NLoS mask generation, torch prior evaluation, mixed-precision neural inference and default numerical `.npz` export after the checkpoint has already been loaded. These optimizations affect runtime and storage, not the scientific definition of the generator.

Additional Results and Qualitative Diagnostics

This one-page appendix keeps only the two diagnostics that are useful to audit the final interpretation. Full split bars, all-city subexpert tables, the full qualitative gallery and long derivations remain in the source archive and generated image folders.



(a) *Full non-residual DL GMM-head model for PL. Thick curves show summed LoS/NLoS GMM densities; faint dotted curves show the individual Gaussian components.*



(b) *Spread-target quantisation: predictions are continuous while the released targets contain staircase-like and sparse integer-valued structures.*

Figure D.1: Compact additional diagnostics. The GMM plot exposed different LoS/NLoS marginals and motivated distribution-aware heads; the quantisation panel records the target-storage limitation discussed in Chapter 4.

Additional source material. The nine-panel gallery, row-level 80-attempt log, CKMGenerator screenshot manual, all-city subexpert tables and split bars are omitted from this compact appendix but remain in `appendices.tex` and the generated figure folders. The Try 78/79/80 mathematical derivations are not repository-only: they are included in Chapter 3 through

prior_detail_try78.tex, prior_detail_try79.tex and prior_detail_try80.tex, which also remain as separate source files for traceability.